

Recurrent Neural Networks (RNNs)

Michel RIVEILL
michel.riveill@univ-cotedazur.fr

Motivation


- ▶ Humans don't start their thinking from scratch every second
 - ▶ Thoughts have persistence
- ▶ Traditional neural networks can't characterize this phenomena
 - ▶ Ex: classify what is happening at every point in a movie
 - ▶ How a neural network can inform later events about the previous ones
- ▶ **Recurrent neural networks** address this issue
 - ▶ Some applications
 - ▶ NER - Naming Entity Recognition
 - Same word may have a different label depending on the context.
 - **Apple** CEO Tim Cook eat an **apple**
 - ▶ Forecasting - Time-series Prediction
- ▶ **How?**
 - ▶ **Add state to artificial neurons**

What are RNNs?

- ▶ Main idea is to make use of sequential information
- ▶ How RNN is different from neural network?
 - ▶ Vanilla neural networks (MLP) **assume** all inputs and outputs are independent of each other
 - ▶ But for many tasks, that's a very bad idea
- ▶ What RNN does?
 - ▶ Perform the same task for every element of a sequence (that's what **recurrent** stands for)
 - ▶ Output depends on the previous computations!
- ▶ Another way of interpretation – RNNs have a “**memory**”
 - ▶ To store previous computations


From vanilla NN to recurrent NN

- ▶ Vanilla cell
 - ▶ $y = F(U.X)$




From vanilla NN to recurrent NN

- ▶ Vanilla cell
 - ▶ $y = F(U.X)$
- ▶ Recurrent cell → use 2 weights matrix
 - ▶ Add an internal variable: h
 - ▶ The output depends to the current entry and the previous internal variable:
 - ▶ $h_t = F(W.h_{t-1} + U.X_t)$
 - ▶ Could be rewritten on $h_t = F(W.[h_{t-1}, X_t])$




From vanilla NN to recurrent NN

- ▶ Vanilla cell
 - ▶ $y = F(U.X)$
- ▶ Recurrent cell → use 2 weights matrix
 - ▶ $h_t = F(W.[h_{t-1}, X_t])$
- ▶ Recurrent layer, step by step
 - ▶ at each time step
 - ▶ A new entry is being supplied
 - ▶ And a new output (h_t) is calculated using:
 - The new input X_t
 - The output of the previous step h_{t-1}
 - ▶ $h_1 = F(W.[h_0, X_1])$
 - ▶ $h_2 = F(W.[h_1, X_2])$
 - ▶ $h_3 = F(W.[h_2, X_3])$
 - ▶ ...




From vanilla NN to recurrent NN

- ▶ Vanilla cell
 - ▶ $y = F(U.X)$
- ▶ Recurrent cell
 - ▶ $h_t = F(W.[h_{t-1}, X_t])$
- ▶ Recurrent neural networks are “unrolled” programmatically during training and prediction
 - ▶ All neurons share the same weight matrix




Remember

From




To




Remember


From



To

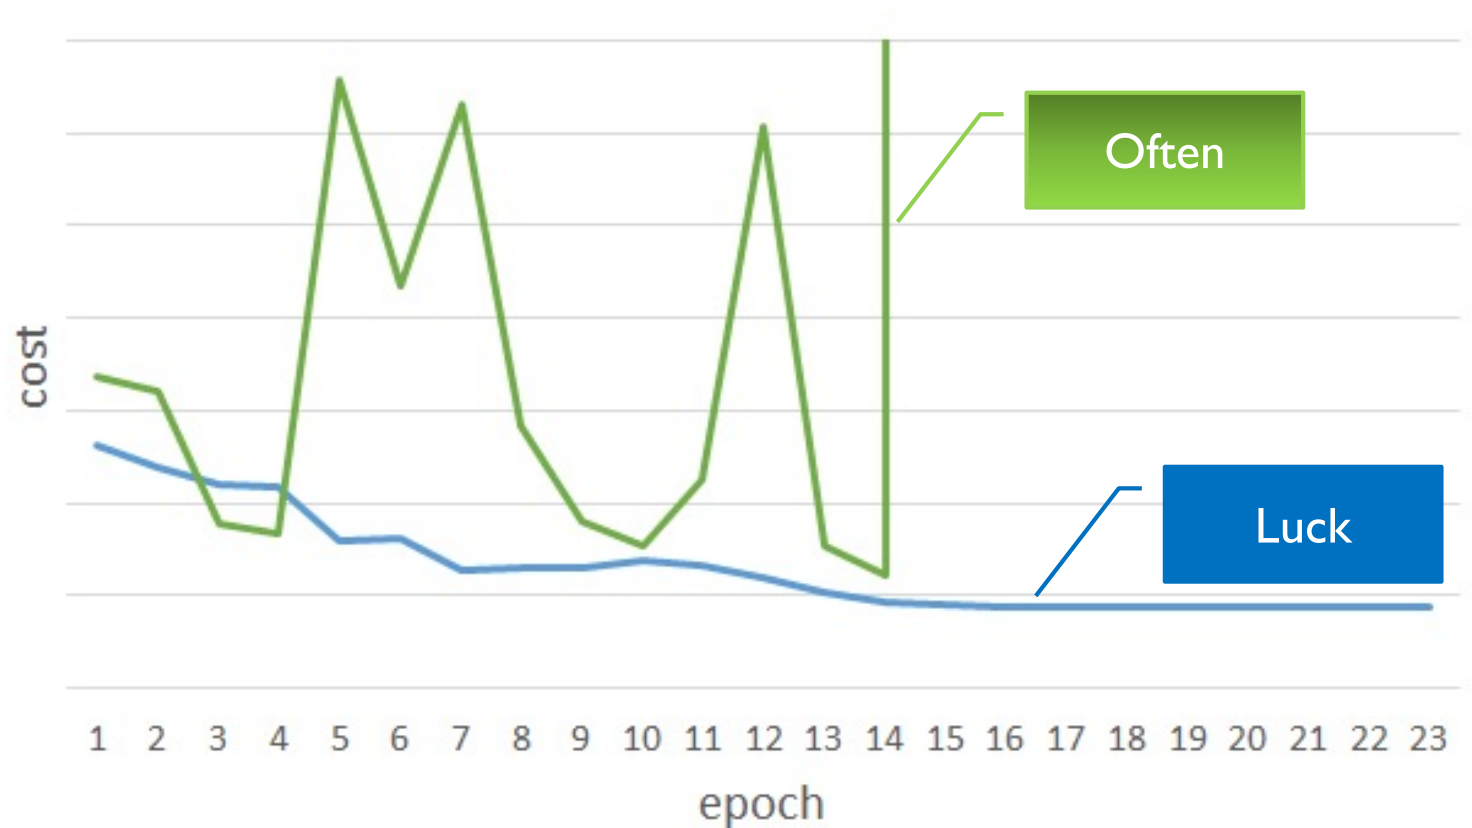


RNN in action



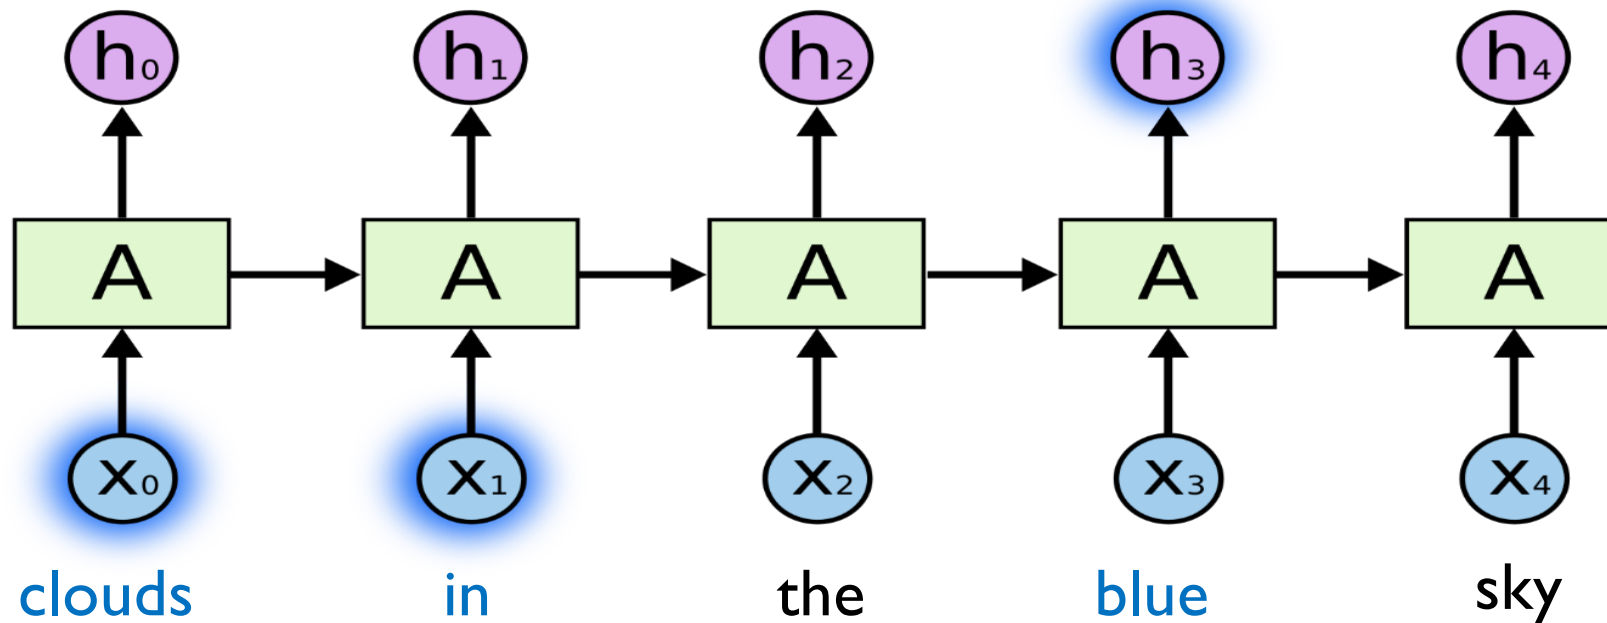
Problems with naive RNN

- ▶ RNNs do not learn easily
- ▶ Unfolding the network for learning leads to vanishing gradient problems!




From vanilla RNN...

- ▶ The context is close to the word to be predicted
 - ▶ Few iterations separate them.
 - ▶ No problem



... to LSTM (Long Short Term Memory)

- ▶ The context is far from the word to predict
 - ▶ Many iterations separate them!
 - ▶ Possible gradient problem




I grew up in France...


I speak French...

From Vanilla to LSTM Cells

- ▶ It is necessary to prevent the gradient from disappearing...
- ▶ Normally, the network memory is
 - ▶ $h_t = \tanh(W \square [h_{t-1}, x_t])$
 - ▶ Involves a single level of processing
 - ▶ Creating the risk of the evanescent gradient.




Dealing with the vanishing gradient problem → LSTM cell



(crédit : image modifiée de Michaël Nguyen)


LSTM cell

- ▶ Cellule composée de trois “portes” : ce sont des zones de calculs qui régulent le flot d’informations (en réalisant des actions spécifiques)
 - ▶ Forget gate (porte d’oubli)
 - ▶ Input gate (porte d’entrée)
 - ▶ Output gate (porte de sortie)
- ▶ Hidden state (état caché)
- ▶ Cell state (état de la cellule)
 - ▶ Like residual

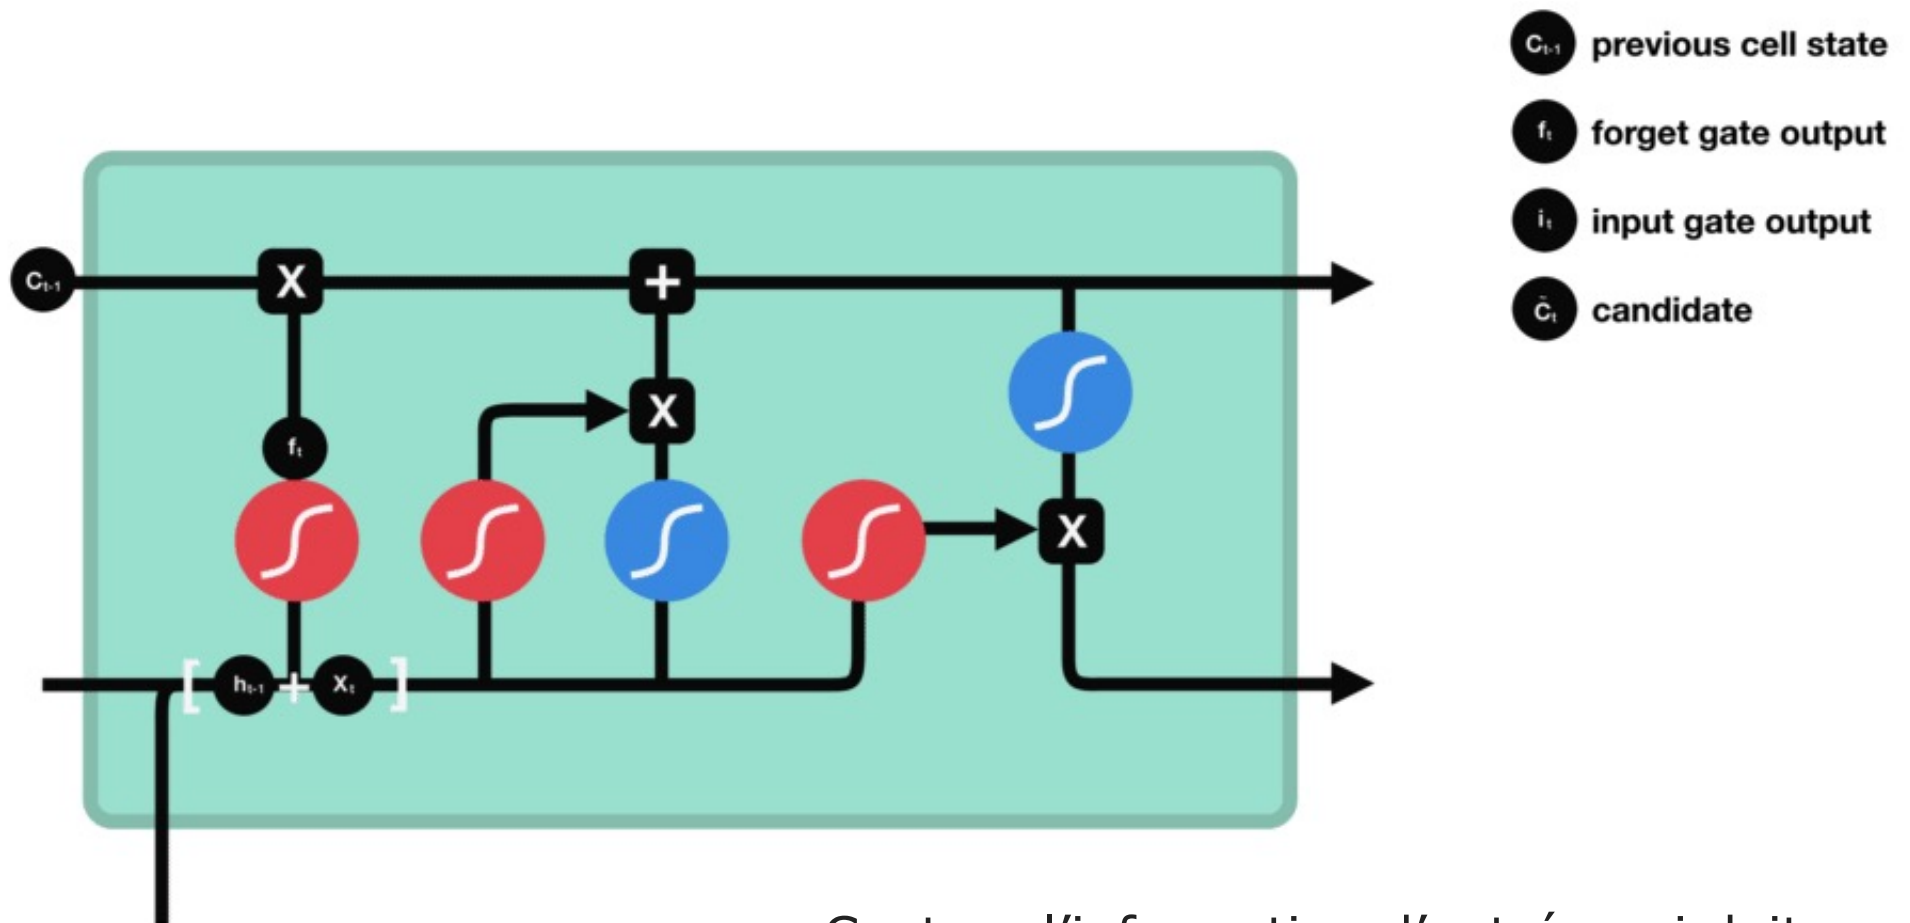


(crédit : image modifiée de Michaël Nguyen)

LSTM cell (porte oubli / forget get)

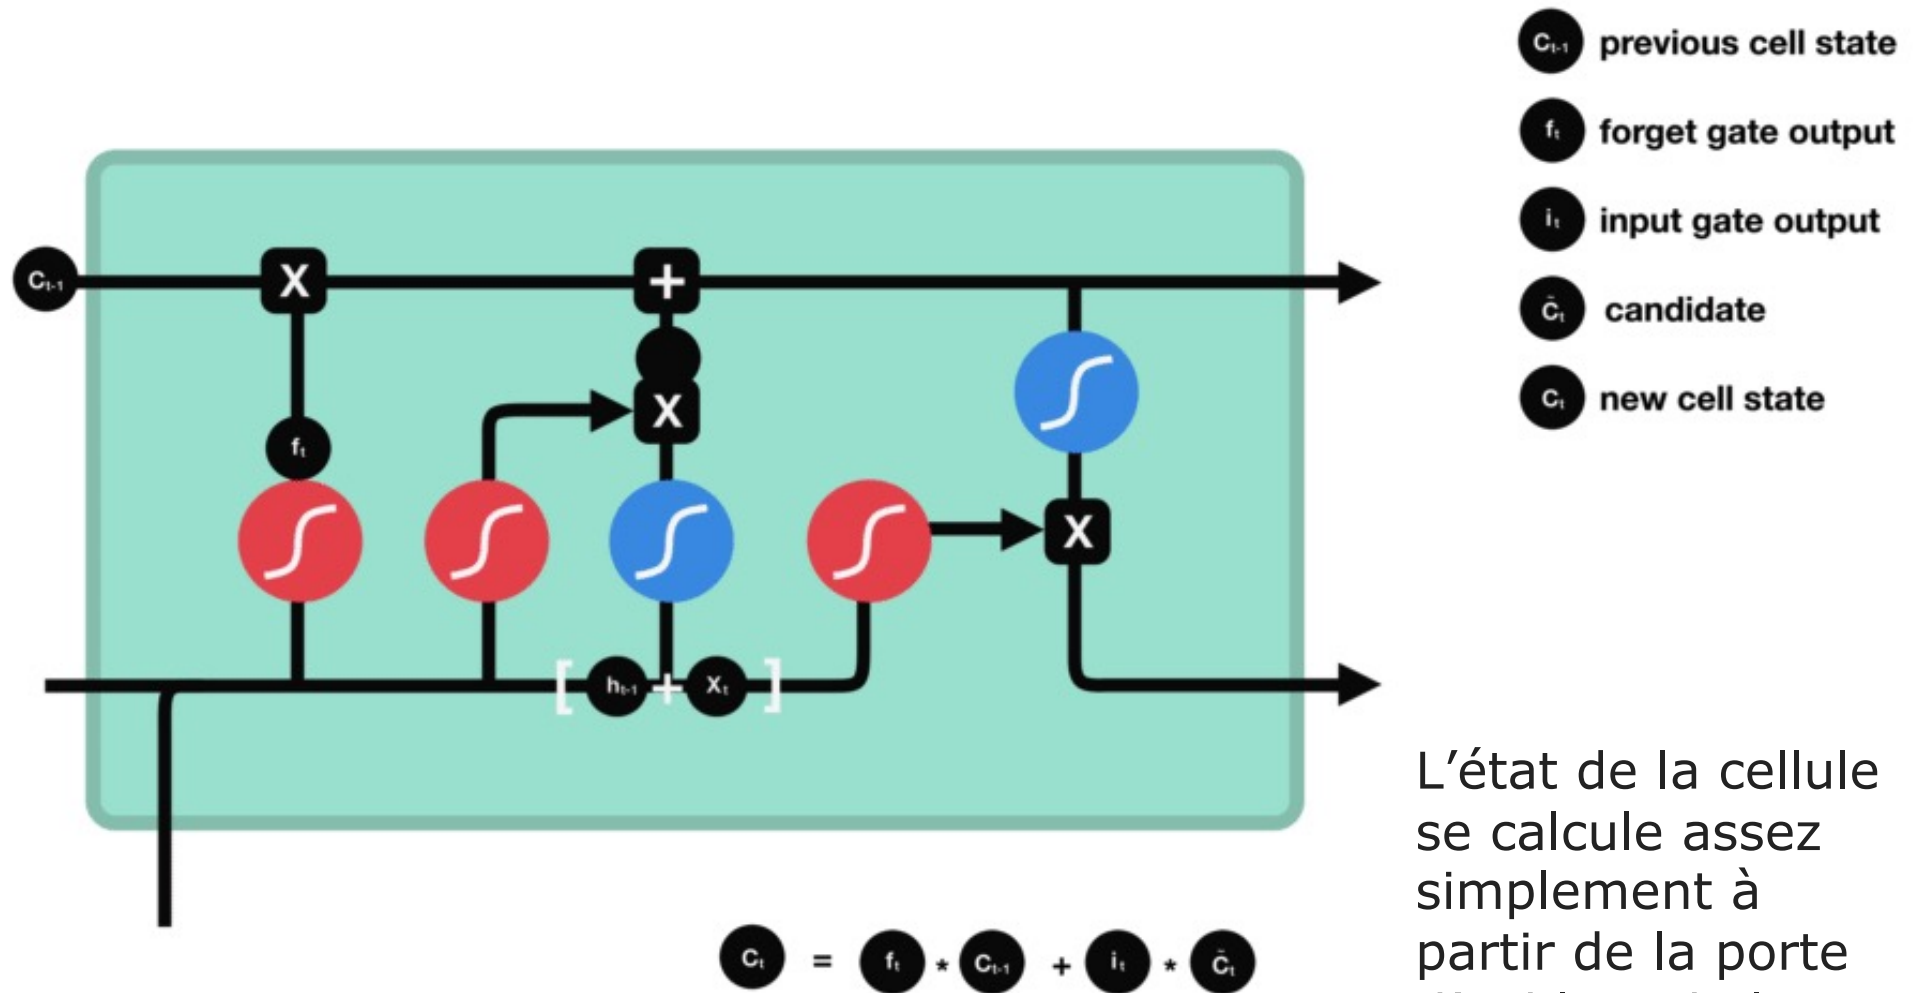


LSTM cell (porte entrée / input get)




Capture l'information d'entrée qui doit être incluse dans l'état de la cellule

LSTM cell (état de la cellule / cell state)



L'état de la cellule se calcule assez simplement à partir de la porte d'oubli et de la porte d'entrée

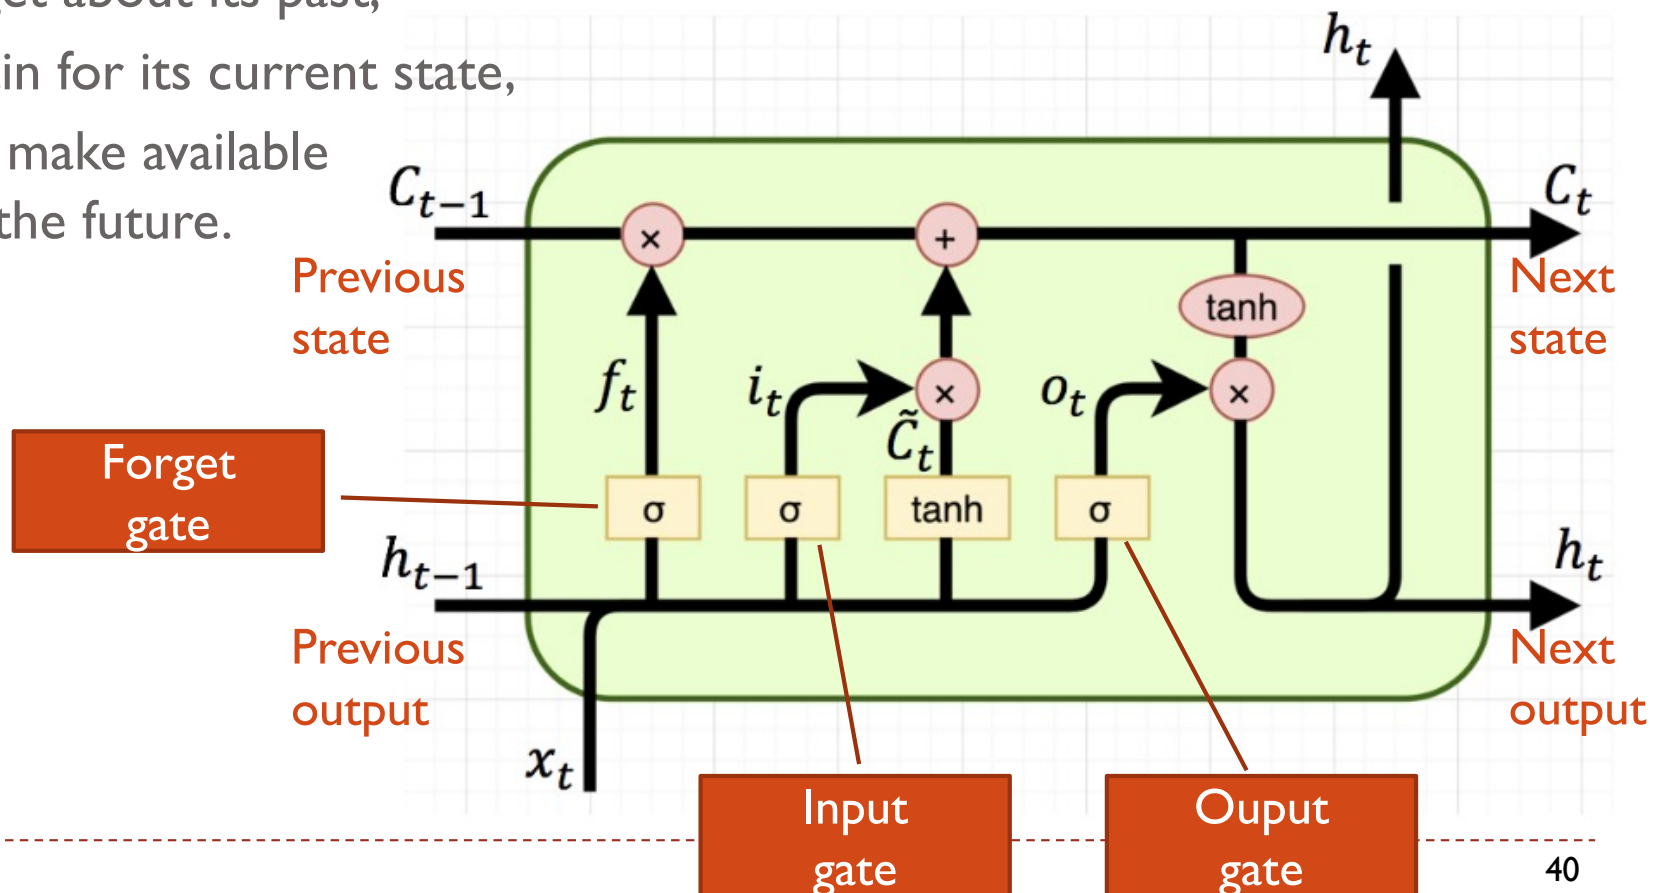
LSTM cell (porte de sortie / output gate)



La porte de sortie décide quel sera le prochain état caché. Il contient des informations sur les entrées précédentes du réseau et sert aux prédictions.

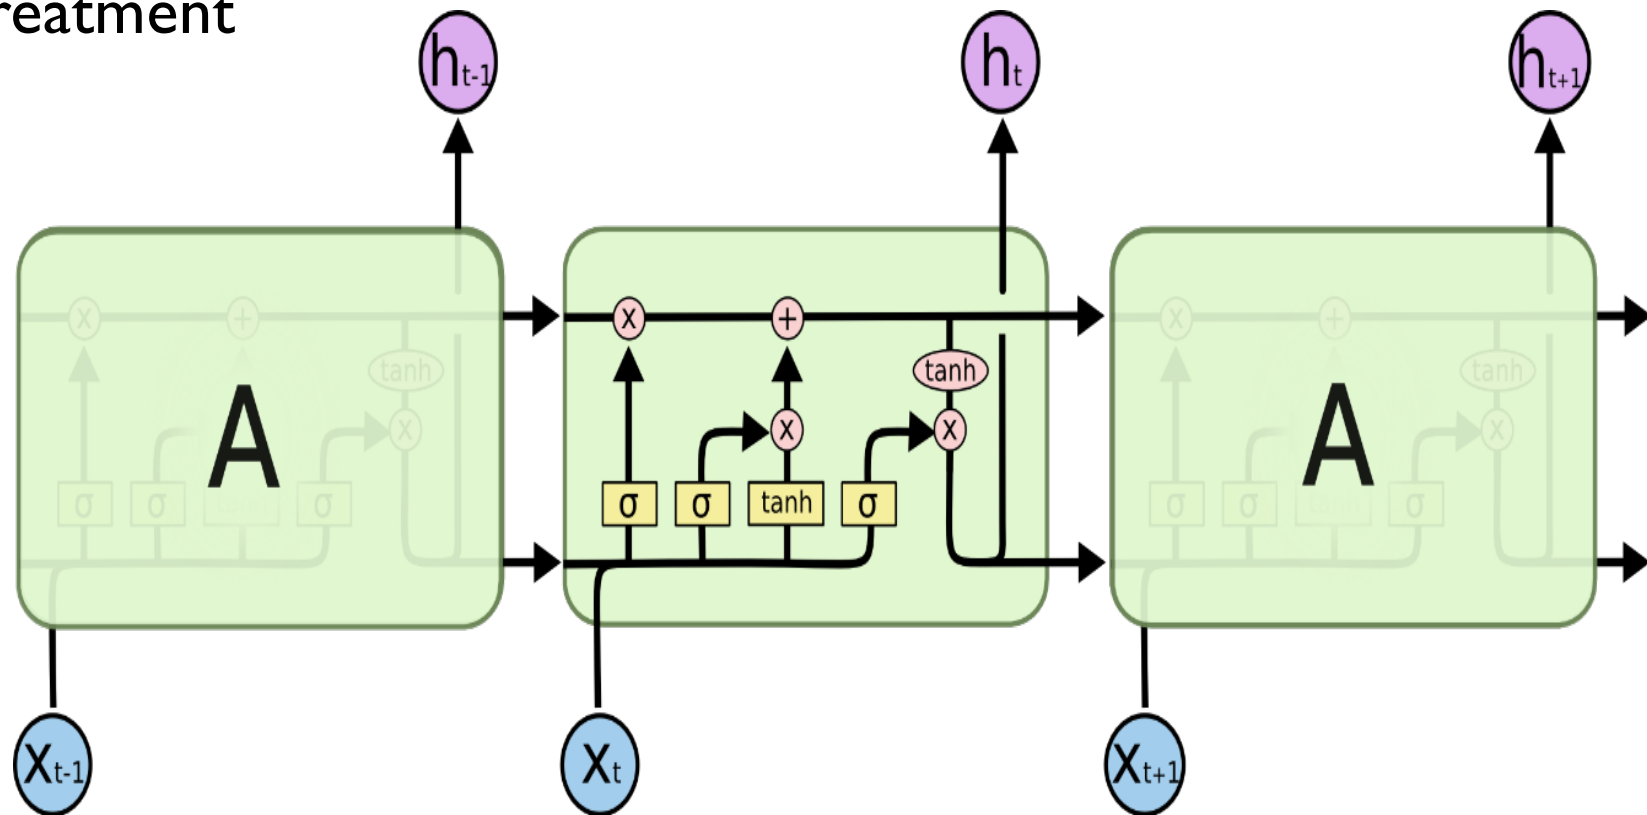
LSTM Cells

- ▶ Adds a context memory that affects the information flow and its processing (cell state).
- ▶ Three gates decide what a cell should
 - ▶ forget about its past,
 - ▶ retain for its current state,
 - ▶ and make available for the future.



LSTM Cells

- Concretely, recurrence in an LSTM cell involves 4 levels of treatment



Neural Network
Layer

Pointwise
Operation

Vector
Transfer

Concatenate

Copy



Naïve RNN vs LSTM

▶ Naïve RNN

- ▶ Reuse at each step the previous output

▶ LSTM

- ▶ At each step 3 gate control the use of Input value, Cell state and previous output




c changes slowly → c^t is c^{t-1} added by something

h changes faster → h^t and h^{t-1} can be very different

GRU – gated recurrent unit

- ▶ GRU = a light LSTM Cell



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$


$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$


- It combines the **forget** and **input** into a single **update gate**.
 - It also **merges** the **cell state** and **hidden state**.
- This is simpler than LSTM.

Bi-directional RNNs

- ▶ RNNs can process the input sequence in forward and in the reverse direction



- Popular in speech recognition, could be used also with text




RNN cell in Keras




Keras Long Short-Term Memory Cell

- ▶ **from tf.keras.layers import LSTM**
- ▶ Main params
 - ▶ **Units:** dimension of output space
 - ▶ **return_sequences:** True or False
 - ▶ If False return only the last output
 - ▶ If True return the full sequence of the output sequence
 - Output sequence = hidden state (the vocabulary change regarding documentation)
 - ▶ **return_state:** True or False
 - ▶ If True return 3 values
 - The full output sequence or only the last one (depend on return_sequences)
 - The last output sequence
 - The cell state
 - ▶ If False return nothing
 - ▶ **stateful:** True or False
 - ▶ If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
 - ▶ You have to put **shuffle=False** in a fit method




return_sequences = False



return_sequences = True

with `return_sequences=False`,
Dense layer is applied only
once at the last cell

If `return_sequences=True`
Dense layer is applied to every
timestep



Keras Gated Recurrent Unit Cell

- ▶ **from keras.layers import GRU**
- ▶ Main params (**similar to LSTM**)
 - ▶ **Units:** dimension of output space
 - ▶ **return_sequences:** True or False
 - ▶ If False return only the last output
 - ▶ If True return the full sequence of the output sequence
 - Output sequence = hidden state (the vocabulary change regarding documentation)
 - ▶ **return_state:** True or False
 - ▶ If True return 3 values
 - The full output sequence or only the last one (depend on return_sequences)
 - The last output sequence
 - The cell state
 - ▶ If False return nothing
 - ▶ **stateful:** True or False
 - ▶ If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
 - ▶ **You have to put shuffle=False in a fit method**

A basic example

- ▶ `inputs = Input(shape=(SEQUENCE_SIZE,))`
- ▶ `embedding = Embedding(VOCABULARY_SIZE,
 EMBEDDING_SIZE,
 input_length=SEQUENCE_SIZE)(inputs)`
- ▶ `output = LSTM(16, return_sequences=False,
 activation='relu')(embedding)`
- ▶ `predictions = Dense(nb_classes,
 activation='softmax')(output)`
- ▶ **Fit by batch**
 - ▶ `Model.fit(X, y, ...)`. ← all item have the same length
- ▶ **Fit by item**
 - ▶ `For i in range(len(X)):` ← could be different length
 - ▶ `Model.fit(X[i], y[i], ...)`

Some use of RNN

→ Text Classification / Sentiment analysis

Affect a label to a text

- ▶ Classify a

- ▶ restaurant review from Yelp!

- ▶ movie review from IMDB

- ...

- as positive or negative

- ▶ Inputs:

- ▶ Multiple words, one or more sentences


- ▶ Outputs:

- ▶ Positive / Negative classification


- ▶ “The food was really good”

- ▶ “The chicken crossed the road because it was uncooked”


Sentiment analysis



Sentiment analysis

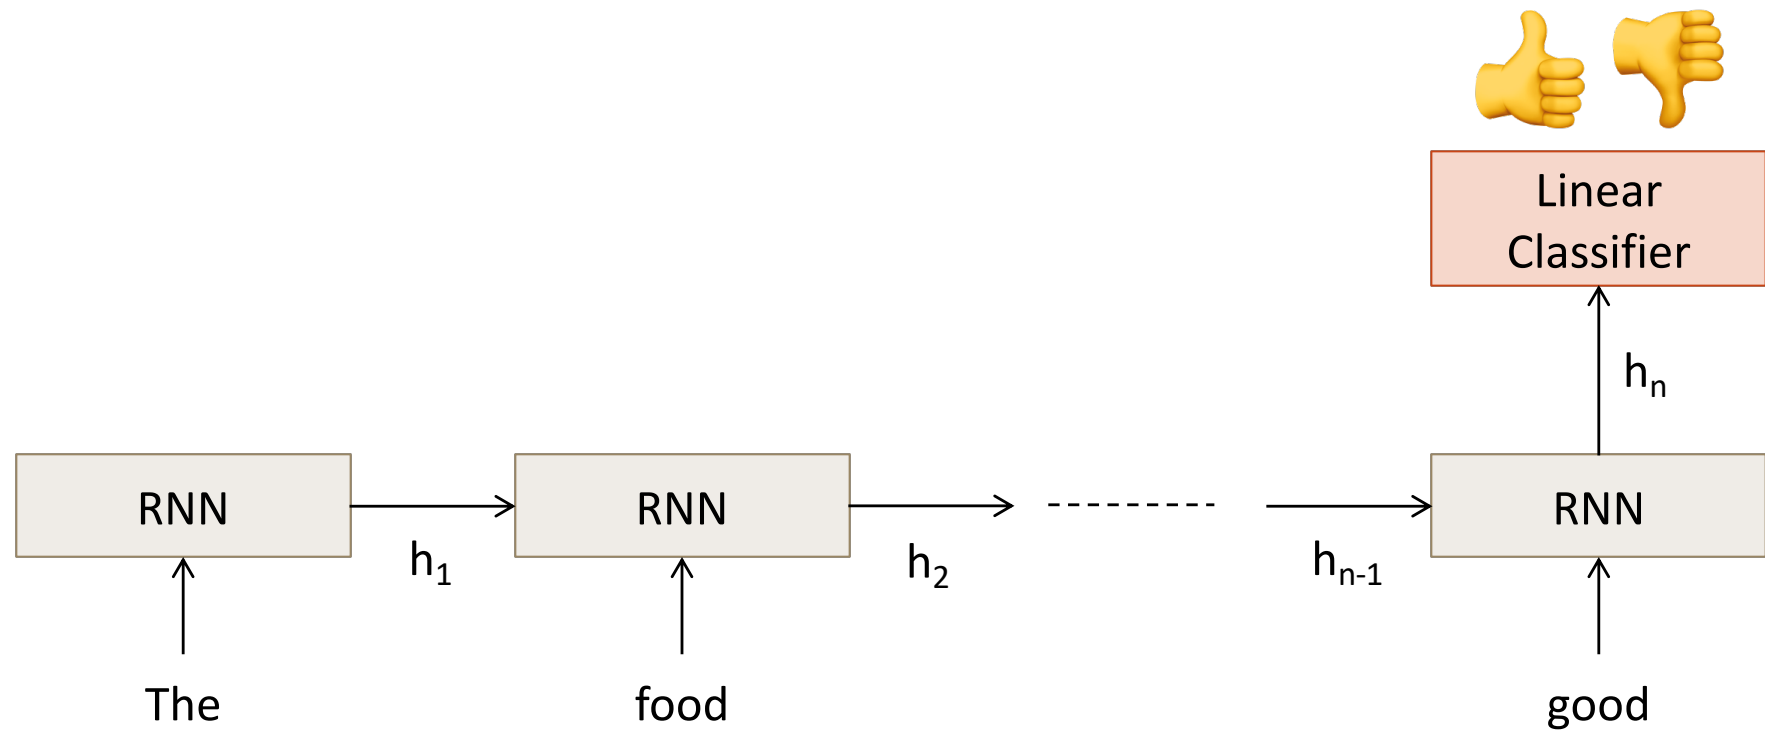


Sentiment analysis



Sentiment analysis - solution 1

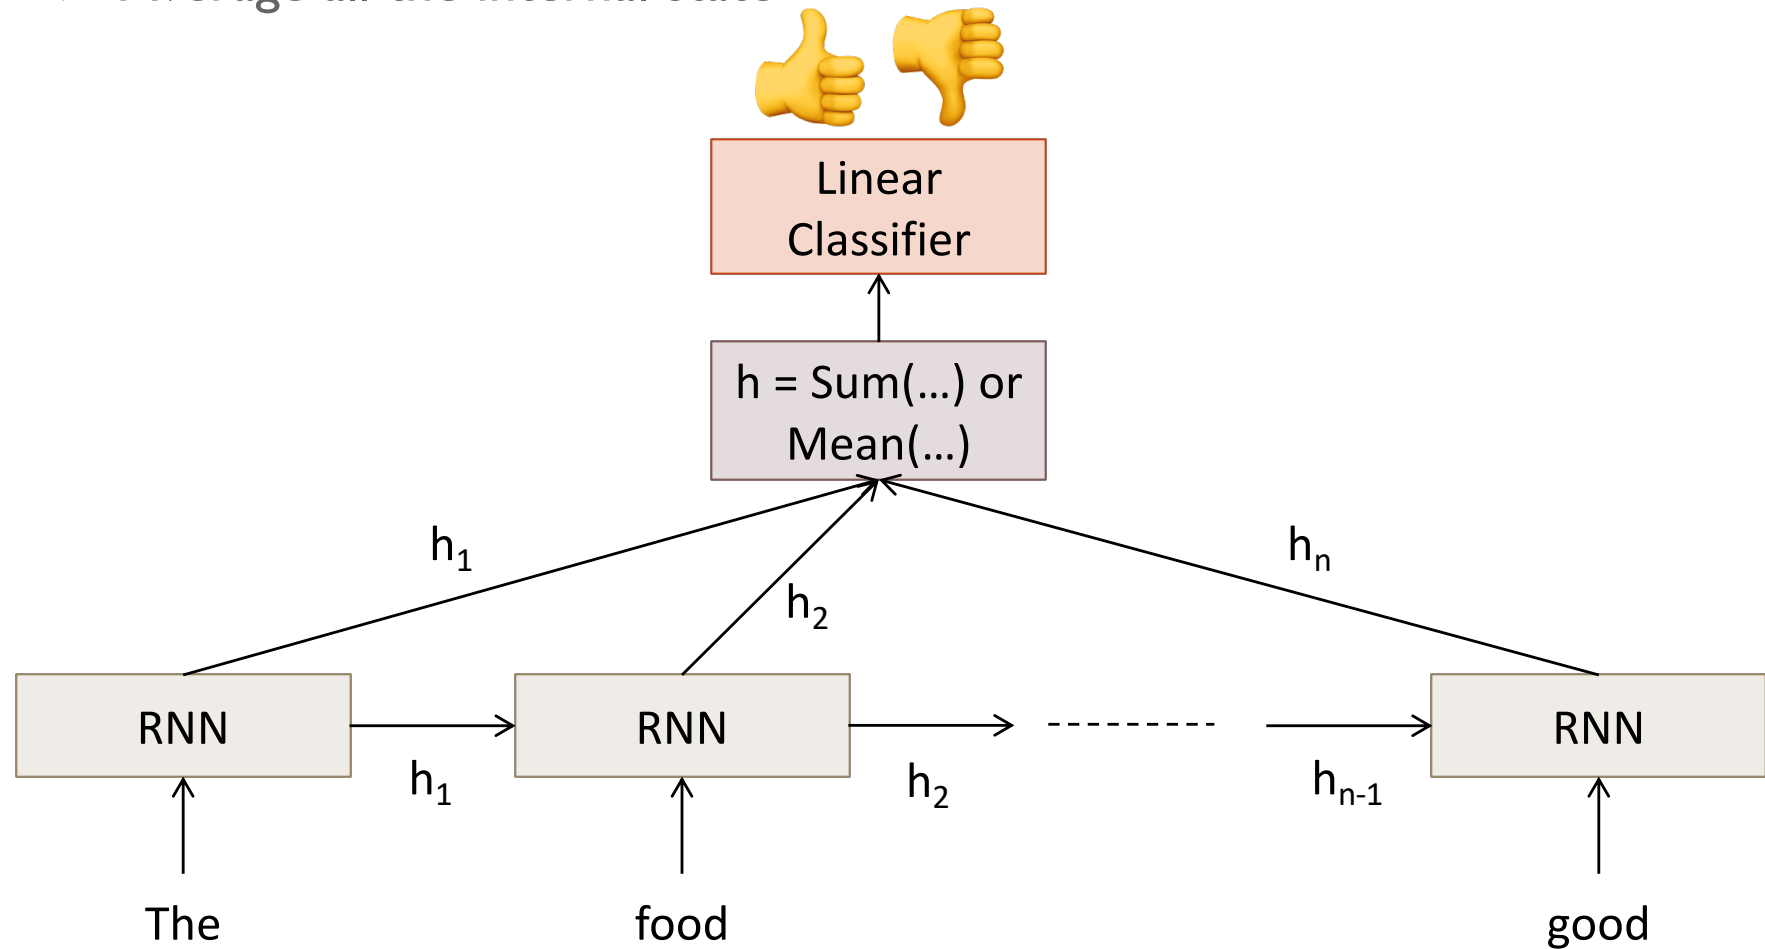
- ▶ retrieve only the last state



Sentiment analysis – solution2

- ▶ Other possible architecture

- ▶ Average all the internal state



Some use of RNN

→ Named Entity Recognition / Part of Speech Tagging

- ▶ Affect a label to each word
 - ▶ **find** and **classify** names in text
 - ▶ Could be an entity : number, country, person, ... (NER)
 - ▶ Could be a function : noun, verb, adverbs, ... (POS)

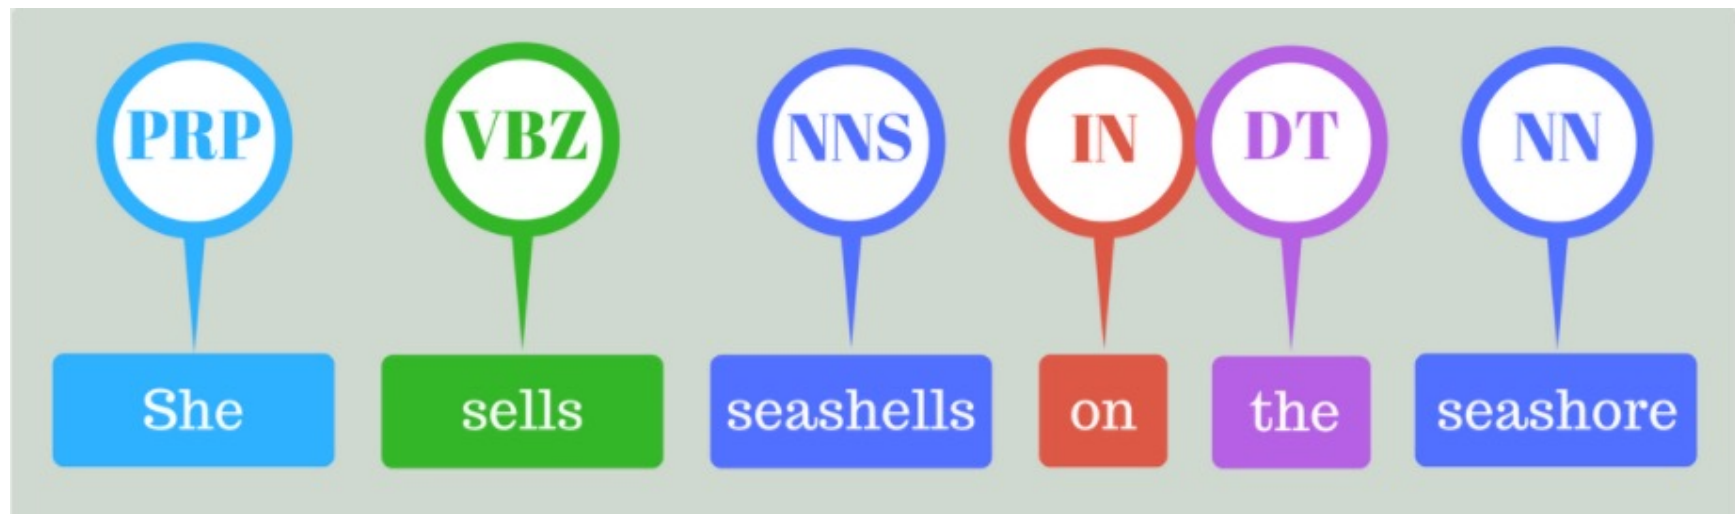
In fact, the **Chinese** **NORP** market has the **three** **CARDINAL** most influential names of the retail and tech space – **Alibaba** **GPE**, **Baidu** **ORG**, and **Tencent** **PERSON** (collectively touted as **BAT** **ORG**), and is betting big in the global **AI** **GPE** in retail industry space. The **three** **CARDINAL** giants which are claimed to have a cut-throat competition with the **U.S.** **GPE** (in terms of resources and capital) are positioning themselves to become the 'future **AI** **PERSON** platforms'. The trio is also expanding in other **Asian** **NORP** countries and investing heavily in the **U.S.** **GPE** based **AI** **GPE** startups to leverage the power of **AI** **GPE**. Backed by such powerful initiatives and presence of these conglomerates, the market in APAC AI is forecast to be the fastest-growing **one** **CARDINAL**, with an anticipated **CAGR** **PERSON** of **45%** **PERCENT** over **2018 - 2024** **DATE**.

To further elaborate on the geographical trends, **North America** **LOC** has procured **more than 50%** **PERCENT** of the global share in **2017** **DATE** and has been leading the regional landscape of **AI** **GPE** in the retail market. The **U.S.** **GPE** has a significant credit in the regional trends with **over 65%** **PERCENT** of investments (including M&As, private equity, and venture capital) in artificial intelligence technology. Additionally, the region is a huge hub for startups in tandem with the presence of tech titans, such as **Google** **ORG**, **IBM** **ORG**, and **Microsoft** **ORG**.

Some use of RNN

→ Named Entity Recognition / Part of Speech Tagging


- ▶ Affect a label to each word
 - ▶ **find** and **classify** names in text
 - ▶ Could be an entity : number, country, person, ... (NER)
 - ▶ Could be a function : noun, verb, adverbs, ... (POS)



Label representation – BIO tags


- ▶ Labels can be for words or groups of words
 - ▶ Adam Smith works for IBM , London .
- ▶ To represent this, a "BIO" representation is generally used.
 - ▶ B beginning of an entity
 - ▶ I continues the entity
 - ▶ O word outside the entity
- ▶ For example
 - ▶ ['Adam', 'Smith', 'works', 'for', 'IBM', ',', 'London', '.']
 - ▶ Without BIO: [PER, PER, O, O, ORG, O, GEO, O]
 - ▶ With BIO: [B_PER, I_PER, O, O, B_ORG, O, B_GEO, O]

MLP for NER




- Activation function for output: softmax
- Labels are OneHotEncoded

Recurrent neural network for NER

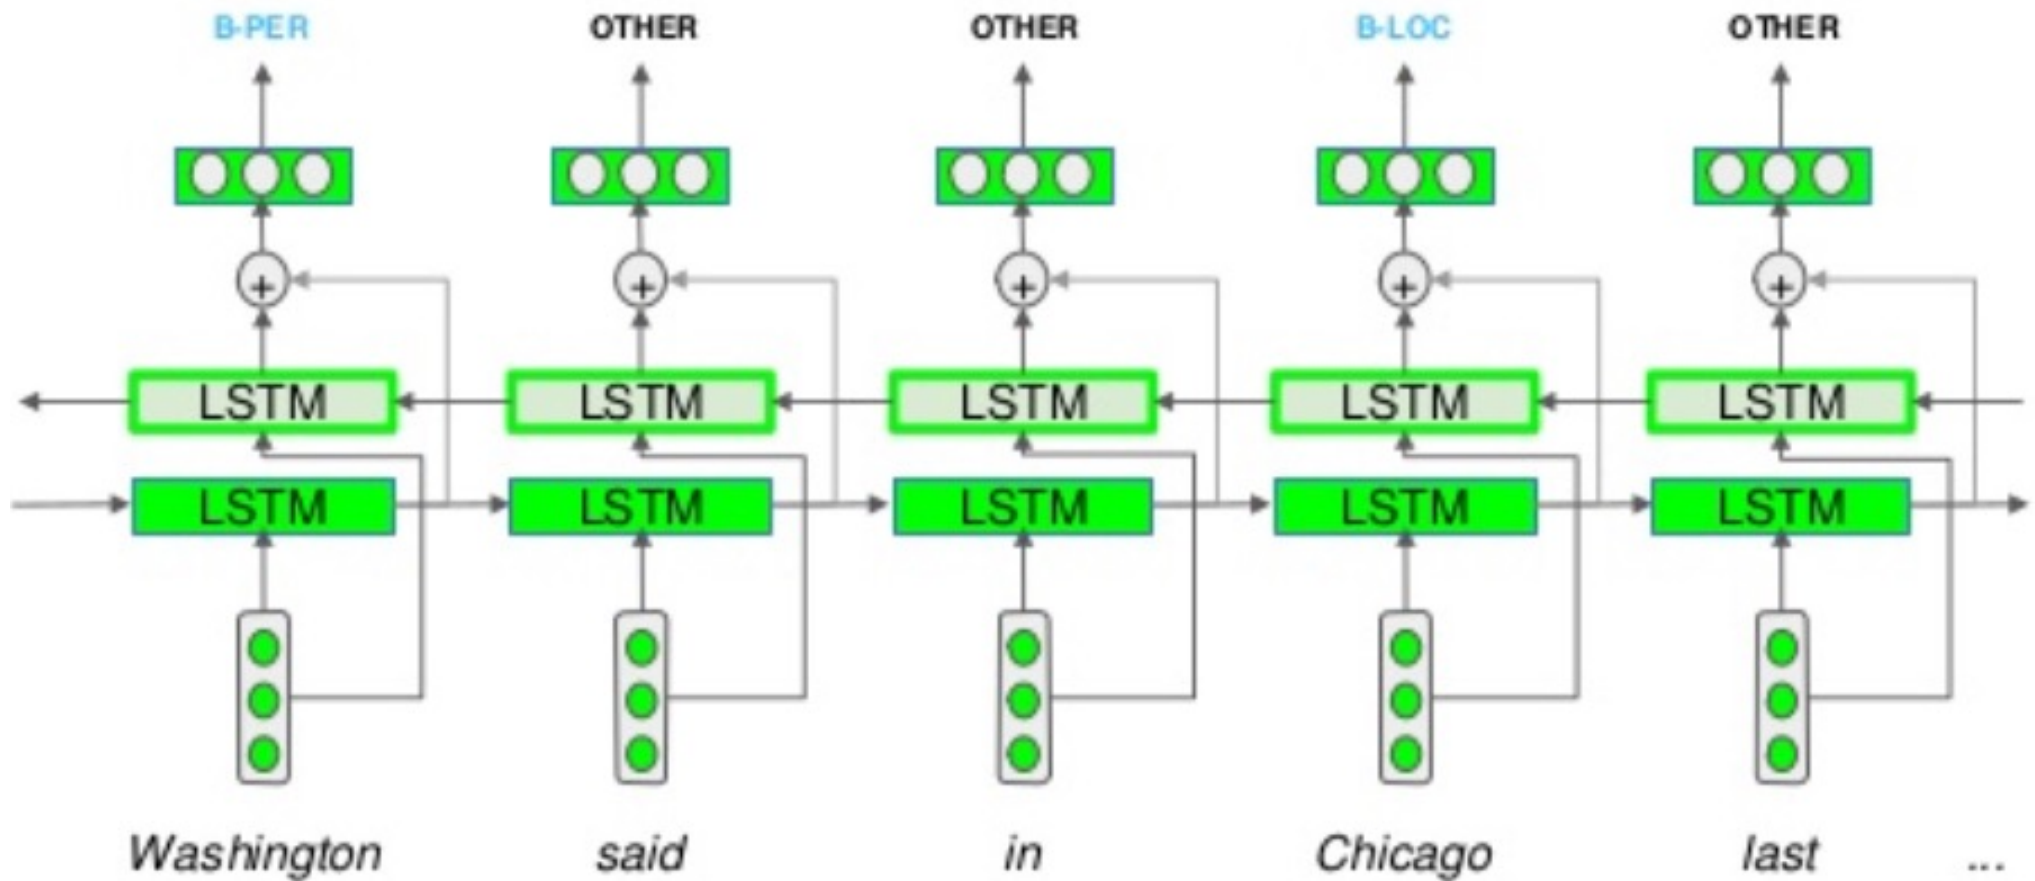


Recurrent neural network for NER (unfolded)




- Activation function for output: softmax
- Labels are OneHotEncoded

Bi directional recurrent neural network for NER



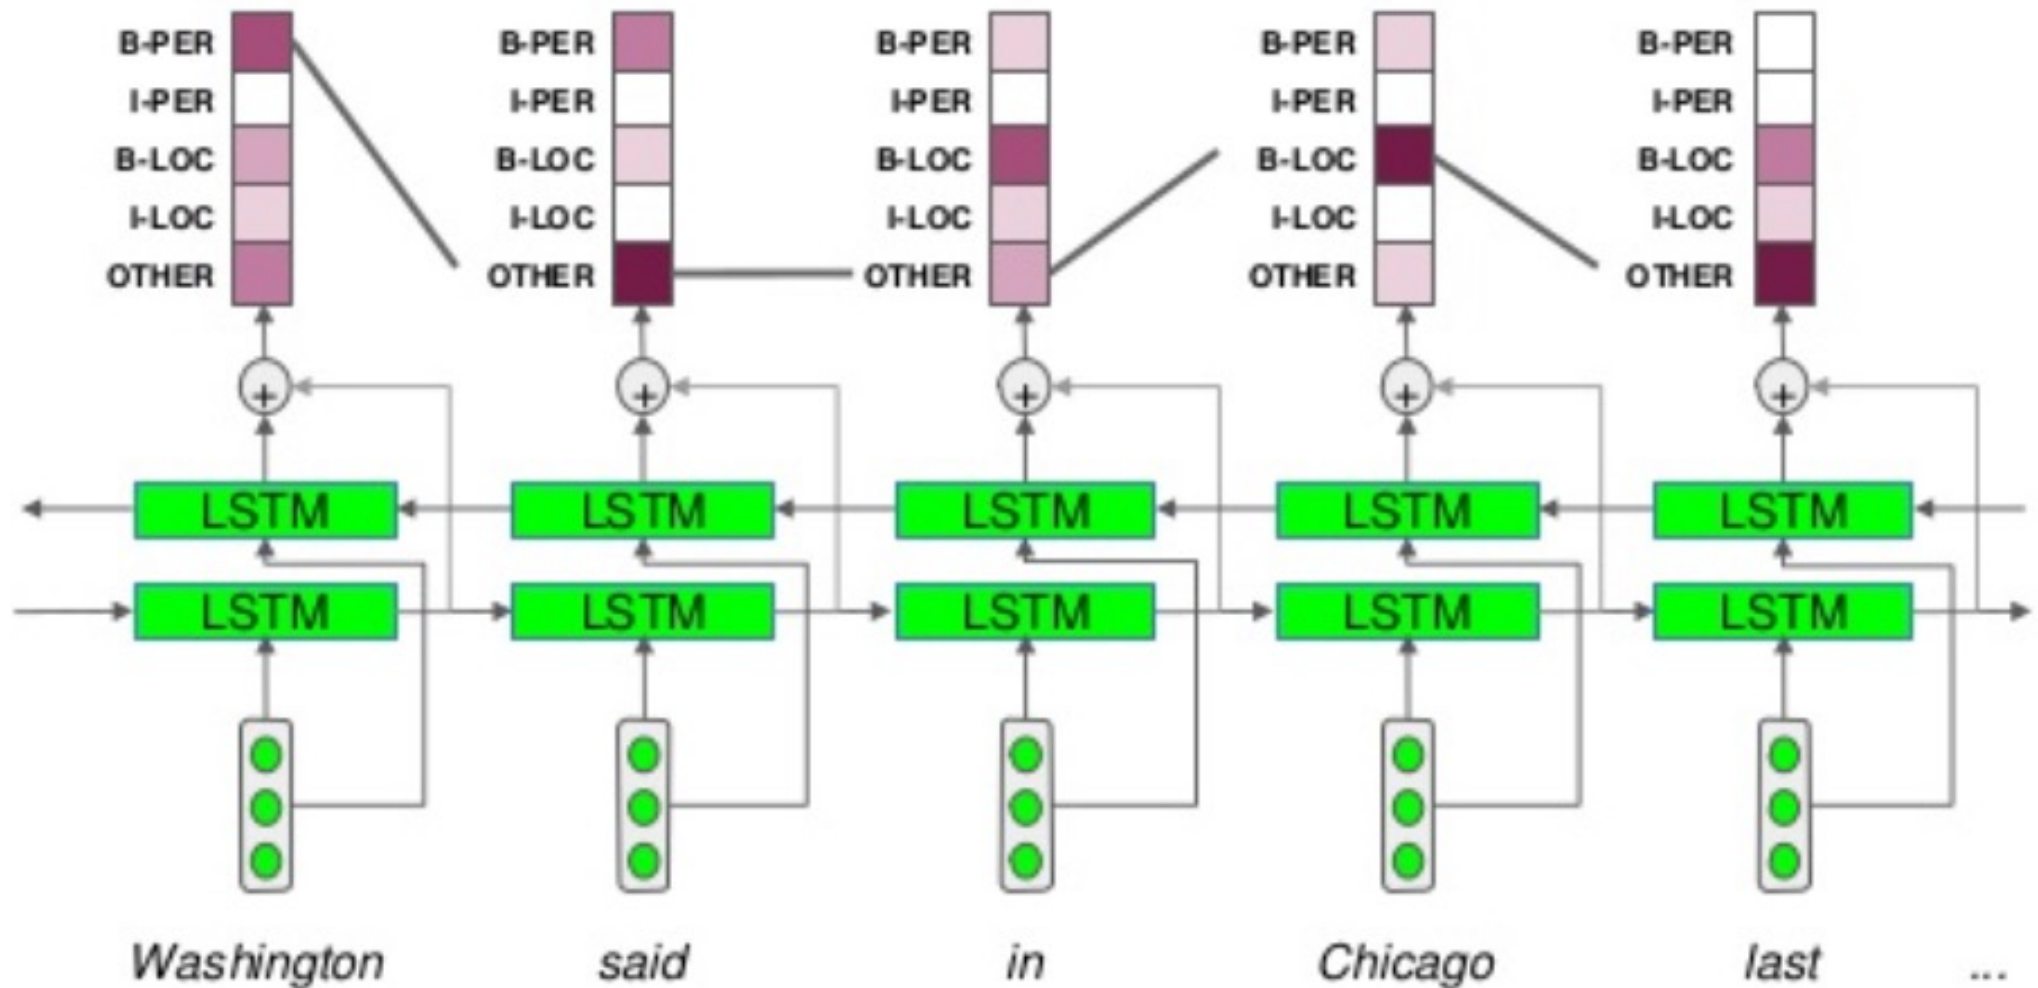
- Activation function for output: softmax
- Labels are OneHotEncoded

Stacked Bi-RNN



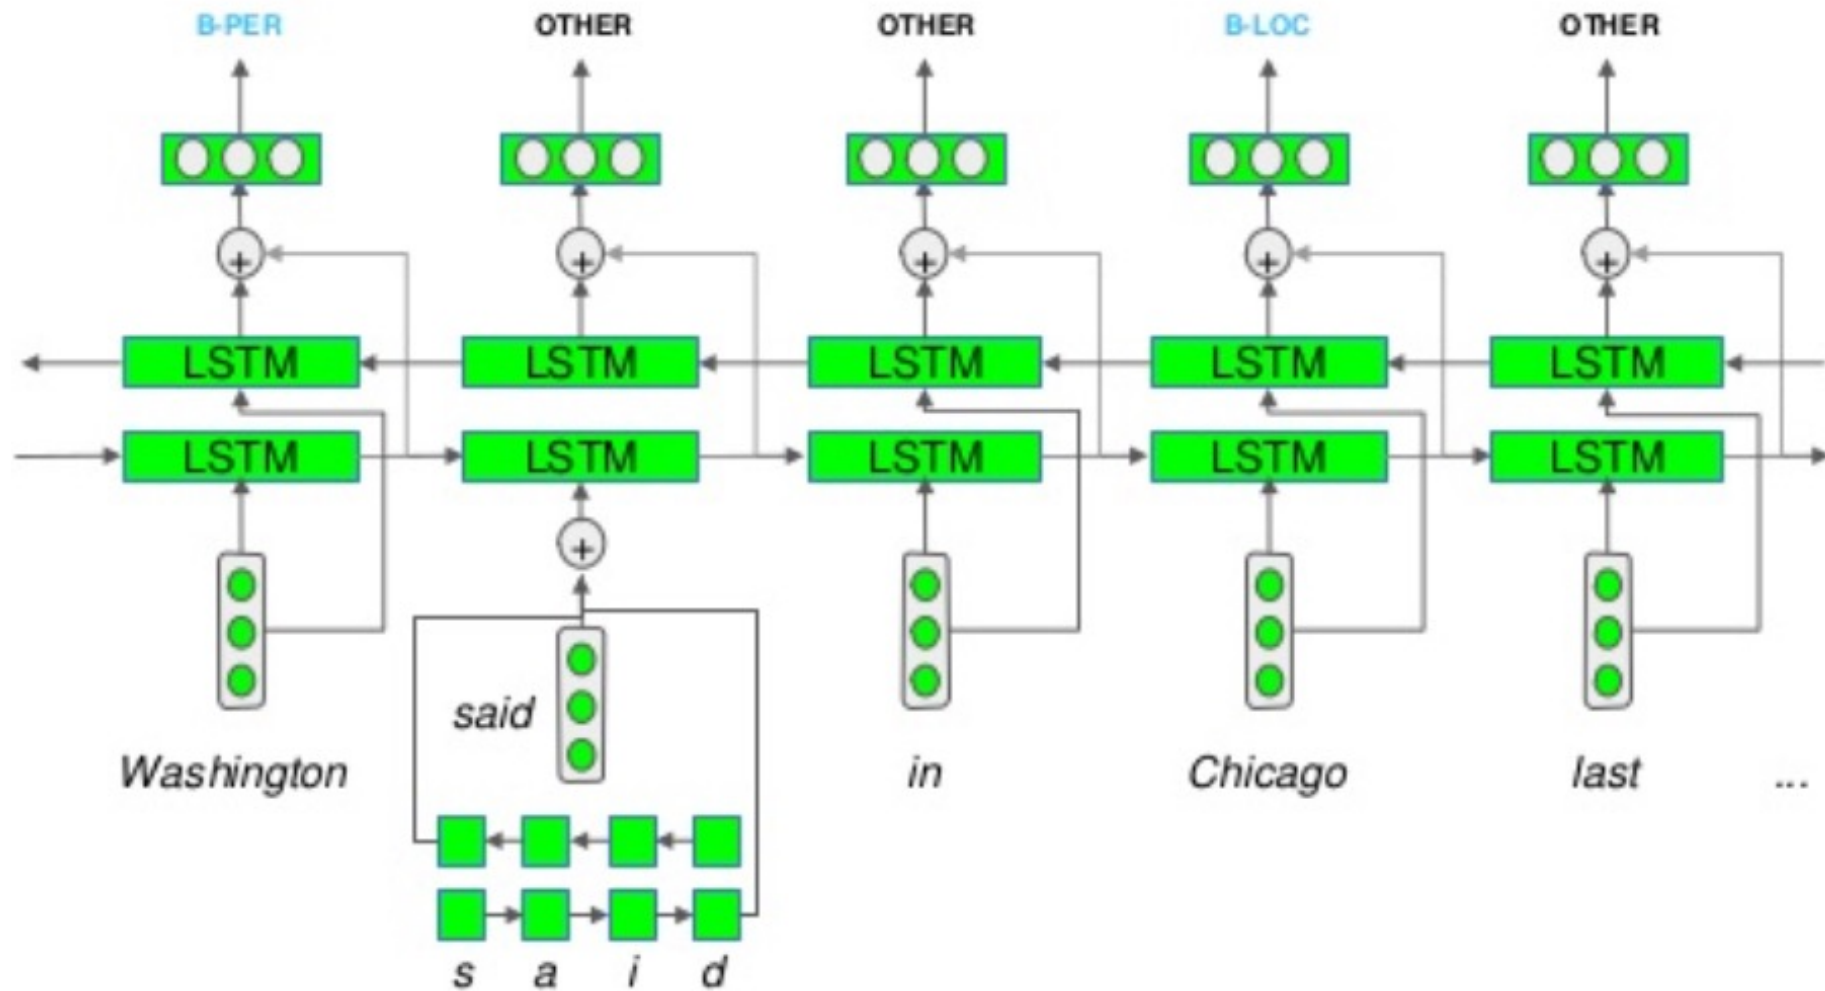
- Activation function for output: softmax
- Labels are OneHotEncoded

Bi-RNN + CRF



- CRF output activation function


Multi-level encoding char encoding + word encoding



Other use of RNN

→ Sequence2sequence model


- ▶ Used for
 - ▶ Translation
 - ▶ Chatbot



Some use of RNN

→ Input – Output Scenarios

Single - Single




One input

One output

→ Feed-forward network

Single -
Multiple




One input

Many output

Image annotation

Multiple -
Single




Many input

Many output

Text classification / Sentiment analysis

Multiple –
Multiple



Many input / Many output
 $\text{length}(\text{input}) \neq \text{length}(\text{output})$

Translation

Chat bot

Many input / Many output
 $\text{length}(\text{input}) = \text{length}(\text{output})$

Named Entity Recognition

Part-Of-Speech tagging

Other Useful Resources / References

- ▶ http://cs231n.stanford.edu/slides/winter1516_lecture10.pdf
- ▶ <http://www.cs.toronto.edu/~rgrosse/csc321/lec10.pdf>
- ▶ R. Pascanu, T. Mikolov, and Y. Bengio, [On the difficulty of training recurrent neural networks](#), ICML 2013
- ▶ S. Hochreiter, and J. Schmidhuber, [Long short-term memory](#), Neural computation, 1997 9(8), pp.1735-1780
- ▶ F.A. Gers, and J. Schmidhuber, [Recurrent nets that time and count](#), IJCNN 2000
- ▶ K. Greff, R.K. Srivastava, J. Koutník, B.R. Steunebrink, and J. Schmidhuber, [LSTM: A search space odyssey](#), IEEE transactions on neural networks and learning systems, 2016
- ▶ K. Cho, B. Van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, [Learning phrase representations using RNN encoder-decoder for statistical machine translation](#), ACL 2014
- ▶ R. Jozefowicz, W. Zaremba, and I. Sutskever, [An empirical exploration of recurrent network architectures](#), JMLR 2015