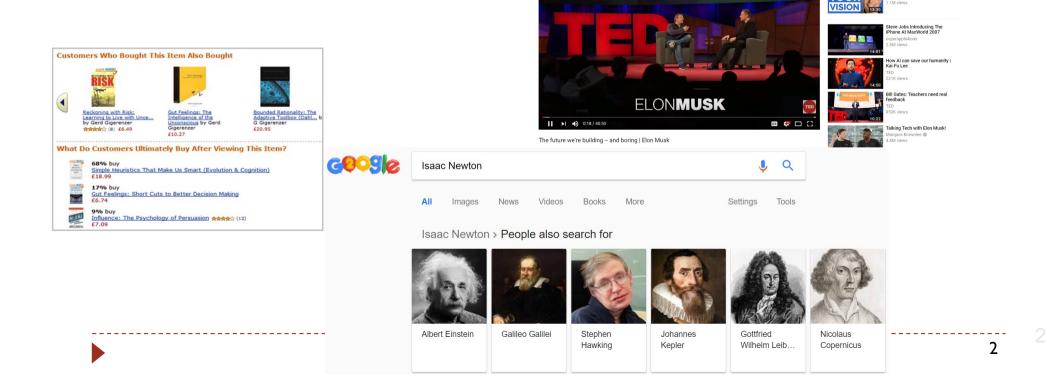
Traditional Recommender Systems

 ${\it Michel.RIVEILL@univ-cotedazur.fr}$

Recommender Systems

- A recommender system (RS) helps users that have no sufficient competence or time to evaluate the, potentially overwhelming, number of alternatives offered by a web site.
 - In their simplest form, RSs recommend to their users personalized and ranked lists of items



The Impact of RecSys

- ▶ 35% of the purchases on Amazon are the result of their recommender system, according to McKinsey.
- During the Chinese global shopping festival of November 11, 2016, Alibaba achieved growth of up to 20% of their conversion rate using personalized landing pages, according to Alizila.
- Recommendations are responsible for 70% of the time people spend watching videos on YouTube.
- ▶ 75% of what people are watching on Netflix comes from recommendations, according to McKinsey

The Age of Recommendation



Search:

User **S**



Items

Recommend:

Items



User

Amazon: A personalized online store

Frequently Bought Together



Price for both: \$158.15

Add both to Cart

Add both to Wish List

One of these items ships sooner than the other. Show details

- This item: Introduction to Data Mining by Pang-Ning Tan Hardcover \$120.16
- Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback \$37.99

Customers Who Bought This Item Also Bought



Data Science for Business: What you need...

> Foster Provost

★★★★☆ 102

#1 Best Seller (in Data Mining

Paperback \$37.99 **Prime**



Data Mining: Practical Machine Learning Tools...

) lan H. Witten ★★★☆☆ 52

Paperback \$40.65 **Prime**



Data Mining: Concepts and Techniques, Third...

Jiawei Han

全章 1757 28 Hardcover

\$60.22 **Prime**



Regression Analysis by Example > Samprit Chatterjee

*** 1 1 1 9 Hardcover

\$92.39 **Prime**



SAS Statistics by Example Ron Cody

全全全全公 10 Perfect Paperback \$44.37 **Prime**



Applied Logistic Regression David W. Hosmer Jr.

金金金金金 9 Hardcover \$62.33 **Prime**



An Introduction to Statistical Learning:...

Gareth James ★★★★★ 56

#1 Best Seller (in Mathematical & Statistical. Hardcover

\$72.79 **Prime**

Page 1 of 15

>

Recommender Problem

A good recommender

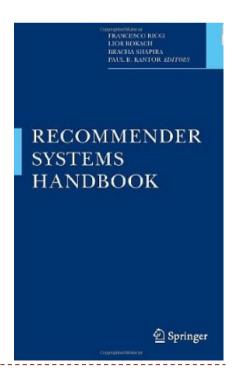
- Show programming titles to a software engineer and baby toys to a new mother.
- Don't recommend items user already knows or would find anyway.
- Expand user's taste without offending or annoying him/her...

Challenges

- Huge amounts of data, tens of millions of customers and millions of distinct catalog items.
- Results are required to be returned in real time.
- New customers have limited information.
- Old customers can have a glut of information.
- Customer data is volatile.

About this lecture

- Will give you a very short introduction to the field of Recommender Systems
- → "FIT" → How do recommender systems (RS) work?
 - ▶ How can you compute recommendations?
 - From the basis
 - Non personalized recommendation
 - To more accurate proposal
 - □ Content-based filtering
 - □ Collaborative filtering
 - ▶ It's just an introduction...
- → "EVALUATE" → How do we measure their success?
 - How can we know that the recommendations are good?



Problem characterization

Given

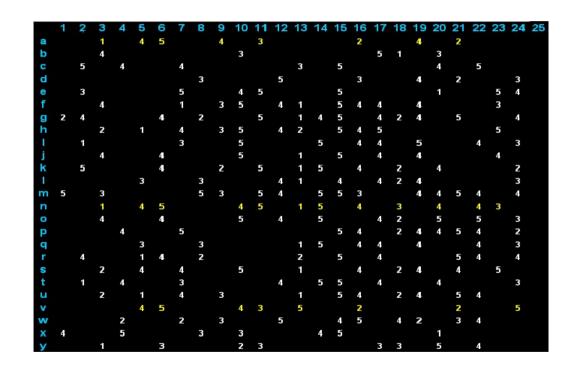
- ▶ The profile of the "active" user
- ▶ The description of a set of items
- And possibly some situational context

Compute

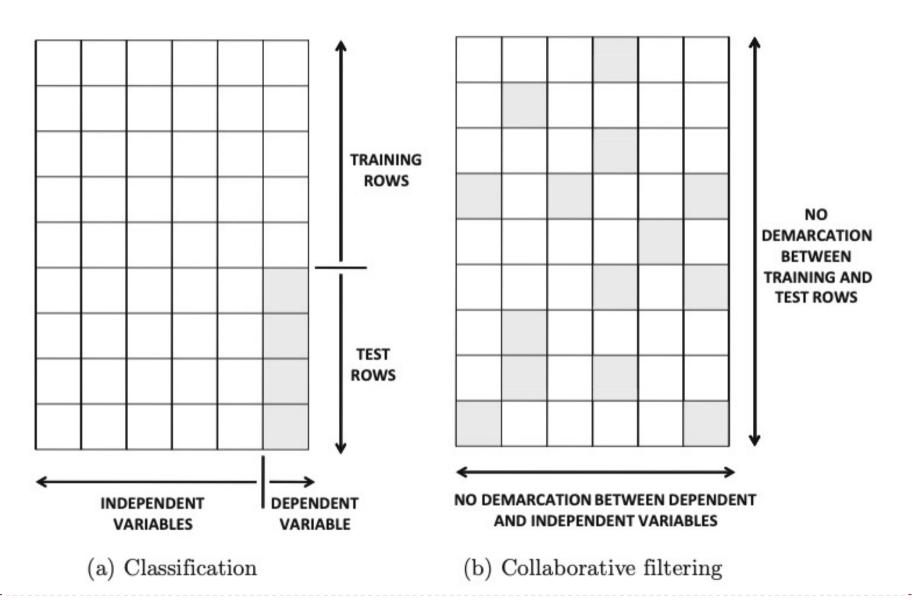
 A relevance (ranking) score for each recommendable item

▶ The problem ...

- ... is to **learn a function** that predicts the relevance score for a given (typically unseen) item
- ... it is a Matrix completion problem



ML dataset vs RS dataset



ML dataset vs RS dataset

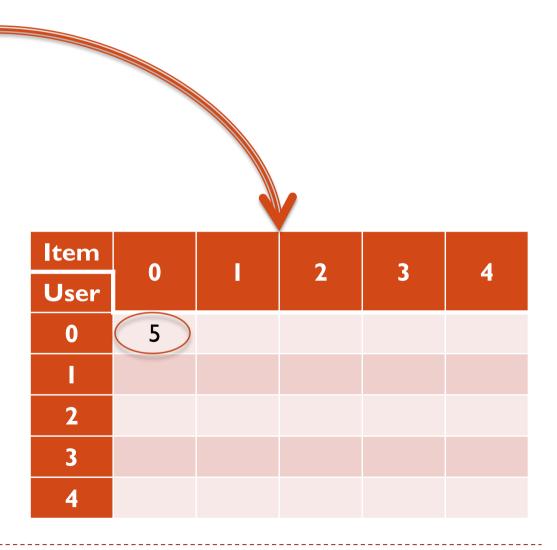
- ▶ Collaborative filtering models are closely related to missing value analysis.
- In traditional ML approach
 - Few missing value
 - Try to impute them
- ▶ In RS approach
 - Missing value >> actual rating
 - Data matrix is very large and sparse.

I. Collect the data

User	ltem	Score	Time
0	0	5	I
2	I	4	2
4	2	5	3
0	2	2	4
2	0	5	5
Etc.			

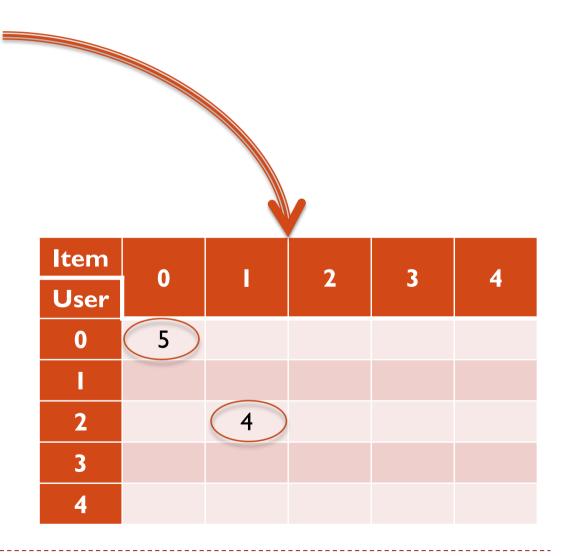
- I. Collect the data
- 2. Construct the pivot matrix

User	ltem	Score	Time
0	0	5	I
2	I	4	2
4	2	5	3
0	2	2	4
2	0	5	5
Etc.			



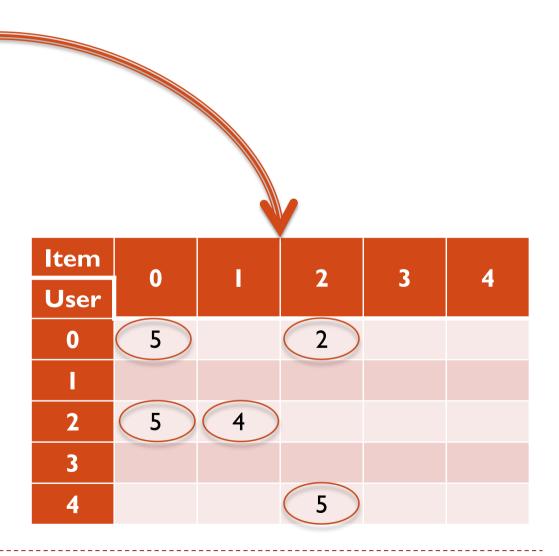
- I. Collect the data
- 2. Eventually construct the pivot matrix

User	ltem	Score	Time
0	0	5	I
2	I	4	2
4	2	5	3
0	2	2	4
2	0	5	5
Etc.			



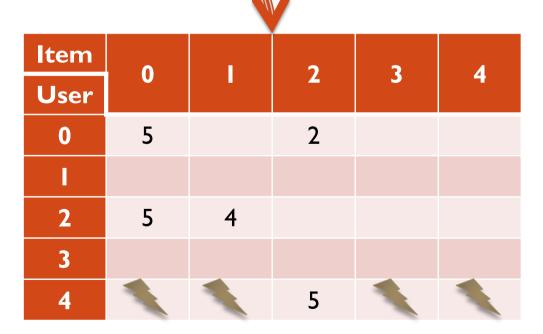
- I. Collect the data
- 2. Eventually construct the pivot matrix

User	ltem	Score	Time
0	0	5	I
2	I	4	2
4	2	5	3
0	2	2	4
2	0	5	5
Etc.			

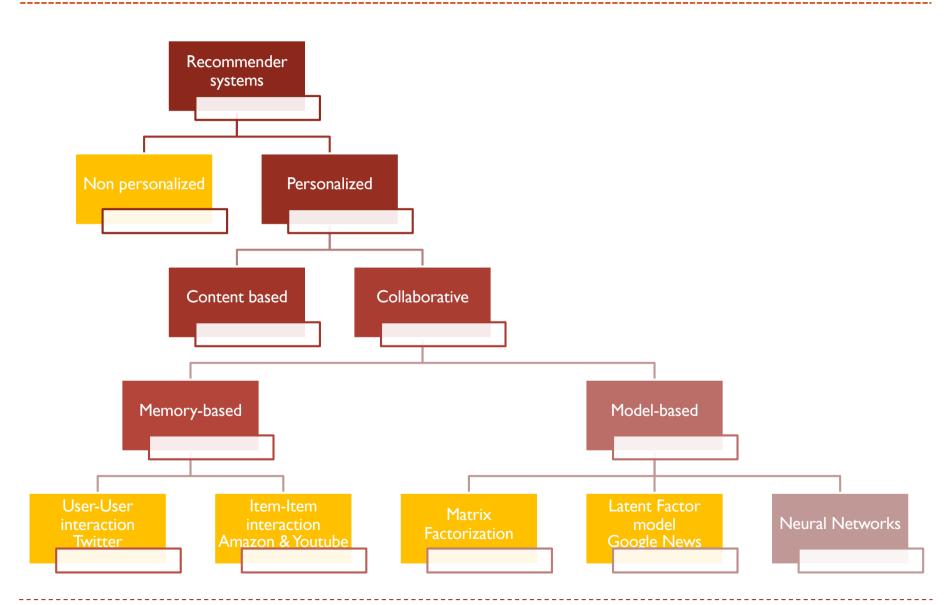


User	ltem	Score	Time
0	0	5	1
2	I	4	2
4	2	5	3
0	2	2	4
2	0	5	5
Etc.			

- I. Collect the data
- 2. Eventually construct the pivot matrix
- 3. Predict missing scores for the desired used
 - Depend of the strategy

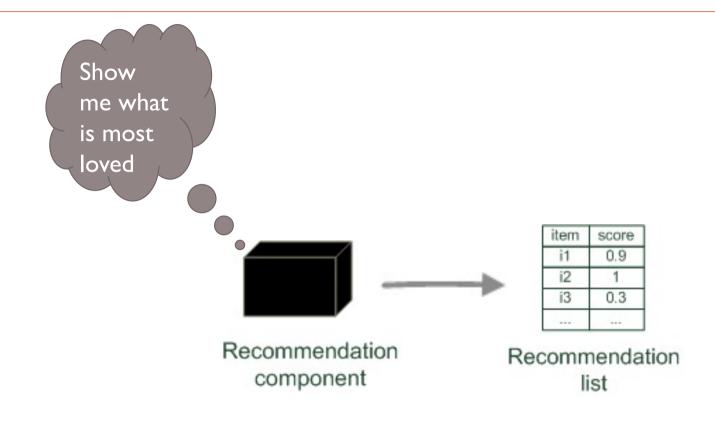


Paradigms of recommender systems



Non personalized recommendation

Non personalized recommendation



Non personalized recommendation Approach 1 = mean rating

- ▶ Goal: provides a list
 - For all user the list is the same
 - But, what do you put at the top of the list?
- A simple approach: rank by score

Score	•	$\sum_{u} r_{ui}$
30016	•	n

ltem			2	3	4	
User	U		2	3	4	
0	5	3	2			
1		5	3	3		
2	5	4	2	I		
3		3		1	2	
4			5	2		
Proposed rankink		3.75				

(3+5+4+3)/4

Non personalized recommendation Approach 1 = mean rating

- ▶ Goal: provides a list
 - For all user the list is the same
 - But, what do you put at the top of the list?
- A simple approach: rank by score
 - Score : $\frac{\sum_{u} r_{ui}}{n}$

ltem	0		2	3	A	
User	U		Z	3	4	
0	5	3	2	1,75	2	
I	5	5	3	3	2	
2	5	4	2	Ι	2	
3	5	3	3	I	2	
4	5	3.75	5	2	2	
Proposed rankink	5	3.75	3	1,75	2	

Ranking list for user 4: [0, 1, 4]

▶ Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{n}$

Item User	0	I	2	3	4	Mean rating r _u
0	5	3	2			
I		5	3	3		3.66
2	5	4	2	Ĭ		1
3		3		ĺ	2	
4			5		EP I: Calculat	
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$				for	e mean rating each user +3+3)/3	

Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{n}$

Item User	0	1	2	3	4	Mean rating r _u
0	5	3	2			3.33
1		5	3	3		3.66
2	5	4	2	I		3
3		3		I	2	2
4			5	2		3.5
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$						

▶ Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{n}$

Item User	0	1	2	3	4	Mean rating r _u
0	5	3	2			3.33
I		5	3	3		3.66
2	5	4	2	Ī		3
3		3		I	2	2
4			5	2		3.5
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$	1.835		compute the	normalized r	ating for each	n movie

((5-3.33)+(5-3))/2 =

▶ Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{r_u}$

ltem						Mean
User	0	- 1	2	3	4	rating r _u
0	5	3	2			3.33
1		5	3	3		3.66
2	5	4	2	I		3
3		3		I	2	2
4			5	2		3.5
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$	1.835	0.75	-0.37	-1.29	0	

▶ Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{r}$

Item User	0	I	2	3	4	Mean rating
Osei						r _u
0	5	3	2			3.33
1	5.5	5	3	3		3.66
2	5	4	2	ĺ		3
3	3.8	3		1	2	2
4			5	2		3.5
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$	1.835	0.75	-0.37	-1.29	0	

STEP 3: add the mean user rating to the item rating

- 3.66+1.835 for item 0/user1
- 2+1.835 for item 3/user1

Goal: provides a list

- For all user the list is the same
- But, what do you put at the top of the list?

A simple approach: rank by score

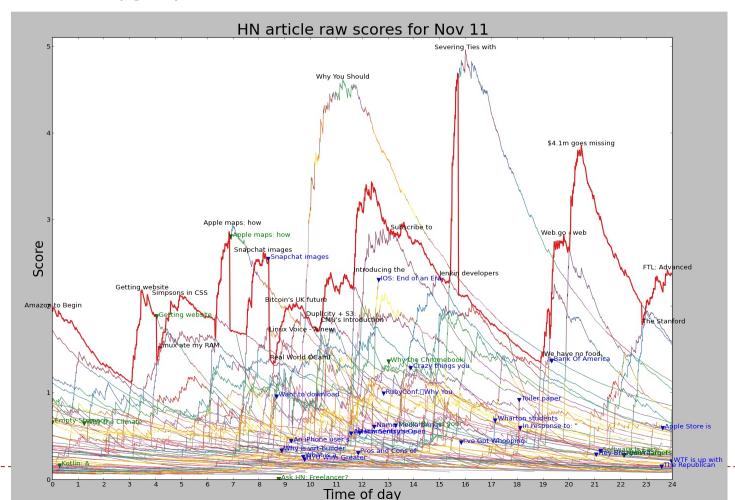
Another proposal for calculating score: normalize each rating by the mean of the rating of each user: $r_u + \frac{\sum_u (r_{ui} - r_u)}{r}$

Item User	0	I	2	3	4	Mean rating r _u
0	5	3	2	3.33-1.29	3.33-0	3.33
1	5.5	5	3	3	3.66	3.66
2	5	4	2	1	3	3
3	3.8	3	2-0.37	I	2	2
4	5.335	4.25	5	2	3.5	3.5
$\frac{\sum_{u}(r_{ui}-r_{u})}{n}$	1.835	0.75	-0.37	-1.29	0	

Ranking list for user 4: [0, 1, 4]

If the time is important....

- Hacker News algorithms
 - http://www.righto.com/2013/11/how-hacker-news-ranking-really-works.html
 - Score: $\frac{(votes-1)^{0.8}}{(age+2)^{1.8}}$



27

Personalized recommendation Collaborative filtering

Memory based

Memory-based collaborative filtering

Use the matrix of user ratings for items of the entire database to find users/or items that are similar to the active user/or items and use their rating to predict new rating

Advantage

- Quality of precision are rather good
- Relatively simple algorithm to implement for any situation
- New data can be added easily and incrementally

Disadvantage

- It depend on human ratings
- Performance decreases when data gets sparse
- Prevent scalability and have problem with large dataset

User-Based Algorithms (Breese et al, UA198)

- $r_{i,j}$ rating of user i on item j
- I_i = items for which user *i* has rated
- Mean rate for user u is

$$\qquad \qquad \mathbf{r}_{u} = \frac{\sum_{j \in I_{u}} r_{uj}}{|I_{u}|}$$



- Predicted vote for "active user" u for item j is weighted sum
 - Without normalization: $\widehat{r_{uj}} = \frac{\kappa \sum_{k \in U} w(u,k) r_{kj}}{\sum_{k \in U} w(u,k)}$
 - With normalization: $\widehat{r_{uj}} = \overline{r_u} + \frac{\kappa \sum_{k \in U} w(u,k) (\mathbf{r_{kj}} \overline{r_k})}{\sum_{k \in U} |w(u,k)|}$

weights of users

- Near 1: users u and k is similar
- Near 0 : users u and k have no similarity

User-Based Algorithms (Breese et al, UA198)

► K-nearest neighbor
$$w(a,b) = \begin{bmatrix} 1, if \ b \in neighbors(a) \\ 0, else \end{bmatrix}$$

▶ Pearson correlation coefficient $\in [-1,1]$:

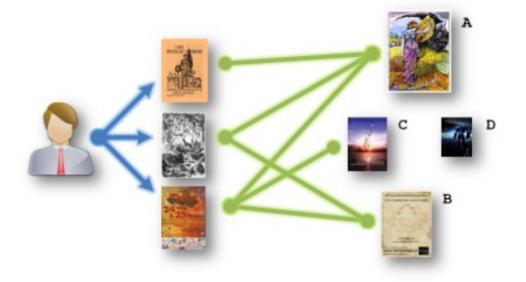
$$w(a,b) = \frac{\sum_{i} (a_{i} - \bar{a})(c_{i} - \bar{b})}{\sqrt{\sum_{i} (a_{i} - \bar{a})^{2}} \sqrt{\sum_{i} (b_{i} - \bar{b})^{2}}}$$

Cosine similarity
$$\in [-1,1]$$
:
$$w(a,b) = \frac{\sum_{i} a_{i}b_{i}}{\sqrt{\sum_{i} a_{i}^{2}} \sqrt{\sum_{i} b_{i}^{2}}}$$

Item-based Approach

- ▶ Similar with user-based approach but is on the item side
- Choosed by Amazon / Youtube







An example: Item-based approach Example: infer - r13

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	8	1	?	2	7
User 2	2	?	5	7	5
User 3	5	4	7	4	7
User 4	7	1	7	3	8
User 5	1	7	4	6	?
User 6	8	3	8	3	7

How to calculate similarity between items?

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	8	1	?	2	7
User 2	2	?:	5	7	5
User 3	5	4	7	4	7
User 4	7	1	7	3	8
User 5	1	7	4	6	?
User 6	8	3	8	3	7

Sim(13, 15)

	Item 3	Item 5
	?	7
→	5	5
→	7	7
→	7	8
	4	?
-	8	7
$ar{r}$	6,75	6,75

Calculate Pearson correlation

•
$$per(I_3, I_5) = \frac{\sum (ri_3 \bar{r}_3)(ri_5 \bar{r}_5)}{\sqrt{\sum (r_{i3} \bar{r}_3)^2} \sqrt{\sum (r_{i5} - \bar{r}_5)^2}}$$

•
$$per(I_3, I_5) = \frac{0.9375}{1,1875} = 0.79$$

Calculate cosine similarity

•
$$\cos(I_3, I_5) = \frac{I_3 \cdot I_5}{\|I_3\| \|I_5\|}$$

•
$$cos(I_3, I_5) = \frac{5*5+7*7+7*8+8*7}{\sqrt{5^2+7^2+7^2+8^2}\sqrt{5^2+7^2+8^2+7^2}}$$

•
$$cos(I_3, I_5) = 0.99$$

How to calculate r_{13} ?

 $ightharpoonup r_{13}$ = ranking for user I and item 3 without normalization

$$r_{13} = \frac{r_{11} * sim(I_1,I_3) + r_{12} * sim(I_2,I_3) + r_{14} * sim(I_4,I_3) + r_{15} * sim(I_5,I_3)}{|sim(I_1,I_3) + sim(I_2,I_3) + sim(I_4,I_3) + sim(I_5,I_3)|}$$

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	8	1	?	2	7
User 2	2	?	5	7	5

$$r_{13} = \frac{8*sim(I_1,I_3) + 1*sim(I_2,I_3) + 2*sim(I_4,I_3) + 7*sim(I_5,I_3)}{|sim(I_1,I_3) + sim(I_2,I_3) + sim(I_4,I_3) + sim(I_5,I_3)|}$$

Generalization (without normalization)

- ▶ User based: $r_{ui} = \frac{\sum_{v} r_{vi} * sim(u,v)}{\sum_{v} |sim(u,v)|}$, $v \in similar User$
 - Calculate sim, only with common item
 - Rating matrix is almost stable \rightarrow pre-calculate sim matrix
 - Trick I:
 - normalize the rating
 - Trick II:
 - ▶ If two users have very few items in common,
 - □ For example : I rating and it is the same--> hight similarity but low confidence
 - Weigh similarity by $\frac{\min(N,|P_{uv}|}{|P_{uv}|}$
 - $\square |P_{uv}|$ is the number of common items shared by user u and v
- ▶ Item based: $r_{ui} = \frac{\sum_{j} r_{uj} * sim(i,j)}{\sum_{j} |sim(i,j)|}$, $j \in similar\ Item$
 - Adapt the previous recommendation

Offline computation: Online Recommendation

Offline Computation:

- **builds a similar-items table** which is extremely time intensive, O(N²M)
 - ▶ M is the number of customers ± 10 millions
 - N is the number of items ± 1 millions
- ▶ In practice, it's closer to O(NM), as most customers have very few purchases
- Sampling customers can also reduce runtime even further with little reduction in quality.

Online Recommendation:

- Given a similar-items table, the algorithm
 - finds items similar to each of the user's purchases and ratings,
 - aggregates those items, and then
 - recommends the most popular or correlated items.
- For User based, exchange Item and User
 - Item similarity is more stable than User similarity
 - N >> M → Item based is more efficient → choosed by Amazon / Youtube

Scalability and Quality: Comparison

User based collaborative filtering:

- little or no offline computation
- impractical on large data sets, unless it uses dimensionality reduction, sampling, or partitioning
- dimensionality reduction, sampling, or partitioning reduces recommendation quality

Item based collaborative filtering:

- scalability and performance are achieved by creating the expensive similar-items table offline
- online component "looking up similar items" scales independently of the catalog size or the number of customers
 - → fast for extremely large data sets
- recommendation quality is excellent since it recommends highly correlated similar items
- unlike traditional collaborative filtering,
 - the algorithm performs well with limited user data,
 - producing high-quality recommendations based on as few as two or three items

User-based or Item-based Approach

PRO

- Few modeling assumptions
- Few tuning parameters to learn
- Easy to explain to users
 - ▶ Dear Amazon.com Customer, We've noticed that customers who have purchased or rated How Does the Show Go On: An Introduction to the <u>Theater</u> by Thomas Schumacher have also purchased Princess Protection Program #I:A Royal Makeover (Disney Early Readers).

CONS

- ▶ Computationally expensive, $O(M^2N)$ for User based in the worst case, where
 - ▶ M is the number of customers ± 10 millions
 - N is the number of items ± 1 millions
- Dimensionality reduction can increase the performance,
 - ▶ BUT, also reduce the quality of the recommendation
- For very large data sets
 - ▶ the algorithm encounters severe performance and scaling issues

Work on scalability: Cluster Models

Approach

- Divide the customer base into many segments and treat the task as a classification problem
- Assign the user to the segment containing the most similar customers
- Uses the purchases and ratings of the customers in the segment to generate recommendations
- Cluster models have better online scalability and performance than collaborative filtering because they compare the user to a controlled number of segments rather than the entire customer base.

Problems

- Quality of the recommendation is low
- The recommendations are less relevant because the similar customers that the cluster models find are not the most similar customers
- To improve quality, it needs online segmentation, which is almost as expensive as finding similar customers using collaborative filtering

Comparative results:

- The MovieLens dataset contains
 - 1 million ratings
 - 6,040 users
 - 3,900 movies.
- The best overall results are reached by the item-by-item based approach. It needs 170 seconds to construct the model and 3 seconds to predict 100,021 ratings.

	User Based	Model Based	Item Based
model construction time (sec.)	730	254	170
prediction time (sec.)	31	1	3
MAE	0.6688	0.6736	0.6382

Table from: Candillier, L., Meyer, F., & Boull'e Marc. (2007). Comparing state-of-the-art collaborative filtering systems [3]

Personalized recommendation Collaborative filtering

Model based

Model-based collaborative filtering

- Find patterns based on training data, and these are used to make predictions for real data
- Extract some information from dataset, and use that as a "model" to make recommendations without having to use complete dataset every time

Advantage

- Handle sparsity better than memory based ones
- Scalable with large dataset
- Improve prediction speed

Disadvantage

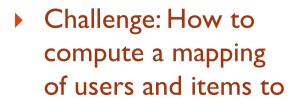
- Expensive model building
- Can lose useful information due to reduction models

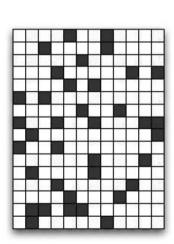
Latent based recommendation

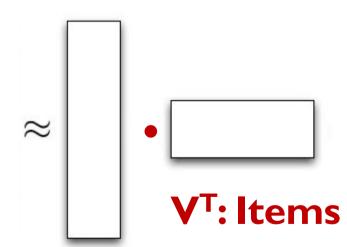
- Models with latent classes of items and users
 - Individual items and users are assigned to either a single class or a mixture of classes
- Could be implemented by
 - Latent Factor Models
 - Items and users described by unobserved factors
 - Matrix factorization or SVD decomposition
 - Probabilistic latent semantic analysis (PLSA)
 - Neural networks

Matrix factorization

- Mathematic concept of matrix factorization
 - Set of Users (U)
 - Set of Items (V)
 - ▶ R = rating matrix
 - ▶ The goal is to discover K latent features, with the input of the 2 matrix U,V
 - ▶ U(|U|*k)
 - V(|V|*k)
 - $R \approx U.V^T = \hat{R}$
 - $\widehat{r_{ui}} = u_u^T v_i = \sum_k u_{uk} v_{ki}$

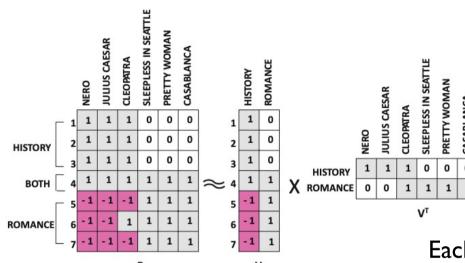




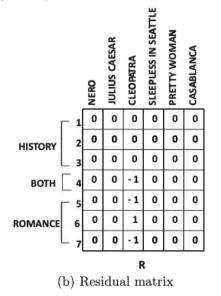


Rating matrix U: Users

Matrix factorization



(a) Example of rank-2 matrix factorization



Each row of U or V is latent factor

The ith row of U is user factor

- → contain k entries corresponding to the affinity of user i towards the k concepts in the ratings matrix
- \rightarrow Here k=2

Each rating r_{ui} in R can be approximately Expressed by $r_{ui} \approx \overline{u_u}$. $\overline{v_i}$

Matrix factorization Leaning algorithms

- Cost function error : MSE
 - Cost = $\sum (error)^2 = \sum (e_{ui})^2 = \sum (actual rating predicted rating)2$ = $\sum_{ui} (r_{ui} - \widehat{r_{uj}})^2 = \sum_{ui} (r_{ui} - u_u^T v_i)^2$
- ▶ Goal: minimize the cost function i.e. minimize the MSE
 - It's an optimization problem
- ▶ Generic solution : gradient descent approach
 - Gradient: $\frac{\partial}{\partial u_{uk}} e_{uj}^2 = -2(r_{uj} \widehat{r_{uj}}) \ v_{jk} = -2e_{uj}v_{jk}$ $\frac{\partial}{\partial v_{ik}} e_{ui}^2 = -2(r_{ui} \widehat{r_{ui}}) \ u_{uk} = -2e_{uj}u_{uk}$
 - Modification of parameters in the opposite direction of the gradient, with α (learning rate)
 - $\qquad \qquad u'_{uk} \leftarrow u_{uk} + \alpha \ e_{uj} v_{jk}$
 - $v'_{jk} \leftarrow v_{jk} + \alpha \ e_{uj} v_{ik}$
 - $\forall k \in \{1..K\}$

Matrix decomposition through SVD decomposition

- Singular value decomposition (SVD)
 - is a form of matrix factorization
 - For a m*n matrix A, we want to compute 3 matrix U, S,V such as $A = USV^T$
 - ▶ U is an m*m square matrix
 - ▶ S is an $\mathbf{m}^*\mathbf{n}$ diagonal matrix, $S_{i+1,i+1} \le S_{i,i}$
 - V is an n*n matrix
 - in which the columns of U and V are constrained to be mutually orthogonal
 - Mutual orthogonality = the concepts can be completely independent
 - \blacktriangleright U and V respectively contain the eigenvecyors of AA^T and A^TA
 - ▶ S contains the square roots of the eigenvalues
- ▶ For fully specified matrix, it easy to perform SVD decomposition
 - A lot of methods
 - A lot of tools

Truncated SVD decomposition

- In order to approximate the rating matrix,
 - Use truncated SVD decomposition at rank $k \ll \min\{m, n\}$
 - $R \approx Q_k \Sigma_k P_k^T$
- The factorization here is into three matrices rather than two
 - ightharpoonup However, the diagonal matrix Σ, can be absorbed in either the user factors Qk or the item factors Pk
 - By convention
 - $U \approx Q_k \Sigma_k$
 - $V \approx P_k^T$

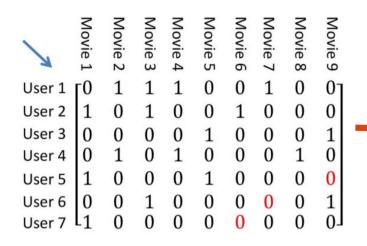
SVD decomposition uses

- ▶ Each row u_i of U is a vector representation of user i
- Each row v_i of V is a vector representation of item j
- ▶ Sd,d assigns a weight to each dimension d of these vectors
- Score $A_{i,j}$ is the weighted do product of u_i and v_j , where the weight of each dimension is value $S_{d,d}$
 - $A_{i,j} = \sum_{d} U_{i,d} * Sd_{jd} * Vd_{jj}$

Example:

1	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6	Movie 7	Movie 8	Movie 9
User 1	Γ0	1	1	1	0	0	1	0	07
User 2	1	0	1	0	0	1	0	0	0
User 3	0	0	0	0	1	0	0	0	1
User 4	0	1	0	1	0	0	0	1	0
User 5	1	0	0	0	1	0	0	0	0
User 6	0	0	1	0	0	0	0	0	1
User 7	L_1	0	0	0	0	0	0	0	07

SVD decomposition uses



USV^{T} , With M = 4

Colors:

Red: best value for users 5,6,7

Green: movies already seen

Black: unseen movies

					*				
ſ	0.00	1.00	1.01	1.00	-0.17	0.17	0.63	0.37	0.177
I	1.12	0.00	1.02	0.00	-0.02	0.69	0.17	-0.17	-0.03
I	0.06	-0.02	0.08	-0.02	0.97	-0.18	-0.05	0.03	0.99
I	0.00	1.00	-0.01	1.00	0.18	-0.17	0.37	0.63	-0.17
I	1.02	0.02	-0.09	0.02	0.86	0.16	-0.12	0.14	0.16
I	-0.11	0.01	0.94	0.01	0.18	0.16	0.22	-0.21	0.88
l	0.78	-0.02	0.08	-0.02	0.24	0.26	-0.05	0.03	-0.18

Recommendation

- Movies 6 and 9 for user 5
- Movies 7 for user 6
- Movie 6 for user 7

SVD decomposition uses

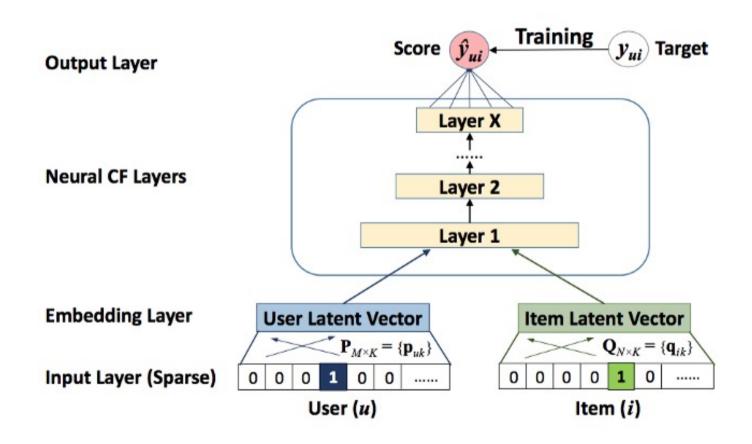
- ▶ Each row ui of U is a vector representation of user i
- Each row v_i of V is a vector representation of item j
- ▶ Sd,d assigns a weight to each dimension d of these vectors
- Score $A_{i,j}$ is the weighted do product of u_i and v_j , where the weight of each dimension is value $S_{d,d}$
 - $A_{i,j} = \sum_{d} U_{i,d} * Sd_{d} * Vd_{j}$

Example:

1	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	Movie 6	Movie 7	Movie 8	Movie 9
User 1	Γ0	1	1	1	0	0	1	0	ر0
User 2	1	0	1	0	0	1	0	0	0
User 3	0	0	0	0	1	0	0	0	1
User 4	0	1	0	1	0	0	0	1	0
User 5	1	0	0	0	1	0	0	0	0
User 6	0	0	1	0	0	0	0	0	1
User 7	L1	0	0	0	0	0	0	0	01

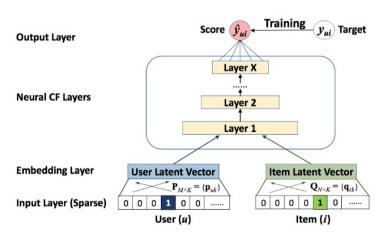
How to use neural network?

- Neural collaborative filtering (NCF)
 - Learn interaction with a deep neural network



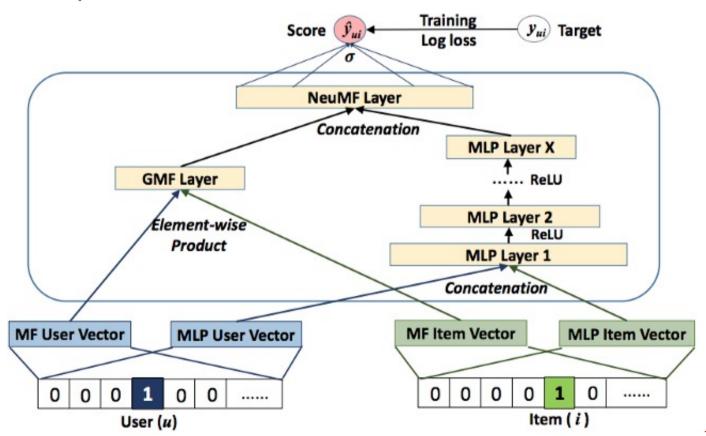
Neural collaborative filtering (NCF)

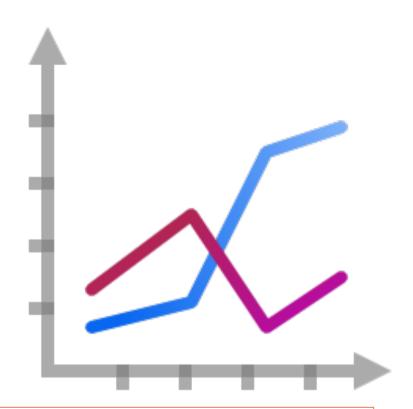
- Input Layer binarise a sparse vector for a user and item identification where:
 - Item (i): I means the user u has interacted with Item(i)
 - User (u): user id
- Embedding layer is a fully connected layer that projects the sparse representation to a dense vector.
 - The obtained user/item embeddings are the latent user/item vectors.
- Neural CF layers use multi-layered neural architecture to map the latent vectors to prediction scores.
- The final output layer returns the predicted score by minimizing the pointwise loss/pairwise loss.



Neural collaborative filtering (NCF)

- According to the construction of layer I, we can build different models
 - concatenation -> Multi-layer Perceptron (MLP)
 - dot-product -> Generalized Matrix Factorisation (GMF)
- It is also possible to combine the models





Evaluating recommender systems

How to evaluate RS?









Test with real users

- A/B tests
- Example measures: sales increase, click through rates



- Controlled experiments
- Example measures: satisfaction with the system (questionnaires)

Offline experiments

- Based on historical data
- ▶ Example measures: prediction accuracy, coverage



Traditional accuracy metrics

- Prediction accuracy metrics include the mean absolute error (MAE), root mean square error (RMSE).
 - Use in regression task
 - Focus on comparing the actual vs predicted ratings
 - Easy to calculate and easy to use
 - They are not relevant for recommander systems because we are more interest in assessing the relevance of the best recommendations than the totality of the recommendations

Record

- MAE: Mean absolute error
 - sklearn.metrics.mean_absolute_error
- MSE: Mean square error
 - sklearn.metrics.mean_squared_error
- RMSE: Root MSE
 - √MSE

MAE =
$$\frac{1}{n} \sum_{j=1}^{n} |y_j - \hat{y}_j|$$

RMSE = $\sqrt{\frac{1}{n} \sum_{j=1}^{n} (y_j - \hat{y}_j)^2}$

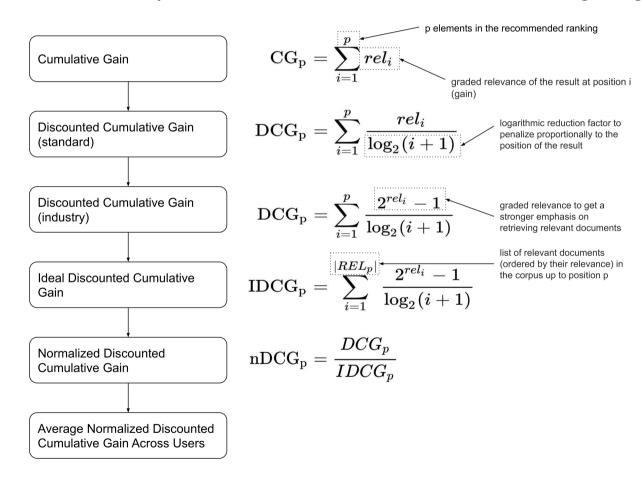
Ranking metrics

- In order to have more suitable metrics, in recommender systems we use mainly
 - MRR: Mean Reciprocal Rank
 - MAP: Mean Average Precision
 - ▶ NDCG: Normalized Discounted Cumulative Gain --> implemented in sklearn
- Others specific metrics available in sklearn
 - ▶ See: multilabel ranking metrics on sklearn documentation.

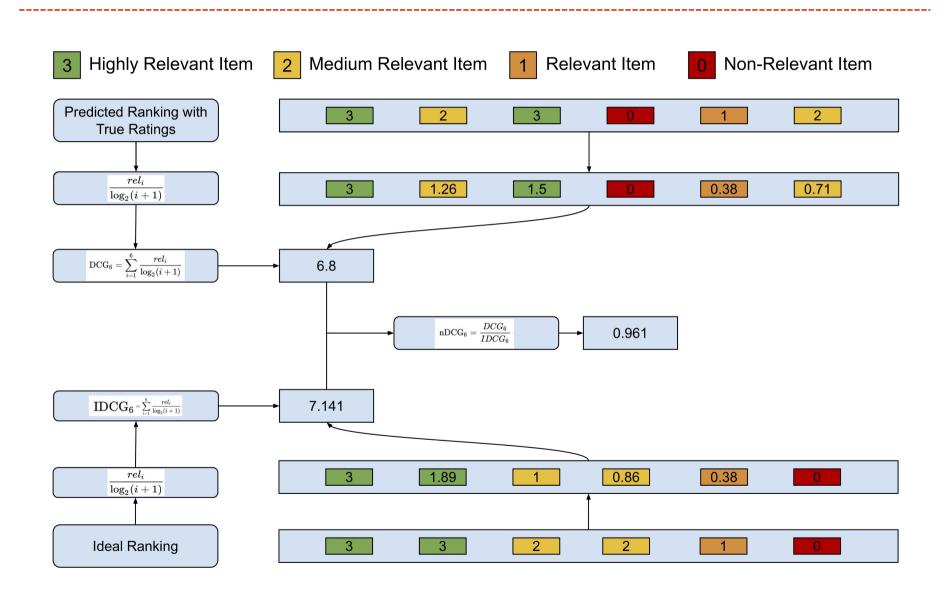
NDCG: Normalized Discounted Cumulative Gain

Goal of NDCG

Place more emphasis on relevant recommendations at the beginning of lists



NDCG in practical



NDCG Pros and Cons

NDCG Pros

The primary advantage of the NDCG is that it takes into account the graded relevance values. When they are available in the dataset, the NDCG is a good fit.

NDCG Cons

- ▶ The NDCG has some issues with partial feedback.
- This happens when we have incomplete ratings.
- ▶ This is case for the majority of recommender systems situations.
- If we had complete ratings there would be no real task to achieve!
- Generally the missing values have a default value of 0