

Table of Contents

- [1 Introduction : a visual intuition of activation functions](#)
- [2 Build an MLP to classify MNIST images](#)

Introduction to the Multi Layer Perceptron

Introduction : a visual intuition of activation functions

Using a regression task on the sinus function, we'll try to get an intuition of the effect of activation functions.

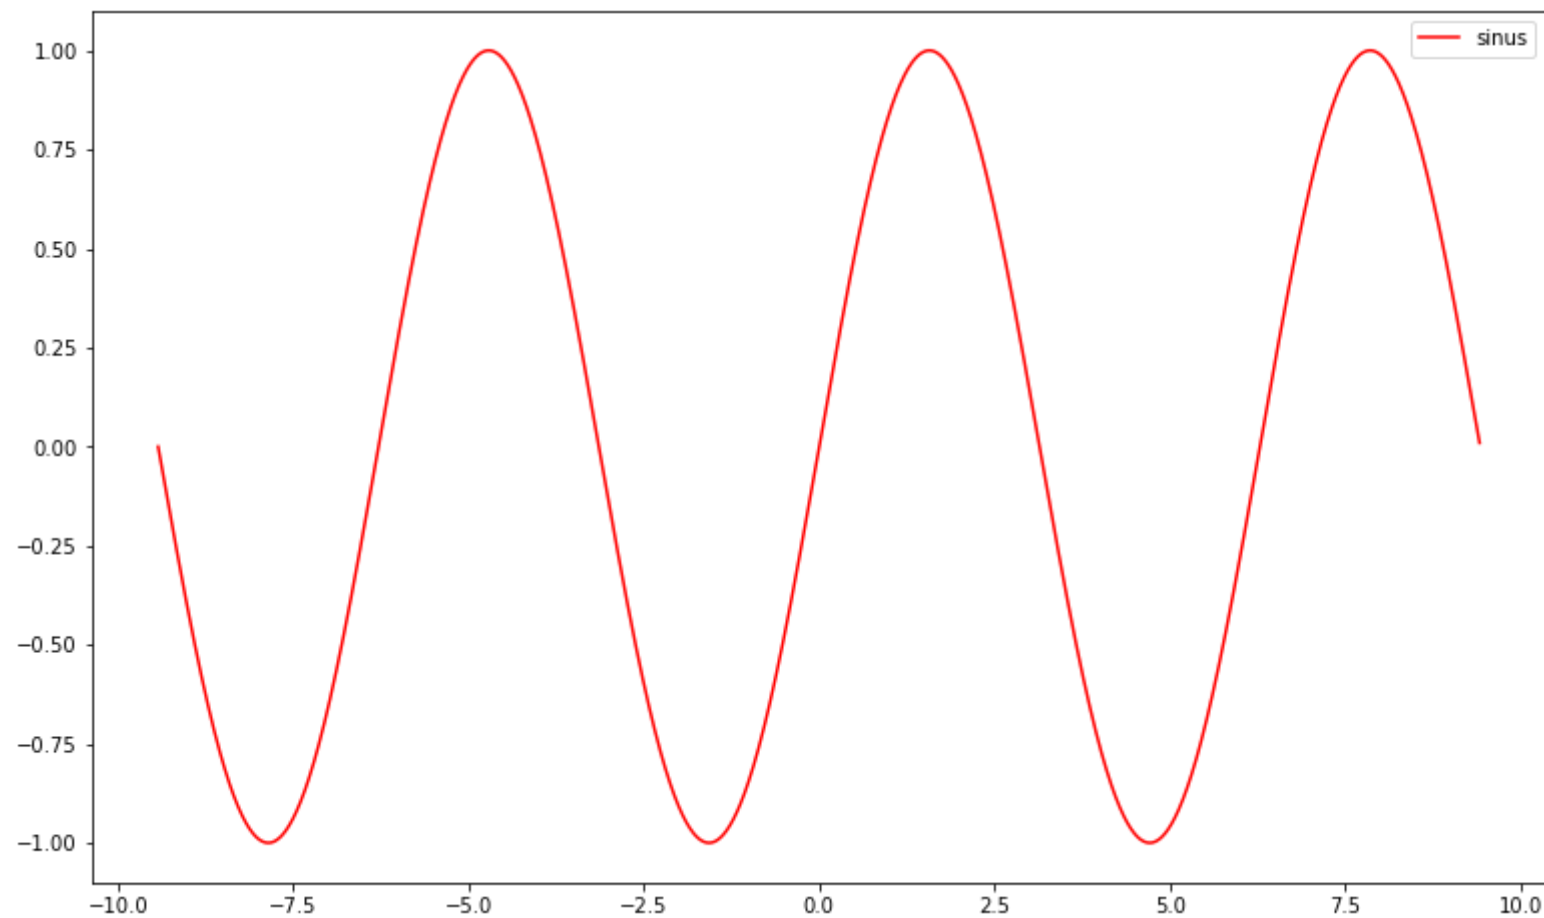
```
In [ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import time
```

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [ ]: import tensorflow as tf
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.utils import plot_model
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [ ]: X = np.arange(-3*np.pi, 3*np.pi, 0.01)
y = np.sin(X)

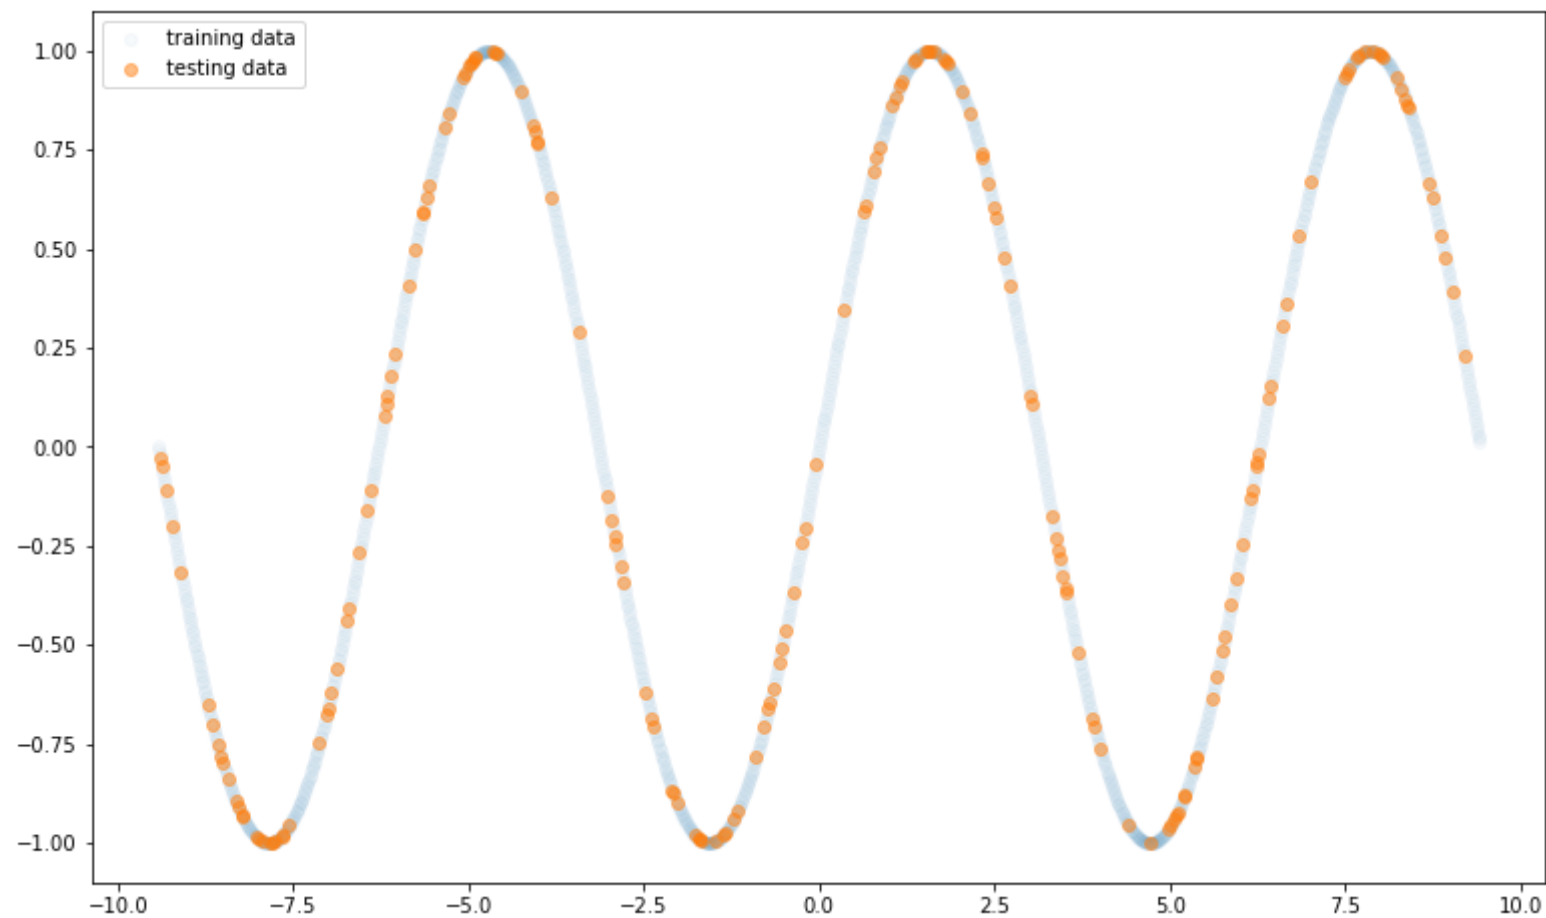
plt.figure(figsize=(13,8))
plt.plot(X, y, label='sinus', color='red')
plt.legend()
plt.show()
```



```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1)

plt.figure(figsize=(13,8))
plt.scatter(X_train, y_train, alpha = 0.03, label='training data')
plt.scatter(X_test, y_test, alpha=0.5, label='testing data')

plt.legend()
plt.show()
```



[TODO – Students]

- Build a model with 1 hidden layer in 1 dimension and train it on x_{train} , y_{train} .
- What activation should we use for output layer ?
- What loss should we use ?
- Try different activations for the hidden layer and plot the predictions obtained on x_{test}
- Plot also a learning curve

```
In [ ]: model = ...
```

```
model.compile(...)\n\nhistory = model.fit(...)
```

```
In [ ]: # Plot leaning cuve\nplt.figure(figsize=(13,8))\nplt.plot(history.history['loss'])\nplt.show()
```

```
In [ ]: def plot_prediction(title, model):\n    y_hat_test = model.predict(X_test)\n\n    plt.figure()\n    plt.scatter(X_test, y_test, label = 'ground_truth', alpha=0.1)\n    plt.scatter(X_test, y_hat_test, label = 'predicted', alpha=0.5)\n    plt.legend()\n\n    plt.title(title)\n    plt.show()
```

```
In [ ]: plot_prediction("my title", model)
```

[TODO – Students]

Try adding layers and increasing the layers dimension to better fit the test data. You can use the following function to quickly build your models.

- Try different `n_layers` (for example 1, 10, 100)
- Try different `hidden_dim` (for example 32, 128, 256, 512)
- Try different batch size
- Try to understand the `patience` parameters of early stopping

```
In [ ]: def build_sin_regression(activation, n_layers, hidden_dim):\n    input = Input(shape=(1,), name='input')\n\n    for i in range(n_layers):
```

```
if i==0:
    x = Dense(input_shape=(1,), units=hidden_dim, activation=activation, name='layer_'+str(i))(input)
else:
    x = Dense(units=hidden_dim, activation=activation, name='layer_'+str(i))(x)

output = Dense(1, activation='linear', name='output')(x)
model = Model(input, output, name = 'sinus_regression')
return model
```

```
In [ ]: model = build_sin_regression(activation = XXX, n_layers = XXX, hidden_dim = XXX)
model.compile(loss='mse', optimizer='adam')
model.summary()
```

```
In [ ]: plot_model(model)
```

```
In [ ]: callbacks_list = [EarlyStopping(monitor='val_loss', min_delta=0.005, patience=20, verbose=2, mode='min', restore_best_weights=True)]

history = model.fit(X_train, y_train, validation_split = 0.1, callbacks=callbacks_list, batch_size=32, epochs=20)
```

```
In [ ]: plt.figure(figsize=(13,8))
plt.plot(history.history['loss'], label="loss")
plt.plot(history.history['val_loss'], label="val_loss")
plt.legend()
plt.show()
```

```
In [ ]: y_hat_test = model.predict(X_test)

plt.scatter(X_test, y_test, label = 'ground_truth', alpha=0.1)
plt.scatter(X_test, y_hat_test, label = 'predicted', alpha=0.5)

plt.legend()
plt.show()
```

Build an MLP to classify MNIST images

Every MNIST data point has two parts: an image of a handwritten digit and a corresponding label. We'll call the images "x" and the labels "y". Both the training set and test set contain images and their corresponding labels; for example, the training images are `mnist.train.images` and the training labels are `mnist.train.labels`.

```
In [ ]: # Load dataset
        from tensorflow.keras.datasets import mnist

        # the data, shuffled and split between a train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()
        X_train.shape, y_train.shape
```

```
In [ ]: # Reshape the image from 3d to 2d (nb_items, other dime)
```

```
In [ ]: # Normalize the data (input between 0 and 1)
```

```
In [ ]: # One hot encode the label
```

```
In [ ]: # Build MLP model
        # You can use the following function
        def build_MLP(input_shape, activation, layers, nb_class):
            input = Input(shape=(input_shape,), name='input')

            for i, hidden_size in enumerate(layers):
                if i == 0:
                    x = Dense(input_shape=(input_shape,), units=hidden_size, activation=activation, name='layer_'+str(i))(input)
                else:
                    x = Dense(units=hidden_size, activation=activation, name='layer_'+str(i))(x)

            output = Dense(nb_class, activation='softmax', name='output')(x)
            model = Model(input, output, name = 'mnist_classifier')
            model.summary()
            return model

        model = build_MLP(## TO BE COMPLETED ##)
```

```
In [ ]:
```

```
# Compile and fit the model
callbacks_list = [EarlyStopping(monitor='val_accuracy', min_delta=0.005, patience=20,
                                verbose=2, mode='min', restore_best_weights=True)
                  ]

model.compile(loss='categorical_crossentropy', metrics=["accuracy"], optimizer='adam')
history = model.fit('' X'', '' y'', validation_split = 0.1, callbacks=callbacks_list,
                    batch_size=32, epochs=20)
```

```
In [ ]: # Print history keys
history.history.keys()
```

```
In [ ]: # Babysit your model
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(26,8))

ax1.plot(history.history['loss'], label="loss")
ax1.plot(history.history['val_loss'], label="val_loss")
ax1.legend()
ax2.plot(history.history['accuracy'], label="accuracy")
ax2.plot(history.history['val_accuracy'], label="val_accuracy")
ax2.legend()
plt.show()
```

```
In [ ]: # Evaluate the model
score = model.evaluate(X_test, y_test_enc)
print('Test loss:', score[0])
print('Test accuracy', score[1])
```

```
In [ ]: # Modify the network in order to obtain better accuracy (better than 0.96)
```