



MSC. DATA SCIENCE & ARTIFICIAL INTELLIGENCE

INTRODUCTION TO MACHINE LEARNING

Dr. Michel Riveill & Dr. Diane LINGRAND

Final project: Petfinder, predicting adoption

Author: Joris LIMONIER

joris.limonier@hotmail.fr

Due: January 15, 2022

Contents

1	Problem description	1
2	Exploratory Data Analysis	1
3	Problem Solution	1
3.1	Overview	1
3.2	Data Preprocessing	1
3.3	Classification	3
3.3.1	General approach	3
3.3.2	Select Models	3
3.3.3	Fine-tune hyperparameters	4
4	Evaluation & critical view	4

List of Figures

1	Complete pipeline diagram	2
---	-------------------------------------	---

List of Tables

1	Data types per column	1
2	Accuracies of first prospect	3
3	Hyperparameters tested for the top 5 classifiers	4

1 Problem description

The problem we are trying to solve consists of predicting whether an animal will be adopted from a shelter within 30 days, given several pieces of information on this animal. This problem is a clean and reduced version of a [Kaggle competition](#) dating back from 2019.

2 Exploratory Data Analysis

We would like to get some basic information of the data set before diving into the machine learning solution.

The training set has shape (8168×16) and the test set has shape (250×16) , where the column names and data types are summarized in Table 1.

CATEGORICAL		NUMERICAL	TEXT	IMAGE
Type	MaturitySize	Age	Description	Images
Gender	FurLength	Fee		
Breed	Vaccinated			
Color1	Dewormed			
Color2	Sterilized			
Color3	Health			

Table 1: Data types per column

Overall, the data set is very clean as it contains 0 NaN values.

3 Problem Solution

3.1 Overview

The problem can be solved by using a pipeling, which is represented in Figure 1. The pipeline consists of two steps:

- Data Pre-processing
- Classification

We will detail these two steps in the following subsections.

3.2 Data Preprocessing

We saw in section 2 that the data is already fairly clean (*e.g.* in terms of NA's). We still need to process the columns, which we do depending on the type of data they contain. As per Figure 1, we use:

- A Categorical Preprocessor
- A Numerical Preprocessor

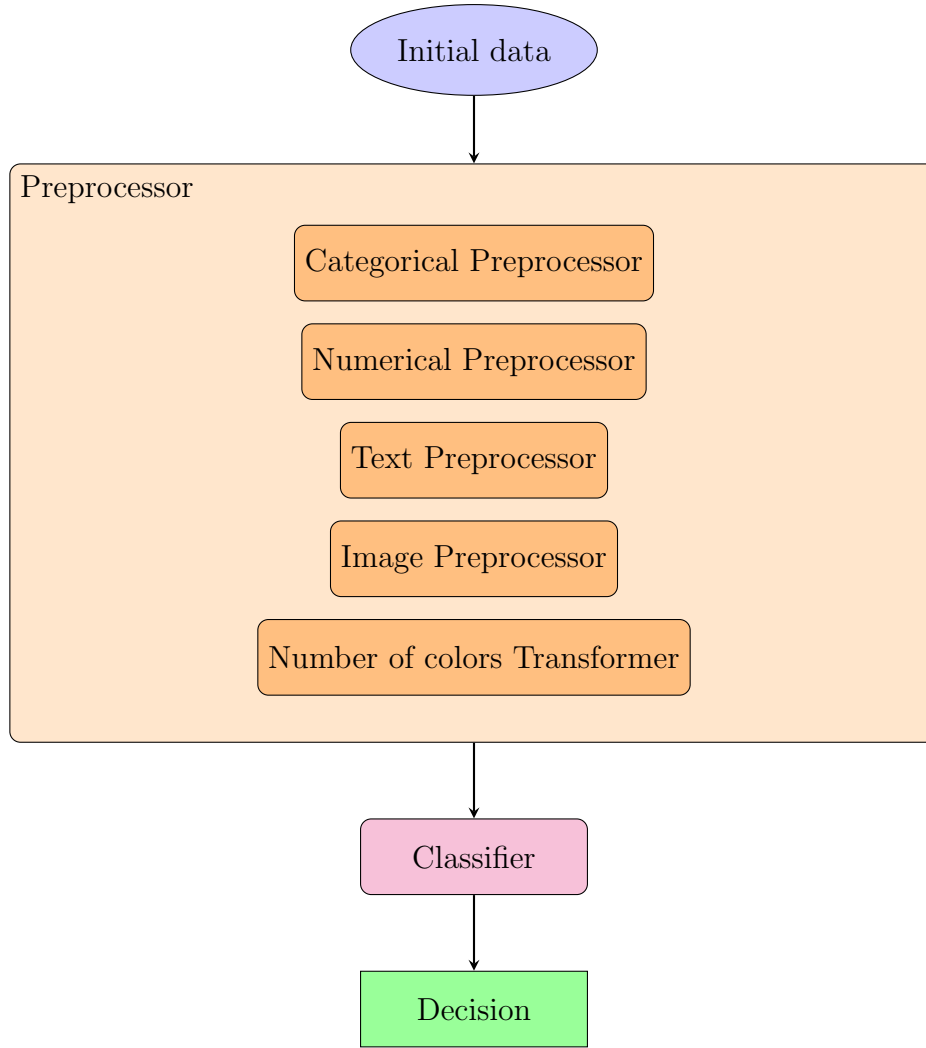


Figure 1: Complete pipeline diagram

- A Text Preprocessor
- A Image Preprocessor
- A Transformer for the number of colors

We detail the above in their respective paragraphs.

Categorical Preprocessor The Categorical Preprocessor is composed of a OneHotEncoder, which encodes categorical features as a one-hot numeric array. That is, for each of the categories, we create a new column with 1 if the initial column is of that category and 0 otherwise.

Numerical Preprocessor The Numerical Preprocessor contains a StandardScaler. Its role is to center and scale each of the numerical variables in order to remove differences in orders of magnitude.

Text Preprocessor The Text Preprocessor contains a TfidfVectorizer, which converts the raw comments from our text feature to a matrix of Term Frequency-Inverse Document Frequency (TF-IDF) features. We used a CountVectorizer initially, but it gave worse results than the TF-IDF, which is why we kept the latter. The CountVectorizer counts the occurrences of each word in the *Description* column. A possible explanation of why the TfidfVectorizer performs better is that the features it produces not-only contain information about the frequency of each term, but they also account for the frequency of each word within the whole document.

Image Preprocessor The Image Preprocessor consists of a custom BOF_extractor. It extracts Scale-Invariant Feature Transforms (SIFTs) and computes Bag Of Features (BOF) on the images from the *Images* column.

Number of colors Transformer The Number of colors Transformer uses Function-Transformer to compute how many colors the animal has (*i.e.* colors different from “Unknown”). It then sets the number of colors as a feature.

3.3 Classification

3.3.1 General approach

Our approach for the Classification part separates in two parts

1. Try various classifiers and evaluate their performance. Then, keep the best 5 classifiers.
2. For the best 5 classifiers, fine-tune their hyperparameters and find the best one.

3.3.2 Select Models

The algorithms tested are presented in Table 2. We did not test XGBoost because it was too slow.

CLASSIFIER	ACCURACY
GradientBoostingClassifier	0.629
RandomForestClassifier	0.623
AdaBoostClassifier	0.612
MLPClassifier	0.602
BernoulliNB	0.600
GaussianNB	0.567
DecisionTreeClassifier	0.559
SVC	0.529
KNeighborsClassifier	0.520
GaussianProcessClassifier	0.509
SGDClassifier	0.509

Table 2: Accuracies of first prospect

3.3.3 Fine-tune hyperparameters

As shown in Table 2, the five best algorithms are GradientBoostingClassifier, RandomForestClassifier, AdaBoostClassifier, MLPClassifier and BernoulliNB. The next step is to select several hyperparameters for each of them and test them with GridSearchCV.

CLASSIFIER	HYPERPARAMETER	VALUES TESTED
GradientBoostingClassifier	learning_rate	0.01, 0.02, 0.05, 0.1, 0.2
	n_estimators	50, 100, 150
	max_depth	1, 3, 5
	min_samples_split	1, 3, 5
	min_samples_leaf	1, 2, 3
	max_features	log2, sqrt
RandomForestClassifier	criterion	gini, entropy
	max_depth	100, 200, 300, 400, None
	min_samples_leaf	1, 2, 3, 4
	min_samples_split	8, 10, 12
	n_estimators	60, 80, 100
AdaBoostClassifier	n_estimators	10, 20, 30, 50, 70
	learning_rate	.5, .6, .7, .8, 1., 1.2
	algorithm	SAMME.R, SAMME
MLPClassifier	hidden_layer_sizes	(50,100,50), (100,)
	learning_rate	constant, invscaling, adaptive
	solver	adam
	activation	identity, logistic, tanh, relu
	beta_1	.8, .9, .99
	beta_2	.99, .999, .9999
BernoulliNB	alpha	.01, .05, .1, .2, .3, .5, 1., 2.
	binarize	0., .4, .5, .6, .65, .7, .8, 1., 2.
	fit_prior	True, False

Table 3: Hyperparameters tested for the top 5 classifiers

4 Evaluation & critical view

Glossary

BOF Bag Of Features. 3

SIFT Scale-Invariant Feature Transform. 3

TF-IDF Term Frequency-Inverse Document Frequency. 3