

Image descriptors

Diane Lingrand



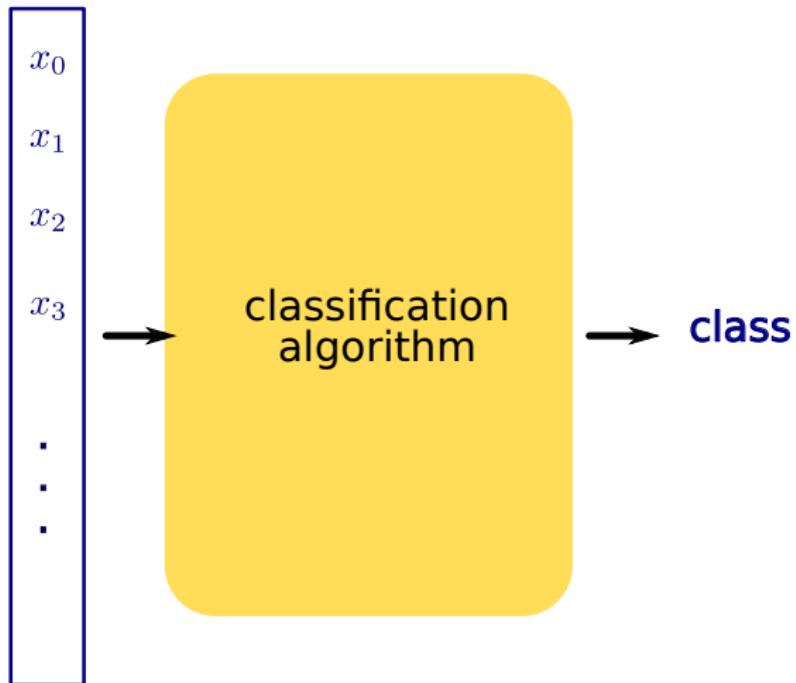
UNIVERSITÉ
CÔTE D'AZUR

Master Data Science M1

2020 - 2021



how to
describe
an image



- Basics on image processing
 - color, edges, filtering, points
- Image descriptors
 - SIFT, SURF, HOG
- Deep features



0.alexandre.hiltcher.009.txt

- RGB = Red, Green, Blue. Usually : 1 byte per component. The most common.
- YUV and variants (linear transformation)
- HSV = Hue Saturation Value
- CIE LAB

$$Y = 0.299 R + 0.587 G + 0.114 B$$

- Betacam ($Y \text{ pb pr}$)
 - $\text{pb} = 0.5 (B - Y) / (1 - 0.114)$
 - $\text{pr} = 0.5 (R - Y) / (1 - 0.299)$
- Digital Video System ($Y \text{ Cb Cr}$)
 - $\text{Cb} = 128 + 112 \text{ pb}$
 - $\text{Cr} = 128 + 112 \text{ pr}$
- YUV
 - $U = 0.493 (B - Y)$
 - $V = 0.877 (R - Y)$
- NTSC YIQ
 - $I = 0.6 R - 0.28 G - 0.32 B$
 - $Q = 0.21 R - 0.52 G + 0.31 B$



- Hunter color space(1948) : Lab
- CIELAB (1976) : L*a*b*
 - L* : lightness (cubic root of relative luminance)
 - a* and b* : distance to grey
 - a* : from green to red
 - b* : from blue to yellow

$$\begin{aligned}L^* &= 116f(Y/Y_n) - 16 \\a^* &= 500(f(X/X_n) - f(Y/Y_n)) \\b^* &= 200(f(Y/Y_n) - f(Z/Z_n))\end{aligned}$$

with

$$f(t) = \begin{cases} t^{1/3} & \text{if } t > \left(\frac{6}{29}\right)^3 \\ \frac{1}{3}\left(\frac{29}{6}\right)^2 t + \frac{4}{29} & \text{else} \end{cases}$$

Using illumination D65, coordinates of white :

$$X_n = 95.047, Y_n = 100.000, Z_n = 108.883$$

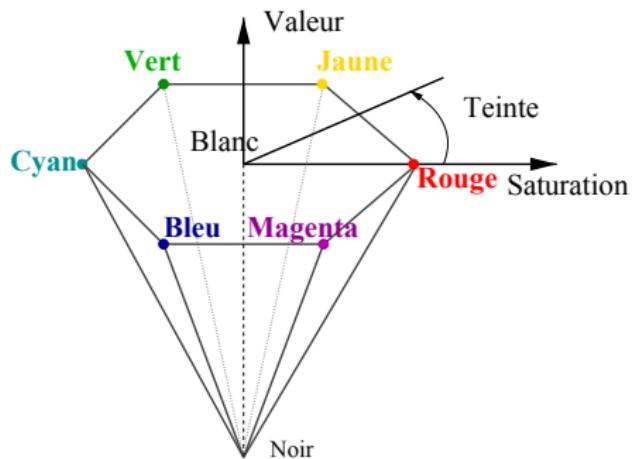
$$V = \max(R, G, B)$$

$$S = \begin{cases} \frac{V - \min(R, G, B)}{V} & \text{if } V \neq 0 \\ 0 & \text{else} \end{cases}$$

$$H = \begin{cases} \frac{60(G-B)}{V - \min(R, G, B)} & \text{if } V = R \\ \frac{120 + 60(B-R)}{V - \min(R, G, B)} & \text{if } V = G \\ \frac{240 + 60(R-G)}{V - \min(R, G, B)} & \text{if } V = B \end{cases}$$

If $H < 0$ then $H = H + 360$.

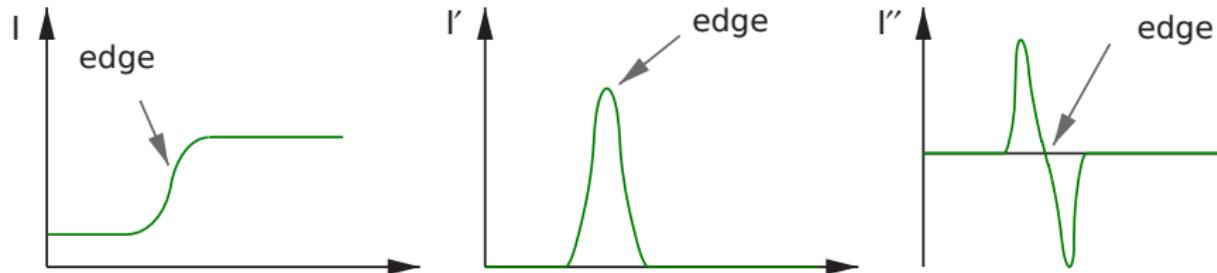
Output : $0 \leq V \leq 1$, $0 \leq S \leq 1$,
 $0 \leq H \leq 360$



How to describe a region

- main color
- color histogram
- texture (Gabor filters, Local Binary Pattern)
- edges, corners

- maximum of derivatives
- derivatives sensitive to noise \Rightarrow need to smooth



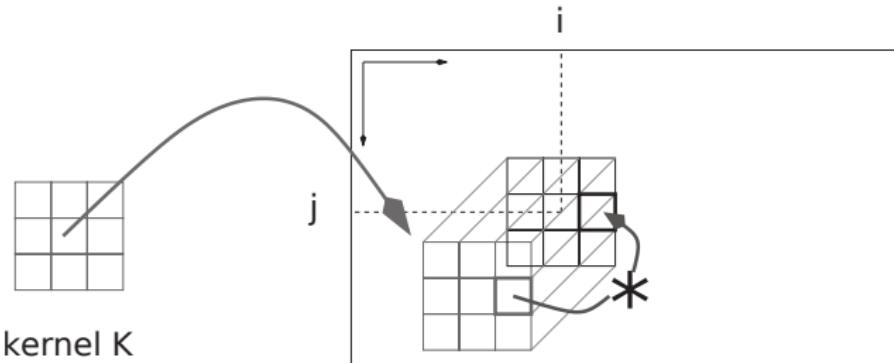
Convolution

- For f and g discrete functions :

$$f * g(x, y) = \sum_{u=-\infty}^{\infty} \sum_{v=-\infty}^{\infty} f(x, y)g(u - x, v - y)$$

- Convolution of an image I_1 by a kernel K of dimensions $(2p + 1) \times (2q + 1)$:

$$I_2[i][j] = \sum_{k=0}^{2p} \sum_{l=0}^{2q} I_1[i - k + p][j - l + q]K[k][l]$$



Average filter smoothing

Dimension 3x3 :

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} / 9$$

Dimension $(2p+1) \times (2p+1)$:

$$\begin{bmatrix} 1 & 1 & \dots & 1 \\ 1 & 1 & \dots & 1 \\ \vdots & & & \ddots \\ 1 & 1 & \dots & 1 \end{bmatrix} / (2p + 1)^2$$



3x3



5x5



7x7



9x9

Gaussian smoothing

$$G_\sigma(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{|x|^2}{2\sigma^2}}$$

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 4 & 1 \\ 0 & 1 & 0 \end{bmatrix} / 8 \quad \text{or} \quad \begin{bmatrix} 1 & 1 & 1 \\ 1 & 8 & 1 \\ 1 & 1 & 1 \end{bmatrix} / 16$$



3x3



5x5



7x7

Sobel detection (1970)

combines smoothing in one direction $[1, 2, 1]$ and derivative in a perpendicular direction using $[1, 0, -1]$

$$\frac{dI}{dx} = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} / 4 \quad \frac{dI}{dy} = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} / 4$$

Angle between horizontal and normal vector :

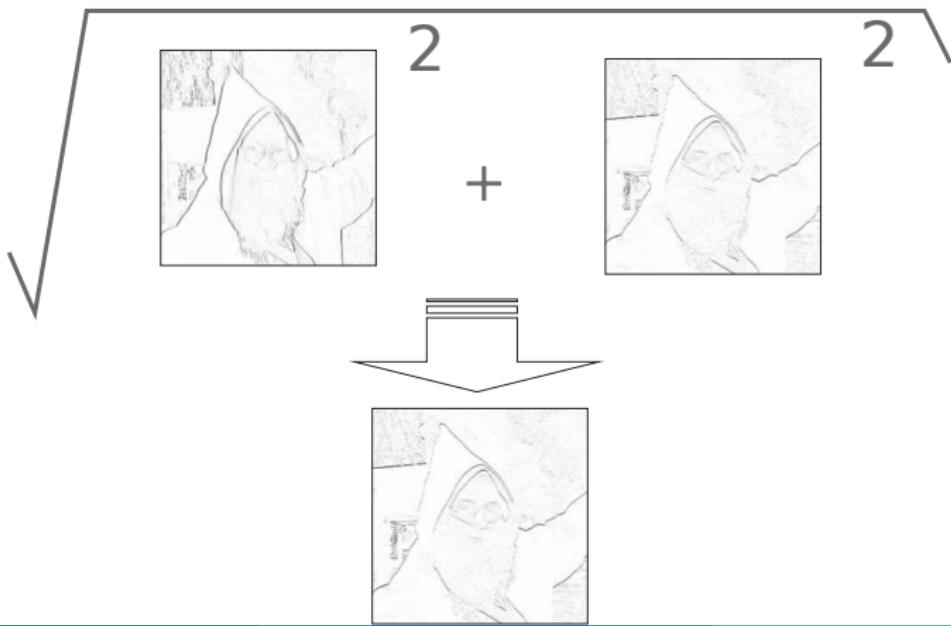
$$\arctan \left(\left(\frac{dI}{dy} \right) / \left(\frac{dI}{dx} \right) \right)$$

Example of Sobel detection

vertical
edges
extraction



horizontal
edges
extraction



Prewitt edge detector(1970)

$$\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} /3, \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & -1 \\ 0 & -1 & -1 \end{bmatrix} /3, \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} /3, \begin{bmatrix} 0 & 1 & 1 \\ -1 & 0 & 1 \\ -1 & -1 & 0 \end{bmatrix} /3$$



Kirsch edge detector(1971)

$$\begin{bmatrix} 5 & -3 & -3 \\ 5 & 0 & -3 \\ 5 & -3 & -3 \end{bmatrix} / 15, \begin{bmatrix} 5 & 5 & -3 \\ 5 & 0 & -3 \\ -3 & -3 & -3 \end{bmatrix} / 15, \begin{bmatrix} 5 & 5 & 5 \\ -3 & 0 & -3 \\ -3 & -3 & -3 \end{bmatrix} / 15,$$
$$\begin{bmatrix} -3 & 5 & 5 \\ -3 & 0 & 5 \\ -3 & -3 & -3 \end{bmatrix} / 15, \begin{bmatrix} -3 & -3 & 5 \\ -3 & 0 & 5 \\ -3 & -3 & 5 \end{bmatrix} / 15, \begin{bmatrix} -3 & -3 & -3 \\ -3 & 0 & 5 \\ -3 & 5 & 5 \end{bmatrix} / 15,$$
$$\begin{bmatrix} -3 & -3 & -3 \\ -3 & 0 & -3 \\ 5 & 5 & 5 \end{bmatrix} / 15, \begin{bmatrix} 5 & 0 & -3 \\ 5 & 5 & -3 \end{bmatrix} / 15$$



D. Lingrand (M1)



Image descriptors

Laplace connexity 4



Laplace connexity 8



Laplace connexity 8

after gaussian smoothing (7x7)



Zeros crossings



Laplacian as difference between two gaussians.

Difference between gaussian smoothing 7×7 and 5×5 :



Corners detection

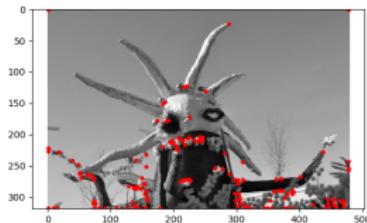
- many approaches (edges points linkage, curvature, models, ...)
- the more famous : Harris and Stephens, 1988

$$\mathcal{O} = \det(\hat{C}(x, y)) - k (\text{trace}(\hat{C}(x, y)))^2$$

$$\hat{C}(x, y) = \begin{pmatrix} \widehat{\left(\frac{\partial I(x,y)}{\partial x} \right)^2} & \widehat{\frac{\partial I(x,y)}{\partial x}} \widehat{\frac{\partial I(x,y)}{\partial x}} \\ \widehat{\frac{\partial I(x,y)}{\partial x}} \widehat{\frac{\partial I(x,y)}{\partial x}} & \widehat{\left(\frac{\partial I(x,y)}{\partial y} \right)^2} \end{pmatrix} \text{ with } k = 0.04$$

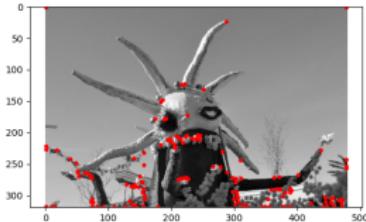
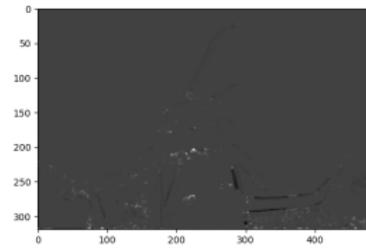
$\widehat{}$: gaussian smoothing

- computation of first derivatives with Gaussian smoothing
- compute \mathcal{O}
- thresholding



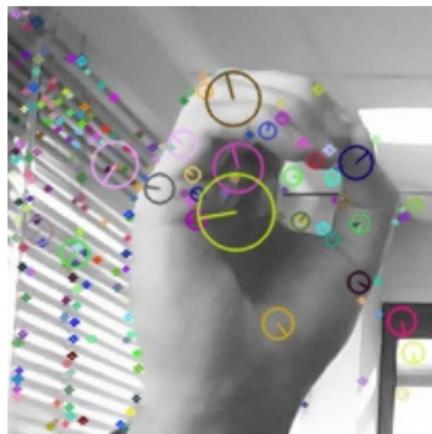
Harris example

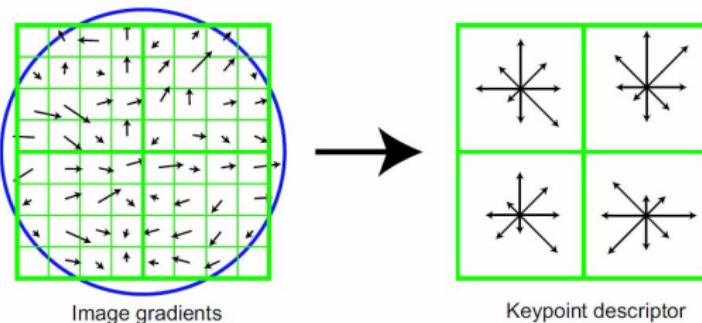
```
from skimage.feature import corner_harris, corner_peaks
img = skimage.io.imread('carnaval4.jpg', as_gray=True)
pts=corner_peaks(corner_harris(img), min_distance=1)
print( pts.shape[0], ' points found')
import matplotlib.pyplot as plt
plt.imshow(img, cmap='gray')
plt.scatter(y=pts[:,0],x=pts[:,1],c='r',s=10)
plt.show()
```



SIFT = Scale Invariant Feature Transform

- Detector
 - multi-scale
 - DOG laplacian (Difference Of Gaussians)
- Descriptor
 - edges orientations in the neighborhood





- vectors of 128 integers
- 4 steps :
 - interest points detection
 - gradients orientation in the neighborhood (16x16 pixels divided in 16 blocks of size 4x4)
 - orientation histogram (quantified on 8 values in blocks of 4x4 pixels)
 - $8 \times 4 \times 4 = 128$
 - normalisation

SIFT descriptor



0.alexandre.hiltcher.006.png
252 descriptors



0.alexandre.hiltcher.009.png
182 descriptors

SIFT implementation in OpenCV

```
import cv2
import numpy as np
img = cv2.imread('carnaval.jpg')
gray= cv2.cvtColor(img,cv2.COLOR_BGR2GRAY)
sift = cv2.xfeatures2d.SIFT_create()
kp = sift.detect(gray,None)
cv2.drawKeypoints(gray,kp,img,\\
    flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
cv2.imwrite('carnavalSIFT.jpg',img)
```

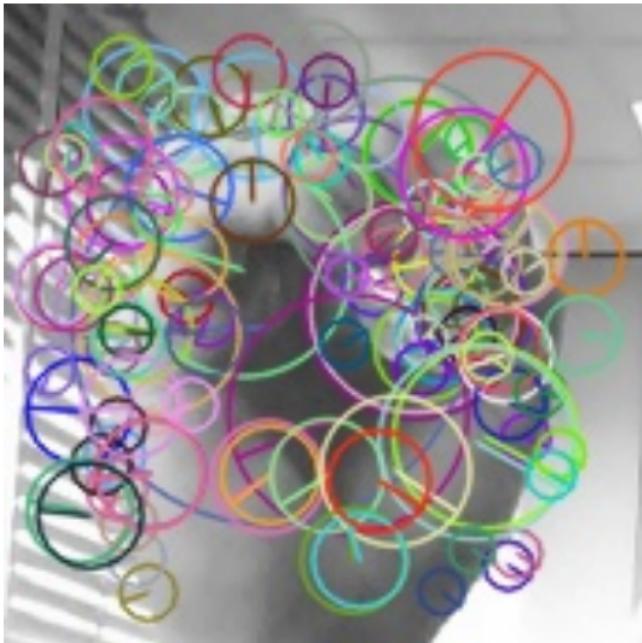


SURF = Speeded-Up Robust Features

default size : 64



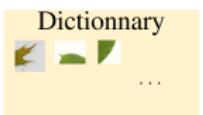
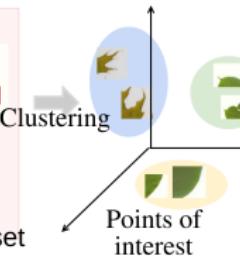
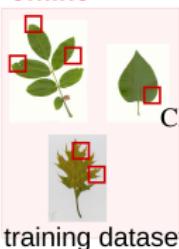
0.alexandre.hiltcher.006.png
184 descriptors



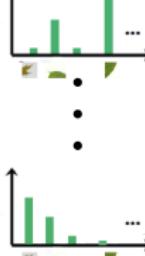
0.alexandre.hiltcher.009.png
121 descriptors

Bag Of Words (BOW)

offline



BoW Histogram

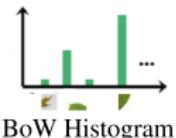
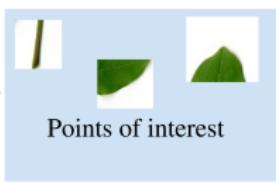
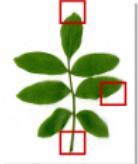


[23 3 7 12 8 1 ...]



[2 43 7 11 19 5 ...]

online

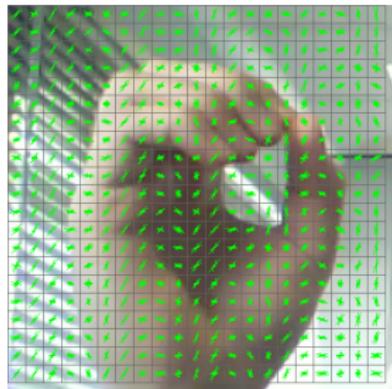


[32 7 53 12 ...]

HOG : Histogram of Gradients

- gradient computation ($[-1 \ 0 \ 1]$ et $[-1 \ 0 \ 1]^T$)
- histogram construction
 - squared cells (from 4×4 to 12×12 pixels)
 - discretisation on 9 angle values
 - pixel votes proportional to gradient amplitude
- blocks construction
 - 1 block = several cells
 - normalisation of blocks
- HOG = concatenation of histograms

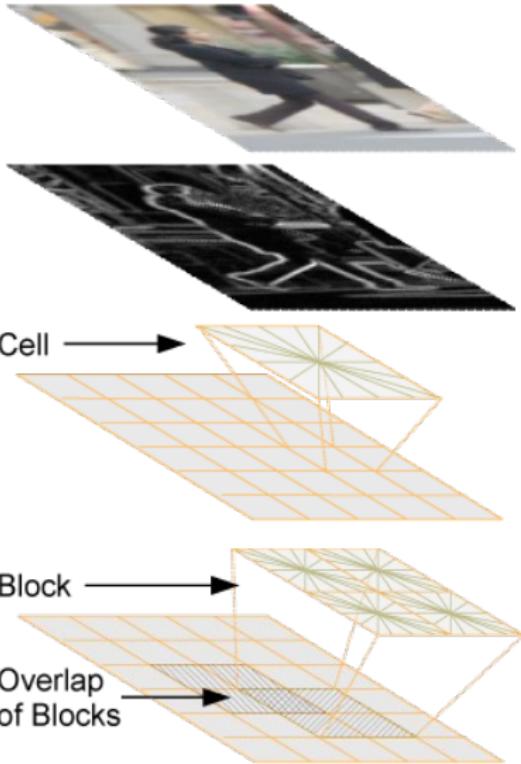
Next slides from presentation by Seeman.



0.alexandre.hiltcher.009.png

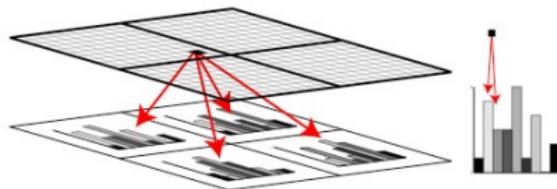
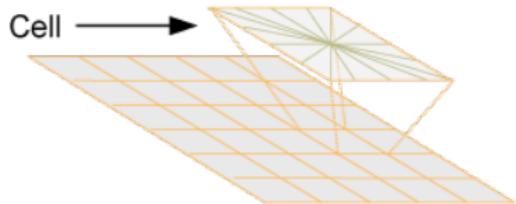
Descriptor

1. Compute gradients on an image region of 64×128 pixels
2. Compute histograms on ‘cells’ of typically 8×8 pixels (i.e. 8×16 cells)
3. Normalize histograms within overlapping blocks of cells (typically 2×2 cells, i.e. 7×15 blocks)
4. Concatenate histograms



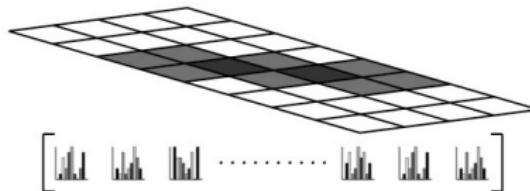
Cell histograms

- 9 bins for gradient orientations (0-180 degrees)
- Filled with magnitudes
- Interpolated trilinearly:
 - Bilinearly into spatial cells
 - Linearly into orientation bins

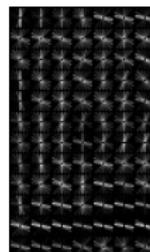
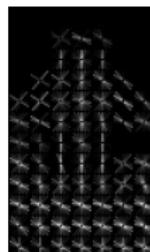
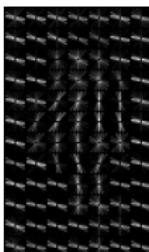


Final Descriptor

- Concatenation of Blocks



- Visualization:



- many convolution filters :
 - many edge detectors (Roberts, Sobel, Kirch, Prewitt, Laplace)
 - smoothing (mean, Gaussian, ...)
 - edge boosters
- which ones to choose ? thresholds ?
- other approaches :
 - Canny, Deriche
 - Deformable models (parametric curves, levelsets)
- CNN is learning the filters coefficients ... and gets the best results

- using already trained CNN
 - Xception (2016)
 - VGG16, VGG19 (2014)
 - ResNet, ResNetV2, ResNeXt (2015-2016)
 - InceptionV3 (2015)
 - InceptionV4, InceptionResNetV2 (2016)
 - MobileNets (2017)
 - DenseNet (2017)
- different options
 - tensor
 - pooling (reducing the size) : average or max
- try them using keras : <https://keras.io/applications/>