

Optimization - Quiz 1

Joris LIMONIER

January 18, 2022

1 Question 1

No, it is not possible to pursue directly this goal because we don't know the true distribution \mathcal{D} . This is a fundamental difference between Machine Learning and Statistics. We know however that our data was sampled from \mathcal{D} , and we know by the law of large numbers that our empirical loss will converge towards the expected value of the loss, as the number of samples increases.

Correction No because you don't know the underlying distribution \mathcal{D} .

2 Question 2

One way to learn a model is by performing a train-test split in order to verify that our function (that we train on the train set) performs well on a set that is never seen before (*i.e.* the test set). We need to find the right model with not too many parameters (otherwise we over-fit our training set), and not too few parameters (otherwise we under-fit and do not learn enough from data). Other solutions, especially in case with small data sets, include K-fold cross-validation. One of its variations consists in disregarding a fold of the data set, while looking only at the $K - 1$ other folds. Then repeat this step with the other folds.

Correction Instead of working with the true loss, work with the empirical loss.