

Theory of Statistical Learning

Part II

Damien Garreau

Université Côte d'Azur

2022

Outline

1. Linear predictors

- Linear classification

- Linear regression

- Ridge regression

- Polynomial regression

- Logistic regression

- Support vector machines

2. Kernel methods

- Positive semi-definite kernels

- Reproducing kernel Hilbert spaces

- More examples

- The kernel trick and applications

- The representer theorem

- Kernel ridge regression

3. Tree-based classifiers

- Partition rules

- Random forests

4. Boosting

1. Linear predictors

1.1. Linear classification

Linear functions

- ▶ $\mathcal{X} = \mathbb{R}^d$, $\mathcal{Y} = \mathbb{R}$
- ▶ thus $x_i = (x_{i,1}, x_{i,2}, \dots, x_{i,d})^\top$
- ▶ we consider no bias term (otherwise *affine*):

$$\{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}.$$

- ▶ **Reminder:** given two vectors $u, v \in \mathbb{R}^d$,

$$\langle u, v \rangle = u^\top v = \sum_{j=1}^d u_j v_j.$$

- ▶ binary classification: 0-1 loss, $\mathcal{Y} = \{-1, +1\}$
- ▶ **Important:** compose h with $\phi : \mathbb{R} \rightarrow \mathcal{Y}$ (typically the sign)

The sign function





Figure: the sign function

Halfspaces

- ▶ thus our function class is

$$\mathcal{H} = \{x \mapsto \text{sign}(w^\top x), w \in \mathbb{R}^d\}.$$

- ▶ gives label +1 to vector pointing in the same direction as w



VC dimension of halfspaces

Proposition: the VC dimension of halfspaces in dimension d is exactly $d + 1$.


► **Consequence:** \mathcal{H} is PAC learnable with sample complexity

$$\Omega\left(\frac{d + \log(1/\delta)}{\varepsilon}\right).$$

Linearly separable data

- ▶ **Important assumption:** data is linearly separable
- ▶ that is, there is a $w^* \in \mathbb{R}^d$ such that

$$y_i = \text{sign}(\langle w^*, x_i \rangle) \quad \forall 1 \leq i \leq n.$$



Linear programming

- ▶ **Empirical risk minimization:** recall that we are looking for w such that

$$\hat{\mathcal{R}}_S(w) = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{y_i \neq \text{sign}(w^\top x_i)}$$

is minimal

- ▶ **Question:** how to solve this?
- ▶ we want $y_i = \text{sign}(w^\top x_i)$ for all $1 \leq i \leq n$
- ▶ equivalent formulation: $y_i \langle w, x_i \rangle > 0$
- ▶ we know that there is a vector that satisfies this condition (w^*)
- ▶ let us set $\gamma = \min_i \{y_i \langle w^*, x_i \rangle\}$ and $\bar{w} = w^* / \gamma$
- ▶ we have shown that there is a vector such that $y_i \langle \bar{w}, x_i \rangle \geq 1$ for any $1 \leq i \leq n$ (and it is an ERM)

Linear programming, ctd.

- ▶ define the matrix $A \in \mathbb{R}^{n \times d}$ such that

$$A_{i,j} = y_i x_{i,j}.$$

- ▶ **Intuition:** observations \times labels
- ▶ remember that we have the ± 1 label convention
- ▶ define $v = (1, \dots, 1)^\top \in \mathbb{R}^n$
- ▶ then we can rewrite the above problem as

$$\text{maximize } \langle u, w \rangle \text{ subject to } Aw \leq v.$$

- ▶ we call this sort of problems **linear programs**¹
- ▶ solvers readily available, e.g., `scipy.optimize.linprog` if you use Python

¹Boyd, Vandenberghe, *Convex optimization*, Cambridge University Press, 2004

The perceptron

- ▶ another possibility: the *perceptron*²
- ▶ **Idea:** iterative algorithm that constructs $w^{(1)}, w^{(2)}, \dots, w^{(T)}$
- ▶ update rule: at each step, find i that is misclassified and set

$$w^{(t+1)} = w^{(t)} + y_i x_i .$$

- ▶ **Question:** why does it work?
- ▶ pushes w in the right direction:

$$y_i \langle w^{(t+1)}, x_i \rangle = y_i \langle w^{(t)} + y_i x_i, x_i \rangle = y_i \langle w^{(t)}, x_i \rangle + \|x_i\|^2$$

- ▶ remember, we want $y_i \langle w, x_i \rangle > 0$ for all i

²Rosenblatt, *The perceptron, a perceiving and recognizing automaton*, tech report, 1957

1.2. Linear regression

Least squares

- ▶ regression \Rightarrow squared-loss function

$$\ell(y, y') = (y - y')^2.$$

- ▶ still looking at linear functions:

$$\mathcal{H} = \{h : x \mapsto \langle w, x \rangle \text{ s.t. } w \in \mathbb{R}^d\}.$$

- ▶ empirical risk in this context:

$$\hat{\mathcal{R}}_S(h) = \frac{1}{n} \sum_{i=1}^n (w^\top x_i - y_i)^2 = F(w).$$

- ▶ also called **mean squared error**
- ▶ empirical risk minimization: we want to minimize $w \mapsto F(w)$ with respect to $w \in \mathbb{R}^d$
- ▶ F is a **convex, smooth** function

Least squares, ctd.

- ▶ let us compute the gradient of F :

$$\begin{aligned}\frac{\partial F}{\partial w_j}(w) &= \frac{1}{n} \sum_{i=1}^n \frac{\partial}{\partial w_j} (w^\top x_i - y_i)^2 \\&= \frac{1}{n} \sum_{i=1}^n 2 \cdot \frac{\partial}{\partial w_j} (w^\top x_i - y_i) \cdot (w^\top x_i - y_i) \\&= \frac{1}{n} \sum_{i=1}^n 2 \cdot \frac{\partial}{\partial w_j} (\cdots + w_j x_{i,j} + \cdots - y_i) \cdot (w^\top x_i - y_i) \\ \frac{\partial F}{\partial w_j}(w) &= \frac{2}{n} \sum_{i=1}^n x_{i,j} \cdot (w^\top x_i - y_i).\end{aligned}$$

Least squares, ctd.

- ▶ it is more convenient to write $\nabla F(w) = 0$ in matrix notation
- ▶ define $X \in \mathbb{R}^{n \times d}$ the matrix such that line i of X is observation x_i
- ▶ one can check that, for any $1 \leq j, k \leq d$,

$$(X^\top X)_{j,k} = \sum_{i=1}^n x_{i,j} x_{i,k}.$$

- ▶ thus

$$\begin{aligned}(X^\top X w)_j &= \sum_{k=1}^d (X^\top X)_{j,k} w_k \\ &= \sum_{k=1}^d \sum_{i=1}^n x_{i,j} x_{i,k} w_k \\ &= \sum_{i=1}^n x_{i,j} w^\top x_i.\end{aligned}$$

Least squares, ctd.

- ▶ thus, if we define

$$A = X^\top X = \sum_{i=1}^n x_i x_i^\top \in \mathbb{R}^{d \times d} \text{ and } b = X^\top y = \sum_{i=1}^n y_i x_i \in \mathbb{R}^d,$$

solving $\nabla F(w) = 0$ is equivalent to solving

$$Aw = b.$$

- ▶ if A is invertible, straightforward:

$$\hat{w} = A^{-1}b$$

- ▶ computational cost: $\mathcal{O}(d^3)$ (inversion is actually a bit less)
- ▶ what happens when A is not invertible?

Singular value decomposition

- ▶ since A is symmetric, it has an eigendecomposition

$$A = VDV^{\top},$$

with $D \in \mathbb{R}^d$ diagonal and V orthonormal

- ▶ define D^+ such that

$$D_{i,i}^+ = 0 \text{ if } D_{i,i} = 0 \text{ and } D_{i,i}^+ = \frac{1}{D_{i,i}} \text{ otherwise.}$$

- ▶ define $A^+ = VD^+V^{\top}$

- ▶ then we set

$$\hat{w} = A^+ b.$$

Singular value decomposition, ctd.

- ▶ why did we do that?
- ▶ let v_i denote the i th column of V , then

$$A\hat{w} = AA^+b \quad (\text{definition of } \hat{w})$$

$$= VDV^\top VD^+V^\top b \quad (\text{definition of } A^+)$$

$$= VDD^+V^\top b \quad (V \text{ is orthonormal})$$

$$A\hat{w} = \sum_{i:D_{i,i} \neq 0} v_i v_i^\top b.$$

- ▶ in definitive, $A\hat{w}$ is the projection of b onto the span of v_i such that $D_{i,i} \neq 0$
- ▶ since the span of these v_i is the span of the x_i and b is in the linear span of the x_i , we have $A\hat{w} = b$
- ▶ cost of SVD: $\mathcal{O}(dn^2)$ if $d > n$ (SVD of X)

Exercise

Exercise: Of course, one does not have to use the squared loss. Instead, we may prefer to use

$$\ell(y, y') = |y - y'|.$$

1. show that, for any $v \in \mathbb{R}^d$,

$$\|v\|_1 = \min_z \mathbf{1}^\top z \quad \text{subject to} \quad z \geq |v|.$$

2. deduce that ERM with the absolute value loss function is equivalent to minimizing the linear function $\sum_{i=1}^n s_i$, where the s_i satisfy linear constraints
3. write this as a linear program, that is, find $A \in \mathbb{R}^{2n \times (n+d)}$, $v \in \mathbb{R}^{d+n}$, and $b \in \mathbb{R}^{2n}$ such that the problem can be written

$$\text{minimize } c^\top v \quad \text{subject to} \quad Av \leq b.$$

Correction of the exercise

1. We have $|v| \geq |v|$ and $\mathbf{1}^\top |v| = \|v\|_1$.
2. In that case, the empirical risk can be written

$$\hat{\mathcal{R}}_S(w) = \frac{1}{n} \sum_{i=1}^n |y_i - w^\top x_i|.$$

We deduce the result from question 1.

3. One possibility is to define $v = (w_1, \dots, w_d, s_1, \dots, s_n)^\top \in \mathbb{R}^{n+d}$,
 $c = (0, \dots, 0, 1, \dots, 1)^\top \in \mathbb{R}^{d+n}$, $b = (y_1, \dots, y_n, -y_1, \dots, -y_n)^\top \in \mathbb{R}^{2n}$, and

$$A = \begin{pmatrix} X & -I_n \\ -X & -I_n \end{pmatrix} \in \mathbb{R}^{2n \times (n+d)},$$

with $X \in \mathbb{R}^{n \times d}$ the matrix whose lines are the x_i s and I_n the identity matrix.

Recap

- ▶ **What happens when we invoke** `sklearn.linear_model.LinearRegression` with default parameters?
- ▶ `fit_intercept` is `True` → assumes that the data is not centered (our maths are not totally accurate)
- ▶ `normalize` is `False` → we are responsible for the normalization of our data
- ▶ behind the scenes, calls `scipy.linalg.lstsq` when fitting, which itself calls LAPACK (Linear Algebra PACKage)³
- ▶ LAPACK is coded in Fortran90

³<http://www.netlib.org/lapack/>

1.3. Ridge regression

Ridge regression

- ▶ same hypothesis class: linear functions

$$\mathcal{H} = \{h : x \mapsto w^\top x, w \in \mathbb{R}^d\}$$

- ▶ squared loss:

$$\ell(y, y') = (y - y')^2.$$

- ▶ **Idea:** regularization:

$$\text{minimize } \left\{ \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2 \right\},$$

with $\|u\|^2 = u_1^2 + \dots + u_d^2$ and $\lambda > 0$ a *regularization parameter*

Exercise

Exercise: Let $(x_1, y_1), \dots, (x_n, y_n) \in \mathbb{R}^d \times \mathbb{R}$ be n given training samples. For any $w \in \mathbb{R}^d$, set

$$F(w) = \frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 + \lambda \|w\|^2 .$$

Notice that F is a convex smooth function.

1. show that the minimizer \hat{w} satisfies

$$(X^\top X + n\lambda I_d) w = X^\top y .$$

2. show that $X^\top X + n\lambda I_d$ is an invertible matrix

Correction of the exercise

1. Let $1 \leq j \leq d$ and let us compute $\partial_j F$:

$$\begin{aligned}\frac{\partial F}{\partial w_j}(w) &= \frac{\partial}{\partial w_j} \left(\frac{1}{n} \sum_{i=1}^n (y_i - w^\top x_i)^2 \right) + \frac{\partial}{\partial w_j} (\lambda(w_1^2 + \dots w_d^2)) \\ &= \frac{2}{n} \sum_{i=1}^n x_{i,j} \cdot (w^\top x_i - y_i) + 2\lambda w_j ,\end{aligned}$$

where we used the derivation for the least squares. We deduce the result by setting to zero and multiplying by n .

Correction of the exercise, ctd.

2. By contradiction, suppose that $X^\top X + n\lambda I_d$ is not invertible. Then

$$\det(X^\top X + n\lambda I_d) = 0.$$

In other words, $-n\lambda$ is an eigenvalue of $X^\top X$. Since $X^\top X$ is a symmetric matrix, its spectrum is $\subseteq \mathbb{R}$. Moreover, it is positive definite, thus all eigenvalues are non-negative. Since $\lambda > 0$, we deduce that $-n\lambda$ cannot be an eigenvalue of $X^\top X$ and we can conclude. \square

Recap

- ▶ **What happens when we invoke** `sklearn.linear_model.Ridge` with default settings?
- ▶ $\alpha = 1 \rightarrow \lambda = 1/n$ with our notation, barely any regularization if n large
- ▶ `fit_intercept` is `True` \rightarrow does not consider centered data (so our analysis is not entirely accurate)
- ▶ `normalize` is `False` \rightarrow we decide whether we normalize our data
- ▶ `solver` is `auto` \rightarrow `sklearn` will decide how to solve the minimization problem depending on the size of the data: **the solution could be not exact!**
- ▶ `tol` = 0.001 \rightarrow tolerance threshold on the residuals

1.4. Polynomial regression

Polynomial regression

- ▶ linear regression is a powerful tool, especially because we can transform the inputs in a non-linear fashion
- ▶ **Example:** polynomial regression in \mathbb{R}
- ▶ inputs $x_1, \dots, x_n \in \mathbb{R}$
- ▶ define the mapping $\phi(x) = (1, x, x^2, \dots, x^p)^\top$
- ▶ then


$$\langle w, \phi(x) \rangle = w_0 + w_1x + w_2x^2 + \dots + w_px^p,$$

and we can find the best coefficients by linear regression

- ▶ `numpy.polyfit` → very handy when we want to fit univariate data

Polynomial regression, ctd.


- ▶ **Example:** data = degree three polynomial + Gaussian noise with small variance
- ▶ fit a degree 3 polynomial:



- ▶ **Remark:** in practice, we do not know the degree of the polynomial!

Polynomial regression, ctd.

- ▶ typical case of under / overfitting:
 - ▶ when degree too low, poor fit
 - ▶ when degree too high, wiggly function ($n + 1 \Rightarrow$ interpolation)



1.5. Logistic regression

Logistic regression

- ▶ classification with $\mathcal{Y} = \{0, 1\}$
- ▶ however, we predict **the probability of belonging to class 1**
- ▶ hypothesis class:

$$\mathcal{H} = \{x \mapsto \phi(\langle w, x \rangle), w \in \mathbb{R}^d\},$$

with ϕ the *logistic function* (aka *sigmoid function*)

$$\phi(z) = \frac{1}{1 + e^{-z}}.$$

- ▶ **Intuition:** squeeze the score between 0 and 1 to transform it into a probability
- ▶ $\mathbb{P}(y = 1 | x) = \phi(w^\top x)$ and $\mathbb{P}(y = 0 | x) = 1 - \phi(w^\top x)$

Logistic function





Figure: the logistic function $\phi : t \mapsto 1/(1 + e^{-t})$.

Logistic loss

- ▶ **Next:** we need to define a loss function
- ▶ for any y, y' , we define the *logistic loss*:

$$\ell(y, y') = -(1 - y) \log(1 - y') - y \log y'.$$



Logistic regression


- ▶ finally, logistic regression = empirical risk minimization with the logistic loss
- ▶ that is, minimize for $w \in \mathbb{R}^d$

$$\hat{\mathcal{R}}_S(w) = \sum_{i=1}^n \{ -(1 - y_i) \log(1 - \phi(w^\top x_i)) - y_i \log \phi(w^\top x_i) \} .$$

- ▶ **Remark (i):** we can show that this is equivalent to maximum likelihood for a certain prior distribution
- ▶ **Remark (ii):** complicated to optimize (see exercise)


Logistic regression in dimension 1

► **Example:** in dimension one:



Logistic regression in dimension 2

- **Example:** in dimension two:



Exercise

Exercise: Recall that we defined the logistic loss by

$$\ell(y, y') = -(1 - y) \log(1 - y') - y \log y'.$$

1. Show that ERM with the logistic loss is equivalent to minimizing

$$F(w) = \sum_{i=1}^n \log(1 + \exp(-\tilde{y}_i \langle w, x_i \rangle)),$$

where $\tilde{y}_i = \text{sign}(y_i - 0.5)$. Deduce that $\hat{\mathcal{R}}$ is a convex function of w .

2. Compute the gradient of $\hat{\mathcal{R}}$ with respect to w . *Hint:* show that $\phi'(z) = \phi(z)(1 - \phi(z))$.
3. Can you solve $\nabla \hat{\mathcal{R}}(w) = 0$? If not, propose a strategy for finding a good w .

Correction of the exercise

1. Let us set $1 \leq i \leq n$. We write

$$\begin{aligned}\ell(y_i, \phi(w^\top x_i)) &= -(1 - y_i) \log(1 - \phi(w^\top x_i)) - y_i \log \phi(w^\top x_i) \\ &= -(1 - y_i) \log \frac{e^{-w^\top x_i}}{1 + e^{-w^\top x_i}} - y_i \log \frac{1}{1 + e^{-w^\top x_i}} \\ &= -(1 - y_i) \log e^{-w^\top x_i} + \log(1 + e^{-w^\top x_i}).\end{aligned}$$

If $y_i = 0$, the last display equals

$$\log(1 + \exp(w^\top x_i)),$$

if $y_i = 1$, it is

$$\log(1 + \exp(-w^\top x_i)).$$

One can check directly that $x \mapsto \log(1 + e^{-x})$ is convex. By composition, F is a sum of convex functions, thus convex.

Correction of the exercise, ctd.

2. Let $1 \leq j \leq d$. We write

$$\begin{aligned}\frac{\partial \hat{\mathcal{R}}(w)}{\partial w_j} &= - \sum_{i=1}^n \frac{\partial}{\partial w_j} \{ (1 - y_i) \log(1 - \phi(w^\top x_i)) + y_i \log \phi(w^\top x_i) \} \\ &= - \sum_{i=1}^n \left\{ \frac{-(1 - y_i)}{1 - \phi(w^\top x_i)} + \frac{y_i}{\phi(w^\top x_i)} \right\} \frac{\partial}{\partial w_j} \phi(w^\top x_i) \\ &= - \sum_{i=1}^n \left\{ \frac{-(1 - y_i)}{1 - \phi(w^\top x_i)} + \frac{y_i}{\phi(w^\top x_i)} \right\} \phi(w^\top x_i) (1 - \phi(w^\top x_i)) x_{i,j} \\ \frac{\partial \hat{\mathcal{R}}(w)}{\partial w_j} &= - \sum_{i=1}^n (y_i - \phi(w^\top x_i)) x_{i,j}.\end{aligned}$$

3. It does not seem possible to solve $\nabla F(w) = 0$ in closed-form, one has to use gradient descent. □

Recap

- ▶ **What happens when we call `sklearn.linear_model.LogisticRegression`?**
- ▶ penalty is $\ell_2 \rightarrow$ **there is regularization by default!** (not much though, $C = 1$)
- ▶ `fit_intercept` is `True` \rightarrow again, our maths are not entirely accurate
- ▶ `solver` is `liblinear` \rightarrow since there is no closed-form, a solver will be used
- ▶ `liblinear` uses coordinate descent
- ▶ will default soon to `lbfgs` (limited memory Broyden-Fletcher- -Goldfarb-Shanno)
- ▶ **do not worry too much about the solvers**, just change if you see that it is not converging

1.6. Support vector machines

Support vector machines


- ▶ classification with $x_1, \dots, x_n \in \mathbb{R}^d$ and $y_i \in \{-1, 1\}$
- ▶ **Recall:** linearly separable means that there exist (w, b) such that

$$\forall i \in [n], \quad y_i(w^\top x_i + b) > 0.$$

- ▶ **Remark:** all halfspaces satisfying this condition are empirical risk minimizers
- ▶ **Question:** which one should we pick?

Some intuition

- **Idea:** choose the one with maximum *margin*



- **Intuitively,** we would prefer the **red** line instead of the **blue** one

Margins

Definition: The margin of a hyperplane with respect to a training set is defined as the minimal distance between a point in the training set and the hyperplane.

- ▶ *Hard-SVM*⁴ = minimizing empirical risk and choosing the max margin hyperplane
- ▶ **Question:** how to put this in equation?
- ▶ first, we need to express the distance between a point and a hyperplane:

Lemma: Assume that $\|w\| = 1$. Then the distance between x and the hyperplane defined by (w, b) is given by $|w^\top x + b|$.

⁴Boser, Guyon, Vapnik, *A training algorithm for optimal margin classifiers*, 5th workshop on computational learning theory, 1992

Proof of the lemma

- ▶ we want to compute

$$\min\{\|x - v\| \quad \text{s.t.} \quad w^\top v + b = 0\}.$$

- ▶ take $v = x - (w^\top x + b)w$:

$$w^\top v + b = w^\top x - (w^\top x + b) \|w\|^2 + b = 0,$$

since $\|w\| = 1$.

- ▶ moreover,

$$\|x - v\| = |w^\top x + b| \|w\| = |w^\top x + b|.$$

- ▶ for now, we have a point v on the hyperplane with distance $|w^\top x + b|$
- ▶ let us show that any other point has a larger distance

Proof of the lemma, ctd.

- ▶ let u such that $w^\top u + b = 0$, then

$$\begin{aligned}\|x - u\|^2 &= \|x - v + v - u\|^2 \\ &= \|x - v\|^2 + \|v - u\|^2 + 2(x - v)^\top (v - u) \\ &\geq \|x - v\|^2 + 2(x - v)^\top (v - u) \\ &= \|x - v\|^2 + 2(w^\top x + b)w^\top (v - u)\end{aligned}$$

- ▶ notice that $w^\top v = w^\top u = -b$, therefore

$$\|x - u\|^2 \geq \|x - v\|^2 .$$



Hard-SVM rule

- ▶ **Consequence of the lemma:** the closest point in the training set has distance $\min_i |w^\top x_i + b|$ to the hyperplane
- ▶ we can rewrite the hard-SVM rule as

$$(\hat{w}, \hat{b}) \in \arg \max_{(w,b), \|w\|=1} \min_i |w^\top x_i + b| \quad \text{s.t.} \quad y_i(w^\top x_i + b) > 0 \quad \forall i.$$

- ▶ **Intuition:** x_i on the right side of the hyperplane if y_i and $w^\top x_i + b$ have the same sign
- ▶ in the separable case, it is possible to show that an equivalent formulation is

$$(\hat{w}, \hat{b}) \in \arg \max_{(w,b), \|w\|=1} \min_i y_i(w^\top x_i + b).$$

Hard-SVM as quadratic programming

- ▶ as in the first linear example, possible to reframe as a standard optimization problem

Lemma: Let (w_0, b_0) be the solution of the following QP:

$$(w_0, b_0) \in \arg \min_{(w, b)} \|w\|^2 \quad \text{s.t.} \quad y_i(w^\top x_i + b) \geq 1 \quad \forall i.$$

Then $\hat{w} = w_0 / \|w_0\|$ and $\hat{b} = b_0 / \|w_0\|$ satisfy the Hard-SVM rule.

- ▶ QP = quadratic programming: objective is a quadratic function and the constraints are linear inequalities

Proof of the lemma

- ▶ let (w, b) be a solution of the Hard-SVM rule
- ▶ define the achieved margin by

$$\gamma = \min_i y_i(w^\top x_i + b).$$

- ▶ by definition, for all $1 \leq i \leq n$, we have

$$y_i(w^\top x_i + b) \geq \gamma,$$

that is

$$y_i \left(\left(\frac{w}{\gamma} \right)^\top x_i + \frac{b}{\gamma} \right) \geq 1.$$

- ▶ thus $(w/\gamma, b/\gamma)$ satisfies the condition of the QP

Proof of the lemma, ctd.

- ▶ in particular,

$$\|w_0\| \leq \left\| \frac{w}{\gamma} \right\| = \frac{1}{\gamma}.$$

- ▶ as a consequence, for all $1 \leq i \leq n$,

$$y_i(\hat{w}^\top x_i + \hat{b}) = \frac{1}{\|w_0\|} \cdot y_i(w_0^\top x_i + b_0) \geq \frac{1}{\|w_0\|} \geq \gamma.$$

- ▶ since $\|\hat{w}\| = 1$, we have shown that (\hat{w}, \hat{b}) is a solution of the Hard-SVM rule



The homogeneous case

- ▶ if we set $b = 0$, Hard-SVM rule becomes

$$\text{minimize}_w \|w\|^2 \quad \text{s.t.} \quad \forall 1 \leq i \leq n, \quad y_i w^\top x_i \geq 1.$$

- ▶ in that case, the solution w_0 is *supported* by the examples exactly at distance $1/\|w_0\|$ from the hyperplane

Theorem (Fritz John): Let $I := \{i \text{ s.t. } |w_0^\top x_i| = 1\}$. Then there exists coefficients α_i , $i \in I$, such that

$$w_0 = \sum_{i \in I} \alpha_i x_i.$$

- ▶ the $\{x_i, i \in I\}$ are called *support vectors*, hence the name

Soft-SVM

- ▶ linearly separable assumption is quite restrictive
- ▶ condition in the QP:

$$\forall 1 \leq i \leq n, \quad y_i(w^\top x_i + b) \geq 1.$$

- ▶ **Natural relaxation:** allow this constraint to be violated for some points in the dataset
- ▶ we introduce *slack variables* ξ_1, \dots, ξ_n and replace the constraint by

$$\forall 1 \leq i \leq n, \quad y_i(w^\top x_i + b) \geq 1 - \xi_i.$$

- ▶ **Intuition:** the ξ_i encode by how much the constraint is violated

Soft-SVM, ctd.

- ▶ **Key idea:** minimize jointly $\|w\|$ and the average of the ξ_i
- ▶ namely, the Soft-SVM rule is

$$\text{minimize}_{w,b,\xi} \left\{ \lambda \|w\|^2 + \frac{1}{n} \sum_{i=1}^n \xi_i \right\} \quad \text{s.t.} \quad \forall 1 \leq i \leq n, \quad y_i(w^\top x_i + b) \geq 1 - \xi_i \text{ and } \xi_i \geq 0,$$

where $\lambda > 0$ is a fixed hyperparameter

Soft-SVM as regularized ERM

- recall the definition of the *hinge loss*:

$$\ell(y, y') = \max\{0, 1 - yy'\}.$$

- we have the remarkable result:

Lemma: Consider

$$(\hat{w}, \hat{b}) \in \arg \min_{(w, b)} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(y_i, w^\top x_i + b) + \lambda \|w\|^2 \right\}.$$

Then (\hat{w}, \hat{b}) satisfies the Soft-SVM rule.

Proof of the lemma

- recall that we want to minimize

$$\lambda \|w\|^2 + \frac{1}{n} \sum_{i=1}^n \xi_i$$

subject to

$$\forall 1 \leq i \leq n, \quad y_i(w^\top x_i + b) \geq 1 - \xi_i \text{ and } \xi_i \geq 0.$$

- fix w, b and let us try and minimize in ξ
- let us fix some i : since $\xi_i \geq 0$, the best assignment would be $\xi_i = 0$ if $y_i(w^\top x_i + b) \geq 1$, and $1 - y_i(w^\top x_i + b)$ otherwise
- in other words,

$$\forall 1 \leq i \leq n, \quad \xi_i = \max\{0, 1 - y_i(w^\top x_i + b)\} = \ell(y_i, w^\top x_i + b).$$



Recap

- ▶ scikit-learn implementation: `sklearn.svm.LinearSVC`
- ▶ then a solver is called, similarly to the logistic regression case
- ▶ default is ℓ_2 regularization (as we have seen)
- ▶ regularization is given by $C = 1$ ($C = 1/\lambda$, beware!)
- ▶ we will see an extension of SVM when we look into kernel methods

2. Kernel methods

2.1. Positive semi-definite kernels

Representation of the data

- ▶ **What we have seen so far:** linear classification / linear regression
- ▶ works well if the data is linearly separable
- ▶ **Problem:** that is not always the case!
- ▶ what if we could transport the data to another space where it is well-behaved?
- ▶ for instance a very high-dimensional space
- ▶ first we define a *kernel*

Positive semi-definite kernels

Definition: a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is called a positive definite kernel if $k(x, x') = k(x', x)$ for any $x, x' \in \mathcal{X}$, and

$$\forall x_1, \dots, x_n \in \mathcal{X}, \forall c_1, \dots, c_n \in \mathbb{R}, \quad \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) \geq 0.$$

- ▶ unlike the name suggests, k has no reason to be *positive*
- ▶ in other words, the Gram matrix $K = (k(x_i, x_j))_{i,j=1}^n$ is positive definite for any input data x_1, \dots, x_n
- ▶ *kernel methods* take this K as input
- ▶ **Remark:** this is *costly*, $\mathcal{O}(n^2)$ whatever we do, with possible dependency in the dimensionality of the data

Fundamental example

- ▶ suppose that $\mathcal{X} = \mathbb{R}$
- ▶ then $k(x, y) := xy$ is a positive definite kernel
- ▶ **Why?** first, we check that $k(x, y) = k(y, x)$
- ▶ second, let $n \geq 1$, $x_1, \dots, x_n \in \mathbb{R}^d$, and $c_1, \dots, c_n \in \mathbb{R}$, then

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j x_i x_j \\ &= \left(\sum_{i=1}^n c_i x_i \right)^2 \\ &\geq 0. \end{aligned}$$

Fundamental example, ctd.

- ▶ we can extend this example:
- ▶ suppose that $\mathcal{X} = \mathbb{R}^d$
- ▶ let $n \geq 1$, $x_1, \dots, x_n \in \mathbb{R}^d$, and $c_1, \dots, c_n \in \mathbb{R}$, then

$$\begin{aligned}\sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j x_i^\top x_j \\ &= \left\| \sum_{i=1}^n c_i x_i \right\|^2 \\ &\geq 0.\end{aligned}$$

- ▶ $k(x, y) := x^\top y$ is usually called the **linear kernel**
- ▶ **Intuition:** kernels are a generalization of inner product

Other examples

- **Polynomial kernel:**

$$\mathcal{X} = \mathbb{R}^d, \quad k(x, y) = (x^\top y + c)^k.$$

- **min kernel:**

$$\mathcal{X} = \mathbb{R}, \quad k(x, y) = \min(x, y).$$

- **Gaussian kernel:**

$$\mathcal{X} = \mathbb{R}^d, \quad k(x, y) = \exp\left(\frac{-\|x - y\|^2}{2\nu^2}\right).$$

- **Exponential kernel:**

$$\mathcal{X} = \mathbb{R}^d, \quad k(x, y) = \exp\left(\frac{-\|x - y\|}{2\nu}\right).$$

- ...

Choosing the bandwidth

- ▶ Gaussian and Laplace kernel: one has to choose the bandwidth parameter ν
- ▶ indeed, if ν is *too large* with respect to the typical value of $\|x_i - x_j\|$, then $K \approx I_n$
- ▶ in the other direction, if ν is *too small*, then $K \approx \mathbf{1}\mathbf{1}^\top$
- ▶ both cases are degenerate: whatever we do with K is not going to work very well
- ▶ one possible solution: **median heuristic**⁵

$$\nu = \text{Med}\{\|x_i - x_j\|, \quad 1 \leq i, j \leq n\}.$$

- ▶ preferable to the mean (too sensitive to extreme values)
- ▶ we can pick other quantiles

⁵Garreau, Jitkrittum, Kanagawa, *Large sample analysis of the median heuristic*, 2017

Exercise

Exercise: Show that the following functions are positive definite kernels:

1. $\mathcal{X} = \mathbb{N}$, $k(x, y) = 2^{x+y}$
2. $\mathcal{X} = \mathbb{R}$, $k(x, y) = \cos(x - y)$
3. $\mathcal{X} = \mathbb{R}^d$, $k(x, y) = (x^\top y)^2$

Correction of the exercise

1. let $x_1, \dots, x_n \in \mathbb{N}$ and $c_1, \dots, c_n \in \mathbb{R}$, then

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j 2^{x_i + x_j} = \left(\sum_{i=1}^n c_i 2^{x_i} \right)^2 \geq 0.$$

2. let $x_1, \dots, x_n \in \mathbb{N}$ and $c_1, \dots, c_n \in \mathbb{R}$, then

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n c_i c_j \cos(x_i - x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j (\cos x_i \cos x_j + \sin x_i \sin x_j) \\ &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j \cos x_i \cos x_j + \sum_{i=1}^n \sum_{j=1}^n c_i c_j \sin x_i \sin x_j \\ &= \left(\sum_{i=1}^n c_i \cos x_i \right)^2 + \left(\sum_{i=1}^n c_i \sin x_i \right)^2. \end{aligned}$$

Correction of the exercise, ctd.

3. let $x, y \in \mathbb{R}^d$:

$$\begin{aligned} (x^\top y)^2 &= \text{trace} (x^\top y x^\top y) \\ &= \text{trace} (y^\top x x^\top y) \\ &= \text{trace} (x x^\top y y^\top) . \end{aligned}$$

We recognize the inner product between matrices.

Important remark

- ▶ recall the linear kernel: as in the exercise, all we used were properties of inner product
- ▶ let $\Phi : \mathcal{X} \rightarrow \mathcal{H}$ be some mapping, \mathcal{H} a Hilbert space with scalar product $\langle \cdot, \cdot \rangle$
- ▶ then $k(x, y) = \langle \Phi(x), \Phi(y) \rangle$ is positive definite:

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j \langle \Phi(x_i), \Phi(x_j) \rangle = \left\| \sum_{i=1}^n c_i \Phi(x_i) \right\|^2 \geq 0,$$

by linearity of the inner product.

Kernel as inner products

- ▶ **Remarkable fact:** the converse statement is true!

Theorem (Aronszajn, 1950): For any kernel k on \mathcal{X} , there exists a Hilbert space $(\mathcal{H}, \langle \cdot, \cdot \rangle)$ and a mapping $\phi : \mathcal{X} \rightarrow \mathcal{H}$ such that

$$\forall x, y \in \mathcal{X}, \quad k(x, y) = \langle \Phi(x), \Phi(y) \rangle .$$

- ▶ **Reminder:** Hilbert space = inner product + *complete* for the associated norm (Cauchy sequences converge in \mathcal{H})
- ▶ not so important for us, needed for good convergence properties
- ▶ **Consequence:** we can think of any kernel as a dot product in the *feature space*

Proof in the finite case

- ▶ assume that $\mathcal{X} = \{x_1, \dots, x_N\}$ is finite of size N
- ▶ any kernel k is entirely defined by the $N \times N$ positive semi-definite matrix $K := (k(x_i, x_j))_{i,j=1}^N$
- ▶ we can diagonalize K in an orthonormal basis (u_1, \dots, u_N) with associated (non-negative) eigenvalues $\lambda_1, \dots, \lambda_N$
- ▶ then we write

$$\begin{aligned} k(x_i, x_j) &= \left(\sum_{\ell=1}^N \lambda_{\ell} u_{\ell} u_{\ell}^{\top} \right)_{i,j} \\ &= \sum_{\ell=1}^N \lambda_{\ell} (u_{\ell})_i (u_{\ell})_j = \langle \Phi(x_i), \Phi(x_j) \rangle, \end{aligned}$$

with

$$\Phi(x_i) := \left(\sqrt{\lambda_1} (u_1)_i \cdots \sqrt{\lambda_n} (u_N)_i \right)^{\top}.$$

2.2. Reproducing kernel Hilbert spaces

Function spaces

- ▶ among all spaces in the previous statement, one of them has interesting properties
- ▶ in particular, it is a **space of functions**
- ▶ *i.e.*, we can map each point $x \in \mathcal{X}$ to a *function* $\Phi(x) = k_x \in \mathcal{H}$
- ▶ **Example:** $\mathcal{X} = \mathbb{R}$, we map each x to the function $t \mapsto xt$
- ▶ \rightarrow space of linear functions
- ▶ more complicated in general...

RKHS

Definition: let \mathcal{X} be a set and \mathcal{H} be a function space forming a Hilbert space with inner product $\langle \cdot, \cdot \rangle$. The function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is called a *reproducing kernel* of \mathcal{H} if

- ▶ \mathcal{H} contains all functions of the form $k_x : t \mapsto k(x, t)$
- ▶ for every $x \in \mathcal{X}$ and $f \in \mathcal{H}$, the *reproducing property* holds:

$$f(x) = \langle f, k_x \rangle.$$

- ▶ if a reproducing kernel exists, then \mathcal{H} is called a *reproducing kernel Hilbert space* (RKHS)

Equivalent definition

Theorem: the Hilbert space $\mathcal{H} \subseteq \mathbb{R}^{\mathcal{X}}$ is a RKHS if, and only if, for any $x \in \mathcal{X}$, the mapping $f \mapsto f(x)$ is continuous.

► *Proof:* \Rightarrow if a reproducing kernel k exists, then

$$|f(x)| = |\langle f, k_x \rangle| \leq \|f\| \cdot \|k_x\|$$

by Cauchy-Schwarz.

- we see that $\|k_x\|^2 = \langle k_x, k_x \rangle = k(x, x)$, thus $|f(x)| \leq \|f\| \cdot k(x, x)^{1/2}$
- thus $f \mapsto f(x)$ is continuous.
- \Leftarrow Riesz theorem.



Important properties

Theorem (uniqueness): if \mathcal{H} is a RKHS, then it has a unique reproducing kernel. Conversely, a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ can be the reproducing kernel of at most one RKHS.

- ▶ this is why we talk about *the* RKHS associated to k

Theorem: a function $k : \mathcal{X} \times \mathcal{X}$ is positive definite if, and only if, it is a reproducing kernel.

- ▶ showing that a kernel is positive definite is enough to get Φ and \mathcal{H} with the reproducing property “for free”

Example

- ▶ **Example:** polynomial kernel of degree 2:

$$k(x, y) = (x^\top y)^2.$$

- ▶ according to the exercise,

$$k(x, y) = \langle xx^\top, yy^\top \rangle_F,$$

we have proven that k is positive definite

- ▶ **Question:** what is the RKHS?
- ▶ we know that \mathcal{H} contains all the functions

$$f(x) = \sum_i a_i k(x_i, x) = \sum_i a_i \langle x_i x_i^\top, x x^\top \rangle = \langle \sum_i a_i x_i x_i^\top, x x^\top \rangle$$

Example, ctd.

- ▶ spectral theorem: any symmetric matrix can be decomposed as $\sum_i a_i x_i x_i^\top$
- ▶ candidate RKHS: set a quadratic functions

$$f_S(x) = \langle S, xx^\top \rangle = x^\top S x,$$

with S symmetric matrix of size $d \times d$

- ▶ inner product on \mathcal{H} :

$$\langle f_S, f_{S'} \rangle = \langle S, S' \rangle_F.$$

- ▶ we can check that \mathcal{H} is a Hilbert space (isomorphic to $\mathcal{S}^{d \times d}$)
- ▶ finally, we check the reproducing property

2.3. More examples

Elementary properties

Proposition: Let $k_i : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ be a (potentially infinite) family of p.d. kernels. Then

- ▶ for any $\lambda_1, \dots, \lambda_p \geq 0$, the sum $\sum_{i=1}^p \lambda_i k_i$ is positive definite
- ▶ for any $a_1, \dots, a_p \in \mathbb{N}$, the product $k_1^{a_1} \cdots k_p^{a_p}$ is positive definite
- ▶ if it exists, the limit $k = \lim_{p \rightarrow +\infty} k_p$ is positive definite

Moreover, let \mathcal{X}_i be a sequence of sets and k_i positive kernels on each \mathcal{X}_i . Then

$$k((x_1, \dots, x_p), (y_1, \dots, y_p)) := \prod_{i=1}^p k_i(x_i, y_i)$$

and

$$k((x_1, \dots, x_p), (y_1, \dots, y_p)) := \sum_{i=1}^p k_i(x_i, y_i)$$

are positive definite kernels.

Taking the exponential

Theorem: if k is a positive definite kernel, then e^k as well.

► *Proof:* we write

$$e^{k(x,y)} = \lim_{n \rightarrow +\infty} \sum_{p=0}^n \frac{k(x,y)^p}{p!},$$

then reason step by step.

- by the product property, $k(x,y)^p$ is a kernel for any $p \geq 0$
- as a positive linear combination of kernels, $\sum_{p=0}^n \frac{k(x,y)^p}{p!}$ is a kernel for all $n \geq 1$
- finally, e^k is a kernel as a limit of kernels. □

Exercise

Exercise: show that the Gaussian kernel k is positive definite, that is,

$$k(x, y) := \exp \left(\frac{-\|x - y\|^2}{2\nu^2} \right),$$

where $\nu > 0$ is a fixed *bandwidth* parameter. *Hint:* decompose the squared norm and use the properties.

Correction of the exercise

- ▶ first recall that

$$\|x - y\|^2 = \|x\|^2 - 2x^\top y + \|y\|^2 .$$

- ▶ we split the kernel in two parts:

$$k(x, y) = e^{-\|x\|^2 - \|y\|^2} \cdot \exp(2x^\top y) .$$

- ▶ the first part is a kernel (scalar product of two feature maps)
- ▶ the second part as well (exponential of a kernel)
- ▶ since the product of kernels is a kernel, we can conclude.

2.4. The kernel trick and applications

The kernel trick

- ▶ input data $x_1, \dots, x_n \in \mathcal{X}$
- ▶ $k : \mathcal{X} \times \mathcal{X}$ kernel with associated RKHS \mathcal{H}
- ▶ we call $\Phi : \mathcal{X} \rightarrow \mathcal{H}$ the feature map
- ▶ **Idea:** imagine that our algorithm only depends on scalar products $x_i^\top x_j$
- ▶ then we can map the x_i to \mathcal{H} and replace the inner products by kernel evaluations, since

$$\langle \Phi(x_i), \Phi(x_j) \rangle = k(x_i, x_j).$$

- ▶ simple “trick” with many, many applications

Example

- ▶ **Example:** computing distances
- ▶ suppose that our algo relies on distance computation
- ▶ that is, $\|x - y\|^2$
- ▶ we can write

$$\begin{aligned}\|\Phi(x) - \Phi(y)\|^2 &= \langle \Phi(x) - \Phi(y), \Phi(x) - \Phi(y) \rangle \\ &= \langle \Phi(x), \Phi(x) \rangle - 2\langle \Phi(x), \Phi(y) \rangle + \langle \Phi(y), \Phi(y) \rangle \\ \|\Phi(x) - \Phi(y)\|^2 &= k(x, x) - 2k(x, y) + k(y, y).\end{aligned}$$

- ▶ in other words,

$$d_{\mathcal{H}}(x, y) = \sqrt{k(x, x) - 2k(x, y) + k(y, y)}.$$

- ▶ as promised, **we do not need to know Φ !**

Exercise

Exercise: let $S = \{x_1, \dots, x_n\}$ be a finite set of points in \mathcal{X} . Compute the distance to the barycenter of S in the RKHS.

Correction

- ▶ first we see that the barycenter is

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i .$$

- ▶ **Beware:** there is not necessarily a point $m \in \mathcal{X}$ such that $\Phi(m) = \mu$
- ▶ then we proceed as before:

$$\begin{aligned} d_{\mathcal{H}}(x, S)^2 &= \|x - \mu\|^2 \\ &= k(x, x) - \frac{2}{n} \sum_{i=1}^n k(x, x_i) + \frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n k(x_i, x_j) . \end{aligned}$$

2.5. The representer theorem

Motivation

- ▶ let us imagine that we take \mathcal{H} as hypothesis class
- ▶ starting from regularized ERM, our optimization problem will look like

$$\arg \min_{f \in \mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)) + \lambda \|f\|^2 \right\}. \quad (\star)$$

- ▶ we penalize by the norm because it is an indicator of the *smoothness* of f
- ▶ **Why?** Cauchy-Schwarz + exercise:

$$|f(x) - f(y)| = |\langle f, k_x - k_y \rangle| \leq \|f\| \cdot \|k_x - k_y\| = \|f\| \cdot d_{\mathcal{H}}(x, y).$$

- ▶ Eq. (\star) is a complicate problem, potentially *infinite-dimensional*
- ▶ **Question:** how to solve it in practice?

The representer theorem

Theorem: let \mathcal{H} be the RKHS associated to k defined on \mathcal{X} . Let $S = \{x_1, \dots, x_n\} \subseteq \mathcal{X}$ be a finite set of points. Let $\Psi : \mathbb{R}^{n+1} \rightarrow \mathbb{R}$ be a function, increasing in the last variable. Then any solution to the minimization problem

$$\arg \min_{f \in \mathcal{H}} \Psi(f(x_1), \dots, f(x_n), \|f\|)$$

admits a representation of the form

$$\forall x \in \mathcal{X}, \quad f(x) = \sum_{i=1}^n \alpha_i k(x_i, x).$$

► **Main consequence:** Eq. (\star) is actually a finite-dimensional problem (!)

Practical use

- ▶ recall that we defined $K := (k(x_i, x_j))_{i,j=1}^n$
- ▶ before turning to concrete examples, we notice that we can simply express the key quantities
- ▶ for instance, for any $1 \leq j \leq n$,

$$f(x_j) = \sum_{i=1}^n \alpha_i k(x_i, x_j) = (K\alpha)_j.$$

- ▶ in the same way,

$$\|f\|^2 = \left\| \sum_{i=1}^n \alpha_i k(x_i, \cdot) \right\|^2 = \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j k(x_i, x_j) = \alpha^\top K \alpha.$$

2.6. Kernel ridge regression

Kernel Ridge Regression⁶ (KRR)

- ▶ regression setting: $(x_i, y_i) \in \mathcal{X} \times \mathcal{Y}$
- ▶ $\mathcal{Y} \subseteq \mathbb{R}$, but \mathcal{X} could be anything
- ▶ we have a kernel k on \mathcal{X}
- ▶ same idea than with ridge regression:

$$\hat{f} \in \arg \min_{f \in \mathcal{H}} \left\{ \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \|f\|^2 \right\}.$$

- ▶ here effect of the regularization is to make \hat{f} smoother

⁶Cristianini and Shawe-Taylor, *An introduction to support vector machines and other kernel-based learning methods*, Cambridge University Press, 2000

Solving KRR

- ▶ representer theorem \Rightarrow

$$\hat{f}(x) = \sum_{i=1}^n \alpha_i k(x_i, x),$$

for some $\alpha \in \mathbb{R}^n$

- ▶ as per the previous remark, we know that

$$(\hat{f}(x_1), \dots, \hat{f}(x_n))^{\top} = K\alpha,$$

and

$$\|\hat{f}\|^2 = \alpha^{\top} K \alpha.$$

- ▶ thus KRR can be re-written as

$$\hat{\alpha} \in \arg \min_{\alpha \in \mathbb{R}} \left\{ \frac{1}{n} (y - K\alpha)^{\top} (y - K\alpha) + \lambda \alpha^{\top} K \alpha \right\}.$$

Solving KRR, ctd.

- ▶ convex, smooth objective \Rightarrow set the gradient to zero
- ▶ $\hat{\alpha}$ has to be solution of

$$0 = \frac{-2}{n} K(y - K\alpha) + 2\lambda K\alpha = \frac{2}{n} K [(K + n\lambda I_n)\alpha - y]$$

- ▶ since $\lambda > 0$, $K + n\lambda I_n$ is invertible
- ▶ the solution is given by

$$\hat{\alpha} = (K + n\lambda I_n)^{-1} y.$$

- ▶ **Question:** what do you think of the uniqueness of the solution?

Exercise

Exercise: show that, when the kernel is linear, the solution given by the kernel ridge regression approach and the vanilla ridge regression coincide.

Hint: show that for any matrices A, B ,

$$A(BA + \gamma I)^{-1} = (AB + \gamma I)^{-1}A,$$

whenever the inverses are well-defined.

Correction

- ▶ for the linear kernel, $K = X^\top X$
- ▶ proof of the hint: assume that the inverse exist
- ▶ then write

$$A(I + BA) = (I + AB)A.$$

- ▶ finally, (right) multiply by $(I + BA)^{-1}$ and (left) multiply by $(I + AB)^{-1}$ both sides

3. Tree-based classifiers

3.1. Partition rules

Introduction

- ▶ let $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \mathbb{R}$
- ▶ in this section, we consider partition-based classifiers:

$$\mathcal{H} = \left\{ h : x \mapsto \sum_{j=1}^p h_j \mathbb{1}_{x \in A_j} \right\},$$

where $a_j \in \mathbb{R}$ and A_1, \dots, A_p form a *partition* of the space

- ▶ that is,

$$A_1 \cup \dots \cup A_p = \mathcal{X} \quad \text{and} \quad A_i \cap A_j = \emptyset \forall i \neq j.$$

- ▶ the A_j s are often called *cells*
- ▶ generally, for practical reasons **the A_j s are rectangles**

ERM for partition rules

- ▶ assume that the partition is fixed and set $A(x)$ = cell containing x
- ▶ **Regression:** with squared loss, ERM rule gives

$$f(x) = \frac{1}{|\{j \text{ s.t. } x_j \in A(x)\}|} \sum_{i=1}^n x_i \mathbb{1}_{y_i \in A(x)},$$


that is the average of the observations on each cell

- ▶ **Classification:** majority vote:

$$f(x) = \begin{cases} 1 & \text{if } |\{i \text{ s.t. } x_i \in A(x) \text{ and } y_i = 1\}| \geq \\ & |\{i \text{ s.t. } x_i \in A(x) \text{ and } y_i = 0\}| \\ 0 & \text{otherwise.} \end{cases}$$


- ▶ thus **ERM** \Leftrightarrow finding the best partition (for a fixed p)
- ▶ **Problem:** this is computationally very hard! p^n possibilities to compare
- ▶ even if we restrict ourselves to rectangles, intractable

Example of a partition-based predictor




Tree structures

- ▶ one possible solution: start from \mathcal{X} and *split* the cells iteratively
- ▶ we obtain a tree-like structure
- ▶ **Remark:** not necessary to split in two, but very common
- ▶ another advantage in doing so: root the new data efficiently



Not a tree



► **Figure:** this partition of $[0, 1]^2$ can not be obtained by recursive binary splitting

Growing trees

- ▶ **Question:** how do we make the splits?
- ▶ **general answer:** take an heuristic that makes sense
- ▶ each heuristic yields a different algorithm, completed with stopping criterion (do a split only if gain greater than γ)
- ▶ complete reference on such procedures: the *yellow book*⁷
- ▶ good splitting rules:
 - ▶ create many cells (enough to capture the local variations of the distribution);
 - ▶ create cells that are large enough (we want sufficiently training data in the cells to compute a relevant average)
- ▶ **Notation:** I current node, I_L (resp. I_R) left (resp. right) node after the split

⁷Devroye, Györfi, Lugosi, *A probabilistic theory of pattern recognition*, 1996

ID3⁹ and C4.5

Definition: Let S be a finite set of points. Then we define the *entropy* of S by

$$H(S) = \sum_{y \in \mathcal{Y}} -p(y) \log_2 p(y),$$

where $p(y)$ is the proportion of elements of S classified as y .

- ▶ easy to see that $H(S) = 0$ means that the node is *pure* = only one label ($0 \log 0 = 0$)
- ▶ C4.5 criterion:⁸ find direction and split that maximizes

$$H(I) - H(I_L) - H(I_R).$$

⁸Quinlan, *C4.5: Programs for Machine Learning*, 1993

⁹Quinlan, *Induction of decision trees*, Machine Learning, 1986

CART trees, classification

- ▶ later supplanted by CART trees¹⁰

Definition: Let S be a finite set of points. We define the *Gini impurity* by

$$G(S) = \sum_{y \in \mathcal{Y}} p(y)(1 - p(y)).$$

- ▶ $G(S) = 0$ if the leaf is pure
- ▶ CART trees: find direction and split that maximizes

$$G(I) - G(I_L) - G(I_R).$$

¹⁰Breiman et al., *Classification and Regression Trees*, 1984

CART trees for regression

- ▶ slightly different in the regression setting: look at the *variance*


$$V(S) = \frac{1}{|S|} \sum_{i \in S} (y_i - \bar{y}_S)^2 = \frac{1}{|S|^2} \sum_{i,j \in S} \frac{1}{2} (y_i - y_j)^2.$$

- ▶ the criterion is the **variance reduction due to the split**:

$$V(I) - \frac{|I_L|^2}{|I|^2} V(I_L) - \frac{|I_R|^2}{|I|^2} V(I_R).$$

- ▶ **Intuition:** maximal if data is homogeneous left and right of the split (then $V(I_L) = V(I_R) = 0$)

Other examples



► **Figure: quad trees**¹¹

¹¹Finkel, Bentley, *Quad trees: a data structure for retrieval on composite keys*, Acta Informatica, 1974

Other examples, ctd.




Figure: comparison-based splits¹²

¹²Haghiry et al., *Comparison-based random forests*, ICML, 2018

When to stop?

- ▶ usually, many direction to try: CART reduces to a random subset of directions
- ▶ also possible to specify T a max height for the tree
- ▶ other strategy: grow the trees to the full extent, and then **pruning**
- ▶ one possibility = **reduced error pruning**
- ▶ starting at the leaves, each node replaced by its most common class
- ▶ if prediction accuracy does not change, ditch the node
- ▶ **Remark:** error computations on a *validation set*

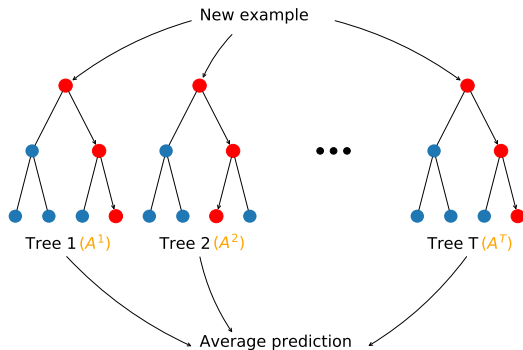
Recap

- ▶ **What happens by default when we invoke the function `sklearn.tree.DecisionTreeClassifier`?** let us look at least at the main options
- ▶ `criterion` is set to Gini → we are using CART trees
- ▶ `splitter` is set to best → looking at the best split at each step
- ▶ `max_depth` is None → splitting until leaves are pure or contain less than `min_samples_split`
- ▶ `min_samples_split` = 2
- ▶ `max_features` is None → no max number of features, log could be a reasonable choice if we have many features
- ▶ `max_leaf_nodes`: None → many leaves, we could also restrict this
- ▶ `min_impurity_decrease` = 0 → continues to split even if very small gain

3.2. Random forests

Random forests

- ▶ one possible problem with tree classifiers: overfitting
- ▶ **Solution:** train many trees and aggregate the prediction
- ▶ **Classification:** majority vote
- ▶ **Regression:** return the mean



Bagging

- ▶ **Additional idea:**¹³ train each tree on a random subsample of the data
- ▶ usual strategy = **bagging**
- ▶ bagging means bootstrap aggregation
- ▶ sample *with replacement* a proportion αn of the training data
- ▶ train the tree classifier on this subset
- ▶ resample for each tree

¹³Breiman, *Random forests*, Machine Learning, 2001

Recap

What happens by default when we invoke the function

`sklearn.ensemble.RandomForestClassifier`?

- ▶ `n_estimators = 100` (T in our notation)
- ▶ `criterion = 'Gini'` → we are using CART trees
- ▶ `max_depth = None` → trees are grown until leaves are pure
- ▶ `max_features = auto` → \sqrt{d} features considered
- ▶ `bootstrap = True` → taking subsamples of the data, but since `max_samples` is set to `None` actually sampling the whole data

4. Boosting

Introduction

- ▶ **Important:** classification setting, $\mathcal{Y} = \{-1, +1\}$
- ▶ **Idea:** aggregate many classifiers together, then majority voting:


$$H(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right),$$

where h_t are classifiers and α_t weights

- ▶ weak classifier = barely better than random guessing
- ▶ **Examples:** linear classifier, small trees,...
- ▶ **Question:** how do we decide which weights to put?
- ▶ different strategies, different algorithms¹⁴


¹⁴Schapire, Freund, *Boosting: foundations and algorithms*, MIT Press, 2012

Non-linearly separable datasets



► **Figure:** moons datasets from sklearn

Weak classifier



► **Figure:** moons datasets from sklearn

4.1. Adaboost

Introduction

- ▶ we first look at AdaBoost¹⁵ (*short for adaptative boosting*)
- ▶ AdaBoost maintains a *distribution* D_t over time
- ▶ start with uniform distribution, then **increase the weights of misclassified examples**
- ▶ at each step, we pick a classifier that minimizes the *weighted error*

$$\begin{aligned}\varepsilon_t &:= \mathbb{P}_{i \sim D_t} (h_t(x_i) \neq y_i) \\ &= \sum_{i=1}^n D_t(i) \cdot \mathbb{1}_{h_t(x_i) \neq y_i}.\end{aligned}$$

- ▶ adjust the weights by multiplying by a quantity depending on ε_t , larger than one if misclassified, smaller if correctly classified

¹⁵Freund and Schapire, *A decision-theoretic generalization of on-line learning and an application to boosting*, Journal of computer and system science, 1997

AdaBoost

Algorithm 1: AdaBoost algorithm

Input: n training examples $(x_1, y_1), \dots, (x_n, y_n)$ where $x_i \in \mathcal{X}$ and $y_i \in \{-1, 1\}$

Initialize the distribution to $D_1(i) = \frac{1}{n}$

for $t = 1$ **to** T **do**

 Train weak learner using distribution D_t

 Get weak hypothesis $h_t : \mathcal{X} \rightarrow \{-1, 1\}$

h_t minimizes the weighed error $\varepsilon_t := \mathbb{P}_{i \sim D_t} (h_t(x_i) \neq y_i)$. Set $\alpha_t := \frac{1}{2} \log \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right)$

 Update, for $i = 1 \dots n$,

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i, \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i. \end{cases}$$

end

Result: final classifier $H(x) := \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$.

Exercise

Exercise: With the notation of the previous slide,

1. show that

$$\mathbb{P}_{i \sim D_{t+1}} (h_t(x_i) \neq y_i) = \frac{\sqrt{\varepsilon_t(1 - \varepsilon_t)}}{Z_t}.$$

2. show that

$$Z_t = 2\sqrt{\varepsilon_t(1 - \varepsilon_t)}.$$

Deduce that

$$\mathbb{P}_{i \sim D_{t+1}} (h_t(x_i) \neq y_i) = \frac{1}{2}.$$

Correction of the exercise

1. We write

$$\begin{aligned}\mathbb{P}_{i \sim D_{t+1}}(h_t(x_i) \neq y_i) &= \sum_{i=1}^n \mathbb{1}_{h_t(x_i) \neq y_i} \cdot D_{t+1}(i) && \text{(total expectation)} \\ &= \sum_{i=1}^n \mathbb{1}_{h_t(x_i) \neq y_i} \cdot \frac{D_t(i)}{Z_t} \cdot e^{\alpha_t} && \text{(definition of } D_{t+1}) \\ &= \frac{1}{Z_t} \sum_{i=1}^n \mathbb{1}_{h_t(x_i) \neq y_i} \cdot D_t(i) \cdot \sqrt{\frac{1 - \varepsilon_t}{\varepsilon_t}} && \text{(definition of } \alpha_t) \\ &= \frac{1}{Z_t} \cdot \sqrt{\frac{1 - \varepsilon_t}{\varepsilon_t}} \cdot \varepsilon_t = \frac{\sqrt{\varepsilon_t(1 - \varepsilon_t)}}{Z_t}.\end{aligned}$$

Correction of the exercise, ctd.

2.

$$\begin{aligned} Z_t &= \sum_{\substack{i=1 \\ h_t(x_i) \neq y_i}}^n D_t(i) \cdot e^{\alpha_t} + \sum_{\substack{i=1 \\ h_t(x_i) = y_i}}^n D_t(i) \cdot e^{-\alpha_t} \\ &= e^{\alpha_t} \sum_{i=1}^n \mathbb{1}_{h_t(x_i) \neq y_i} \cdot D_t(i) + e^{-\alpha_t} \left(1 - \sum_{i=1}^n \mathbb{1}_{h_t(x_i) \neq y_i} \cdot D_t(i) \right) \\ &= \sqrt{\frac{1 - \varepsilon_t}{\varepsilon_t}} \cdot \varepsilon_t + \sqrt{\frac{\varepsilon_t}{1 - \varepsilon_t}} \cdot (1 - \varepsilon_t) \\ Z_t &= 2\sqrt{\varepsilon_t(1 - \varepsilon_t)}. \end{aligned}$$

