# How to work with Text

# Natural Language Processing

NLP = Natural Language Processing

- NLP is a field in machine learning with the ability of a computer to understand, analyze, manipulate, and potentially generate human language.

- Useful to all big data applications
- Especially useful for mining knowledge about people's behavior, attitude and opinions
- Express directly knowledge about our world: Small text data are also useful!

# Main NLP Task

- Document classification
  - Associate a label to a document

- Sentiment analysis
  - Associate a label to a sentence

  - Sentiment analysis is the automated process of analyzing text data and classifying opinions as negative, positive or neutral. Usually, besides identifying the opinion, these systems extract attributes of the expression e.g.:
    - Polarity: if the speaker express a positive or negative opinion,
    - Subject: the thing that is being talked about,
    - Opinion holder: the person, or entity that expresses the opinion

# Sentiment analysis

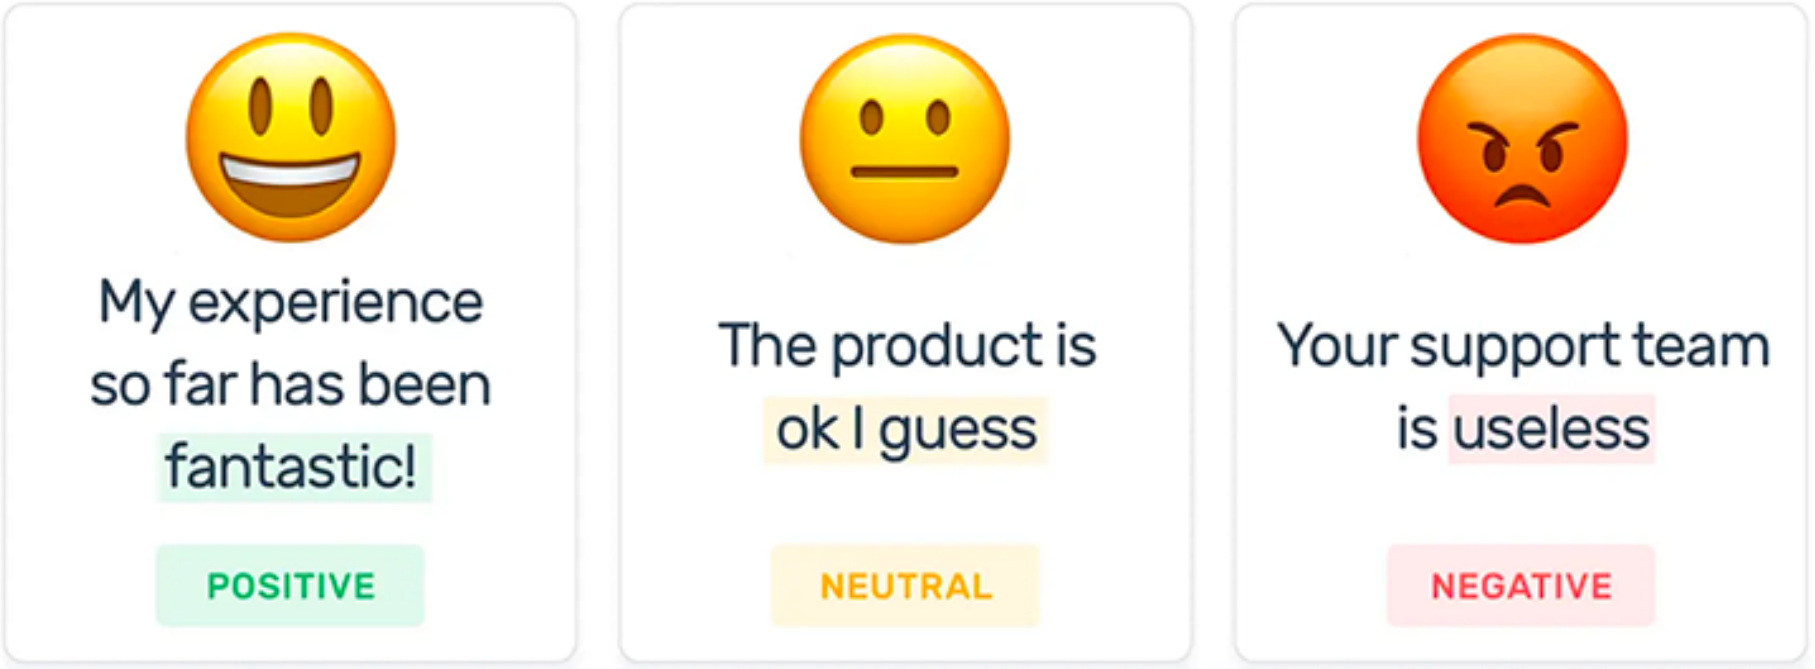

# Main NLP Task

- Part of speech tagging
  - Associate a label to a word

  - Part of Speech Marking (POS), also known as grammatical marking or word category disambiguation, involves marking a word according to its relationship to adjacent and related words in a sentence, word group or paragraph.
  - A simplified form of this definition is commonly taught to school-age children in the identification of words such as nouns, verbs, adjectives, adverbs, and so on.

# Part of speech tagging (POS)

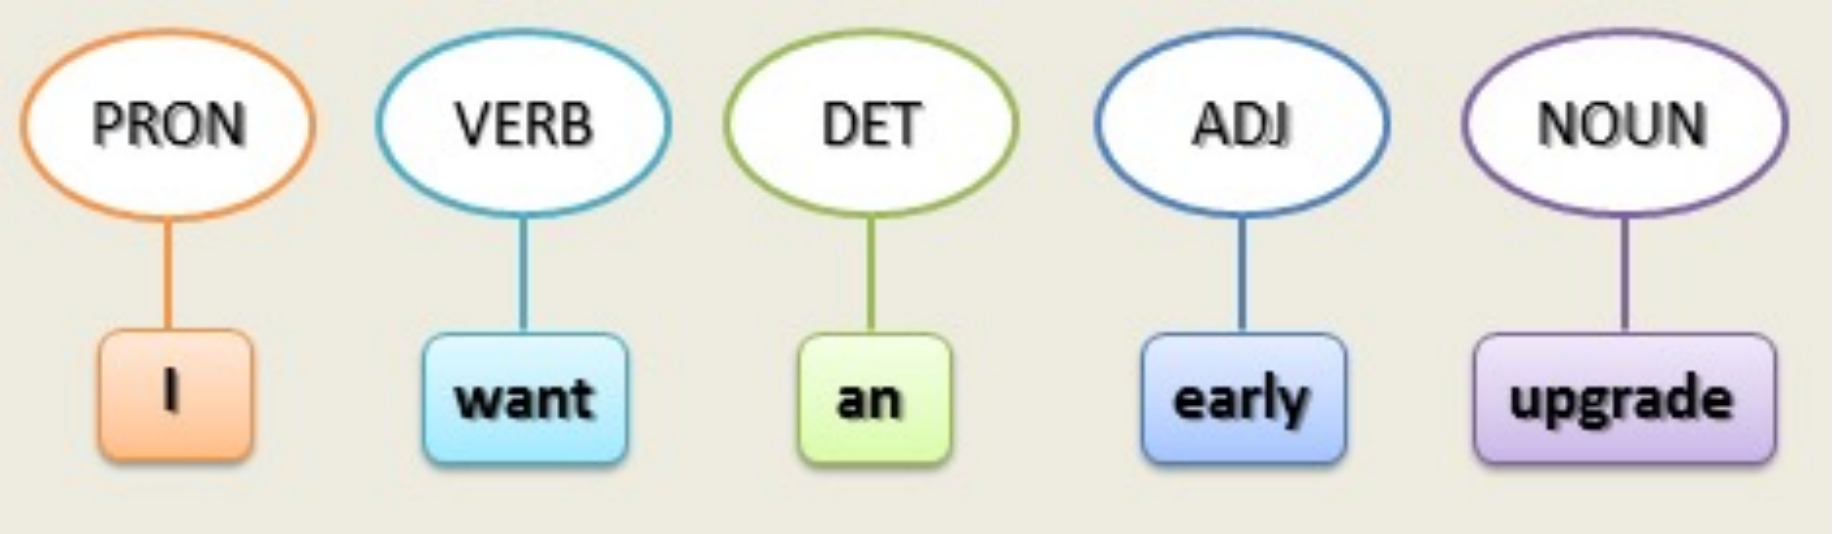

# Main NLP Task

- Naming entity recognition
  - Associate a label to a word

  - Named-entity recognition (NER) seeks to locate and classify named entities in text into pre-defined categories such as the names of persons, organizations, locations, expressions of times, quantities, monetary values, percentages, etc.
  - In a second step, we also try to extract relationships between previously recognized entities

# Naming Entity Recognition (NER)

In fact, the Chinese NORP market has the three CARDINAL most influential names of the retail and tech space – Alibaba GPE , Baidu ORG , and Tencent PERSON (collectively touted as BAT ORG ), and is betting big in the global AI GPE in retail industry space . The three CARDINAL giants which are claimed to have a cut-throat competition with the U.S. GPE (in terms of resources and capital) are positioning themselves to become the 'future AI PERSON platforms'. The trio is also expanding in other Asian NORP countries and investing heavily in the U.S. GPE based AI GPE startups to leverage the power of AI GPE . Backed by such powerful initiatives and presence of these conglomerates, the market in APAC AI is forecast to be the fastest-growing one CARDINAL , with an anticipated CAGR PERSON of 45% PERCENT over 2018 - 2024 DATE .

To further elaborate on the geographical trends, North America LOC has procured more than 50% PERCENT of the global share in 2017 DATE and has been leading the regional landscape of AI GPE in the retail market. The U.S. GPE has a significant credit in the regional trends with over 65% PERCENT of investments (including M&As, private equity, and venture capital) in artificial intelligence technology. Additionally, the region is a huge hub for startups in tandem with the presence of tech titans, such as Google ORG , IBM ORG , and Microsoft ORG .

# Two books...

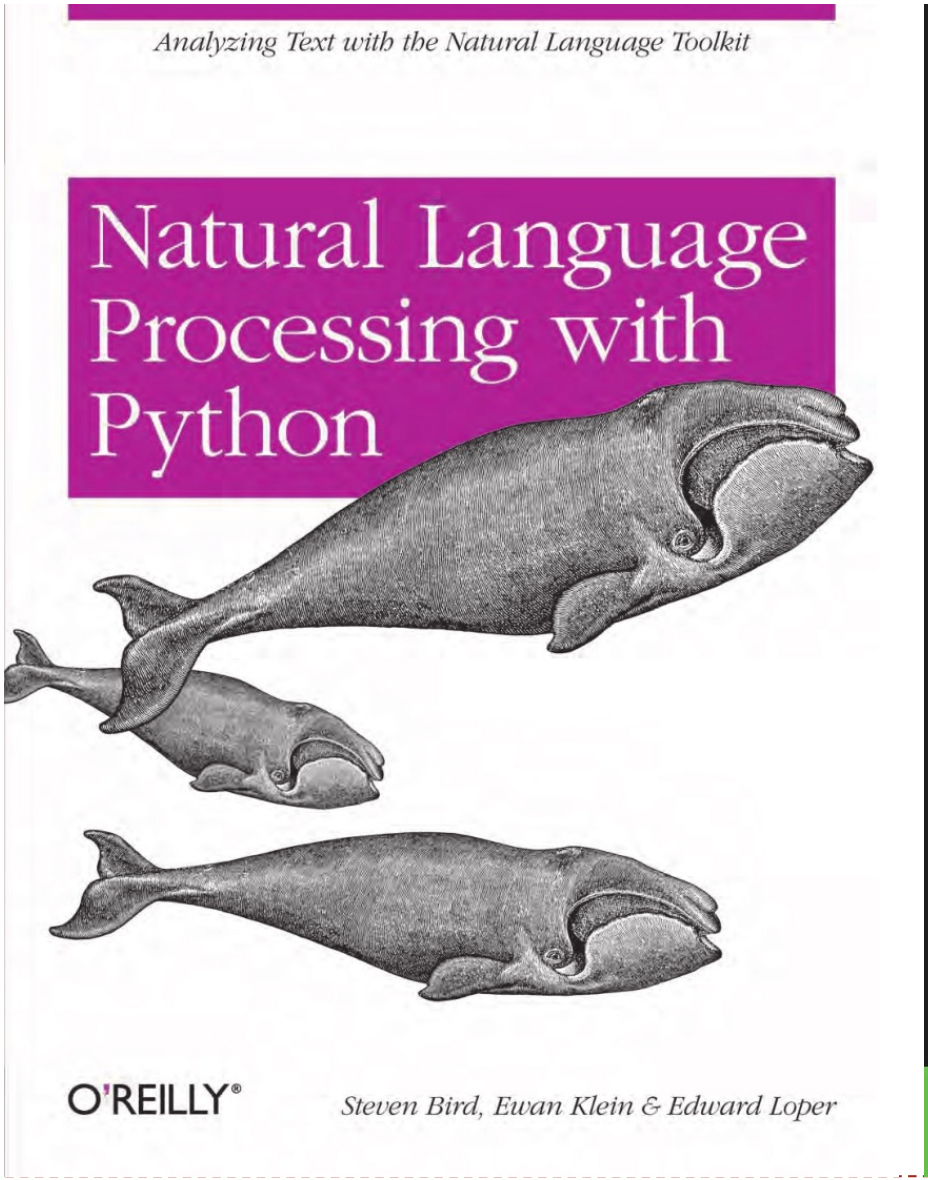
Analyzing Text with the Natural Language Toolkit
Natural Language Processing with Python
O'REILLY®
Steven Bird, Ewan Klein & Edward Loper

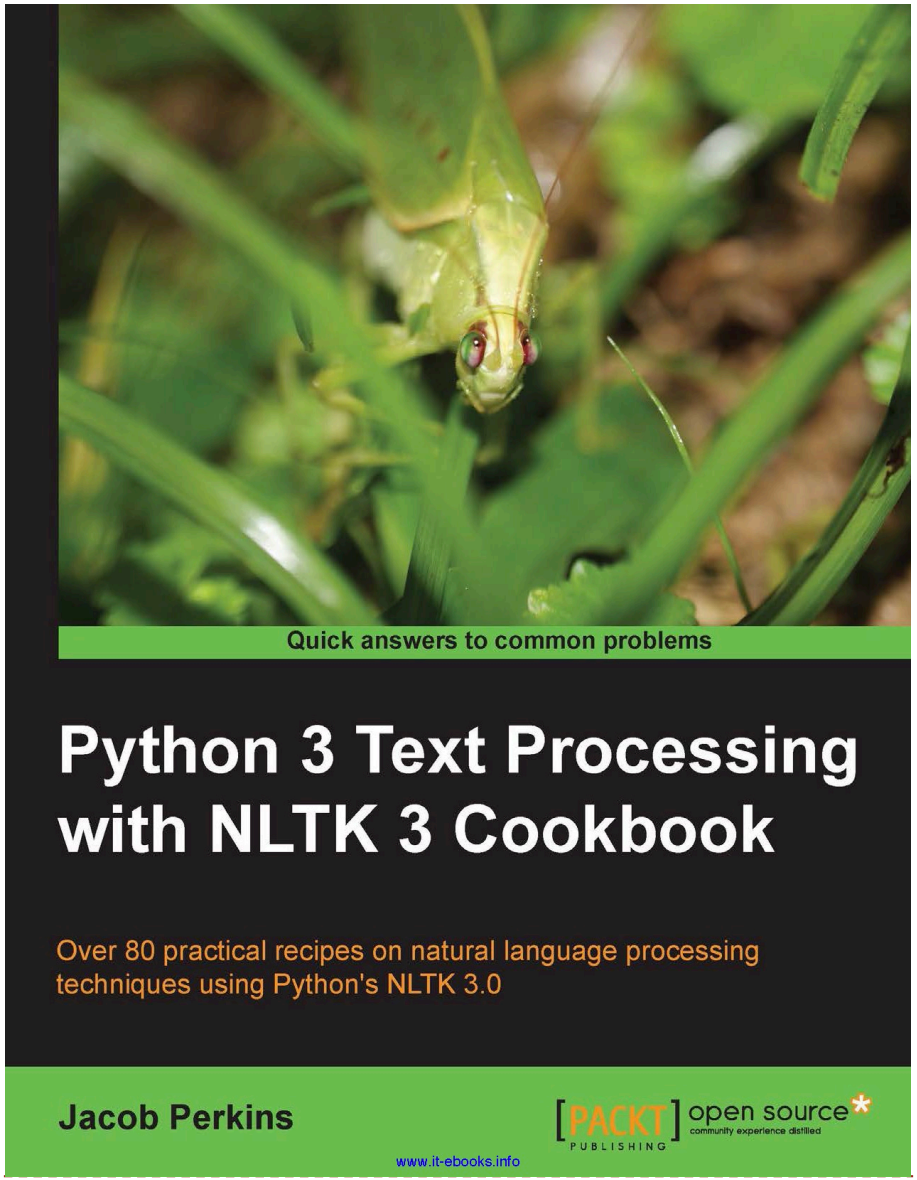
Quick answers to common problems
Python 3 Text Processing with NLTK 3 Cookbook
Over 80 practical recipes on natural language processing techniques using Python's NLTK 3.0
Jacob Perkins
[PACKT] open source
PUBLISHING
www.it-ebooks.info

# And one main library

- Natural Language ToolKit ( NLTK)
  - http://www.nltk.org/
  - A comprehensive Python library for natural language processing and text analytics
  - Originally designed for teaching
  - also adopted in the industry for research and development due to its usefulness and breadth of coverage
- NLTK is often used for rapid prototyping of text processing programs
- Demos of select NLTK functionality and production-ready APIs are available at http://text-processing.com
- In Python: use nltk library (http://www.nltk.org/book/)
  - `!pip install -U nltk` in the Jupyter Notebook
  - or `conda install -c conda-forge nltk`

# The NLTK Pipeline

# Main step for an NLP pipeline

- Text normalization
  - The set of operation depend of the task
- Main goal
  - Replace the list of chars (the original text) by a list of tokens
  - Normalize some representation : data, phone number
  - Try to reduce the vocabulary size
    - Use only lower character
    - Spelling Correction
    - Lemmatization (or Stemming)
    - Replace by synonyms (semantic reduction)

# Tokenization

- Tokenization: process of splitting a string into a list of pieces (tokens).
  - A token is a piece of whole
    - A char is a token in a word
    - A word is a token in a sentence
    - A sentence is a token in paragraph
- Token != Words
  - Tokens
    - Substrings
    - Only structural
    - Data
  - Words
    - Objects
    - Contains a 'sense'
    - Meaning
- Not always an easy task
  - '' between space ? One or two words
    - cats!
    - San Francisco

# Python tokenization

- By default, work with english language
  - from nltk.tokenize import sent_tokenize
  - sent_tokenize(a_text)
    - Return a list of sentences

  - from nltk import word_tokenize
  - word_tokenize(a_text)
    - Return a list of word
    - Ponctuation is a word

- For other language, you have to load specific tokenizer
  - from nltk.data import load
  - spanish_tokenizer = load("tokenizers/punkt/PY3/spanish.pickle")
  - spanish_tokenizer.tokenize(a_text)
- Avalaible tokenizers are on ~/nltk_data/tokenizers/punkt/PY3

# Text normalization

- Text normalization is the process of transforming text into a single canonical form
- Text normalization requires being aware of
  - What type of text is to be normalized
  - how it is to be processed afterwards
  - There is no all-purpose normalization procedure
- Easy part
  - Put the text in lower case
    - lower_text = text.lower()
      - Text → "This is the first sentence. A gallon of milk in the U.S. …"
      - Lower text → "this is the first sentence. a gallon of milk in the u.s. …"
  - Suppress '.' in acronym:
    - U.S.A → USA
- Difficult part
  - Phone number: +33 6 10 20 30 40 or +33.(0)6.10.20.30.40
  - Date: 11/01/2018 or 2018-01-11
  - Etc.

# Stop word removal

- Stopwords are common words that generally do not contribute to the meaning of a sentence
  - Examples: the, as, a
- Most search engines will filter out stopwords fom search queries in order to save space in their index
- NLTK comes with a stopword corpus
  - from nltk.corpus import stopwords
  - stopwords.words('english')
    - ['i', 'me', 'my', 'myself', 'we', 'our', …
  - stopwords.words('french')
    - ['au', 'aux', 'avec', 'ce', 'ces', …
- General use
  - tokens = word_tokenize(text)
    - ['this', 'is', 'the', 'first', 'sentence', '.', 'a', …
  - [t for t in tokens if t not in english_stopwords]
    - ['first', 'sentence', '.',

# Some experiment with stop word removal

- from nltk.corpus import stopwords
- print(stopwords.words('english'))

['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', 'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', 'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some', 'such', 'no', 'nor', **'not'**, 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', 've', 'y', 'ain', **'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn', "wouldn't"**]

# Some experiment with stop word removal

- Let's imagine you are asked to create a model that does sentiment analysis of product reviews. The dataset is fairly small that you label it your self. Consider a few reviews from the dataset.
  1. *The product is really very good. — POSITIVE*
  2. *The products seems to be good. — POSITIVE*
  3. *Good product. I really liked it. — POSITIVE*
  4. *I didn't like the product. — NEGATIVE*
  5. *The product is not good. — NEGATIVE*
- You performed preprocessing on data and removed all stopwords. Now, let us look what happens to the sample we selected above.
  1. *product really good. — POSITIVE*
  2. *products seems good. — POSITIVE*
  3. *Good product. really liked. — POSITIVE*
  4. *like product. — NEGATIVE ?*
  5. *product good. — NEGATIVE ?*
- Scary, right?

# Reduce word forms
## Stemming and Lemmatisation

### Stemming

- The term "stem" generally refers to a crude heuristic process that cuts off the end of words in the hope of achieving a reduction in the forms of a word
- This often results in the removal of suffixes and sometimes prefixes
  - cats, cat → cat
  - looked → look
- May result in an unknown word

### Lemmatization

- Lemmatization generally involves doing things correctly using vocabulary and morphological analysis of words
- Reduce inflections or variant forms to base form
  - am, are, is → be
  - Jack's → Jack
- Always ends up with a known word

# Reduce word forms
# Stem and Lem in Python

- Stemming: many algorythms
  - For example, you can use Porter algorithm
    - porter = nltk.PorterStemmer()
    - stemming_form = porter.stem(token)

- Lemmatization:
  - WNlemma = nltk.WordNetLemmatizer()
  - Lemma_form = WNlemma.lemmatize(token)

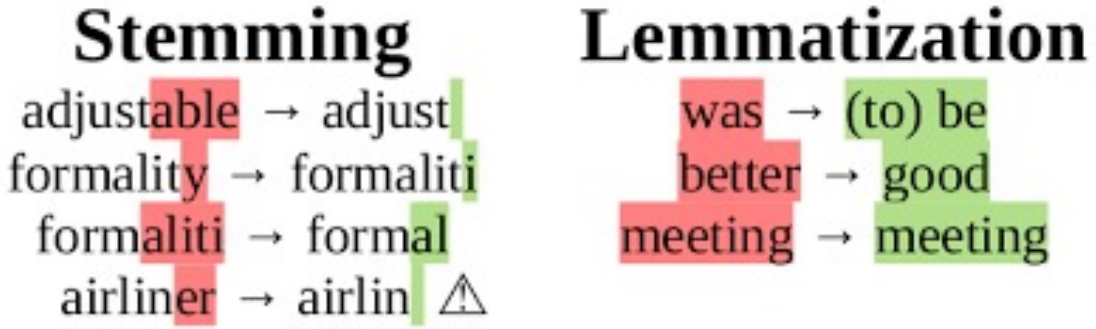

# Stem vs Lem

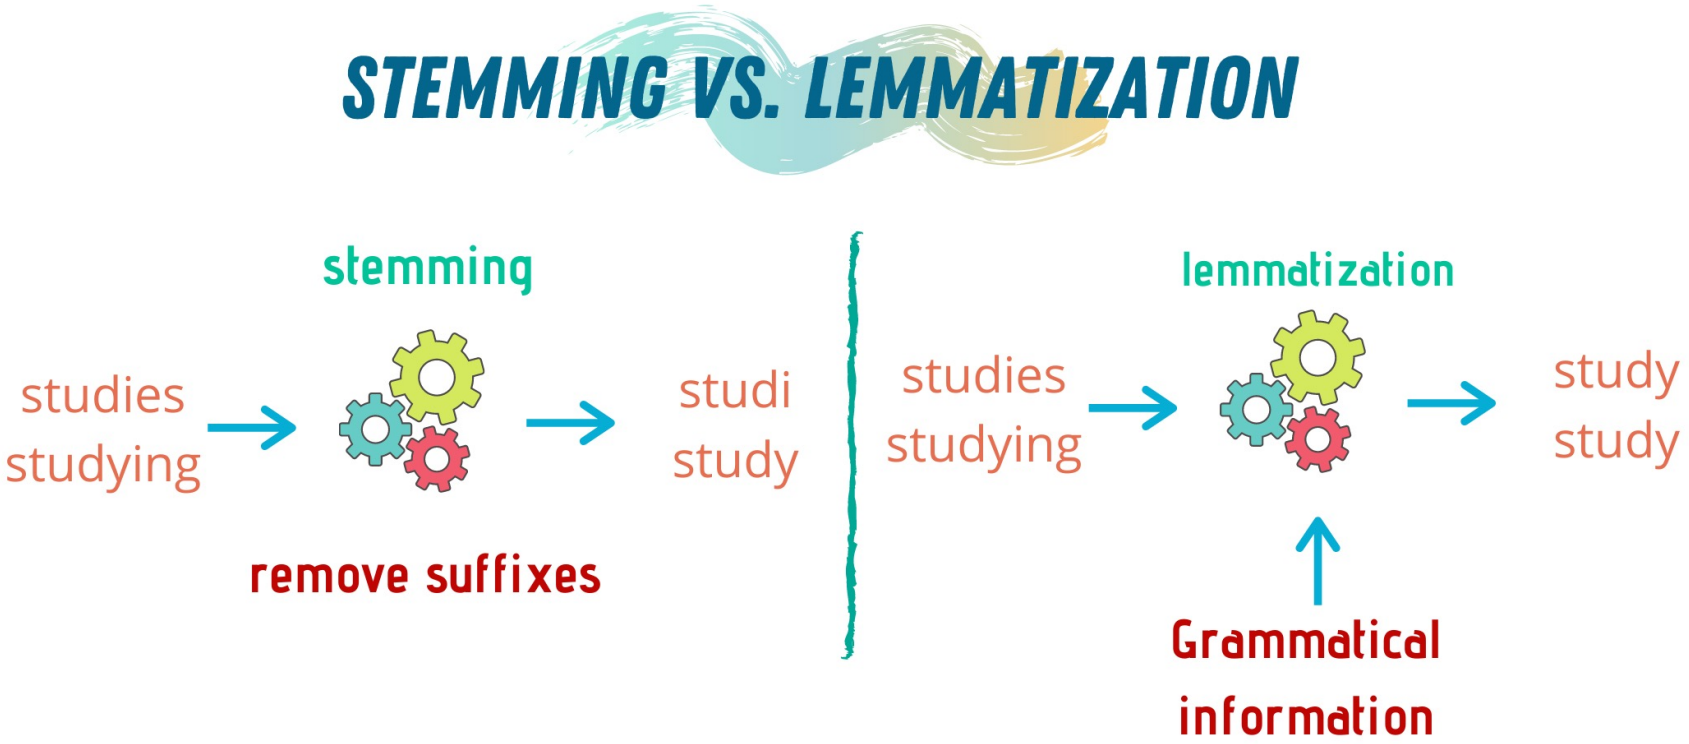
STEMMING VS. LEMMATIZATION
stemming
studies
studying
remove suffixes
studi
study
lemmatization
studies
studying
Grammatical
information
study
study

# Summary
# Text normalization in Python

Tokenization

- Usually depends on the language, sometimes on the task
  - tokens = nltk.word_tokenize(sentence)
  - sentences = nltk.sent_tokenize(paragraph)

Normalization

- Use only one form: lowercase for example
  - lower_text = text.lower()

Reduce vocabulary

- Stop word removal
  - from nltk.corpus import stopwords
  - tokens = [t for t in tokens if t not in stopwords.words('english')]
- Stemming: user Porter algorithm - Several algorithms available
  - porter = nltk.PorterStemmer()
  - stemming_form = porter.stem(token)
- Lemmatization - Several algorithms available
  - WNlemma = nltk.WordNetLemmatizer()
  - lemma_form = WNlemma.lemmatize(token)

# Features extraction

# Bag of Word

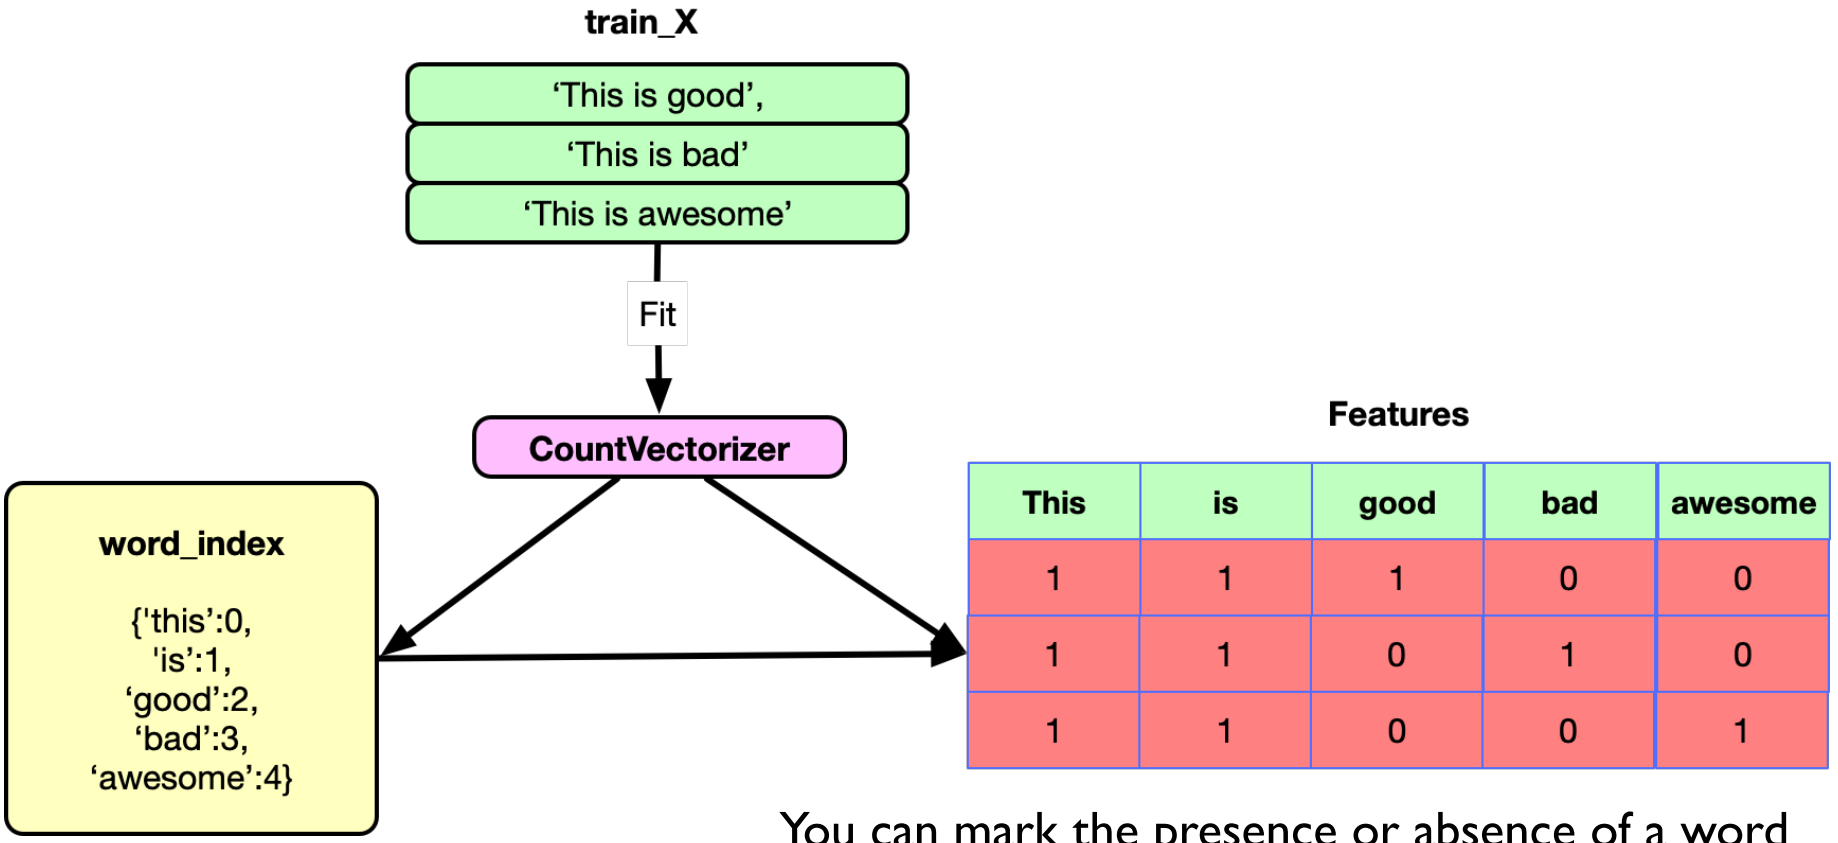


You can mark the presence or absence of a word

Or

You can count the number of times it appears.

# TF-IDF

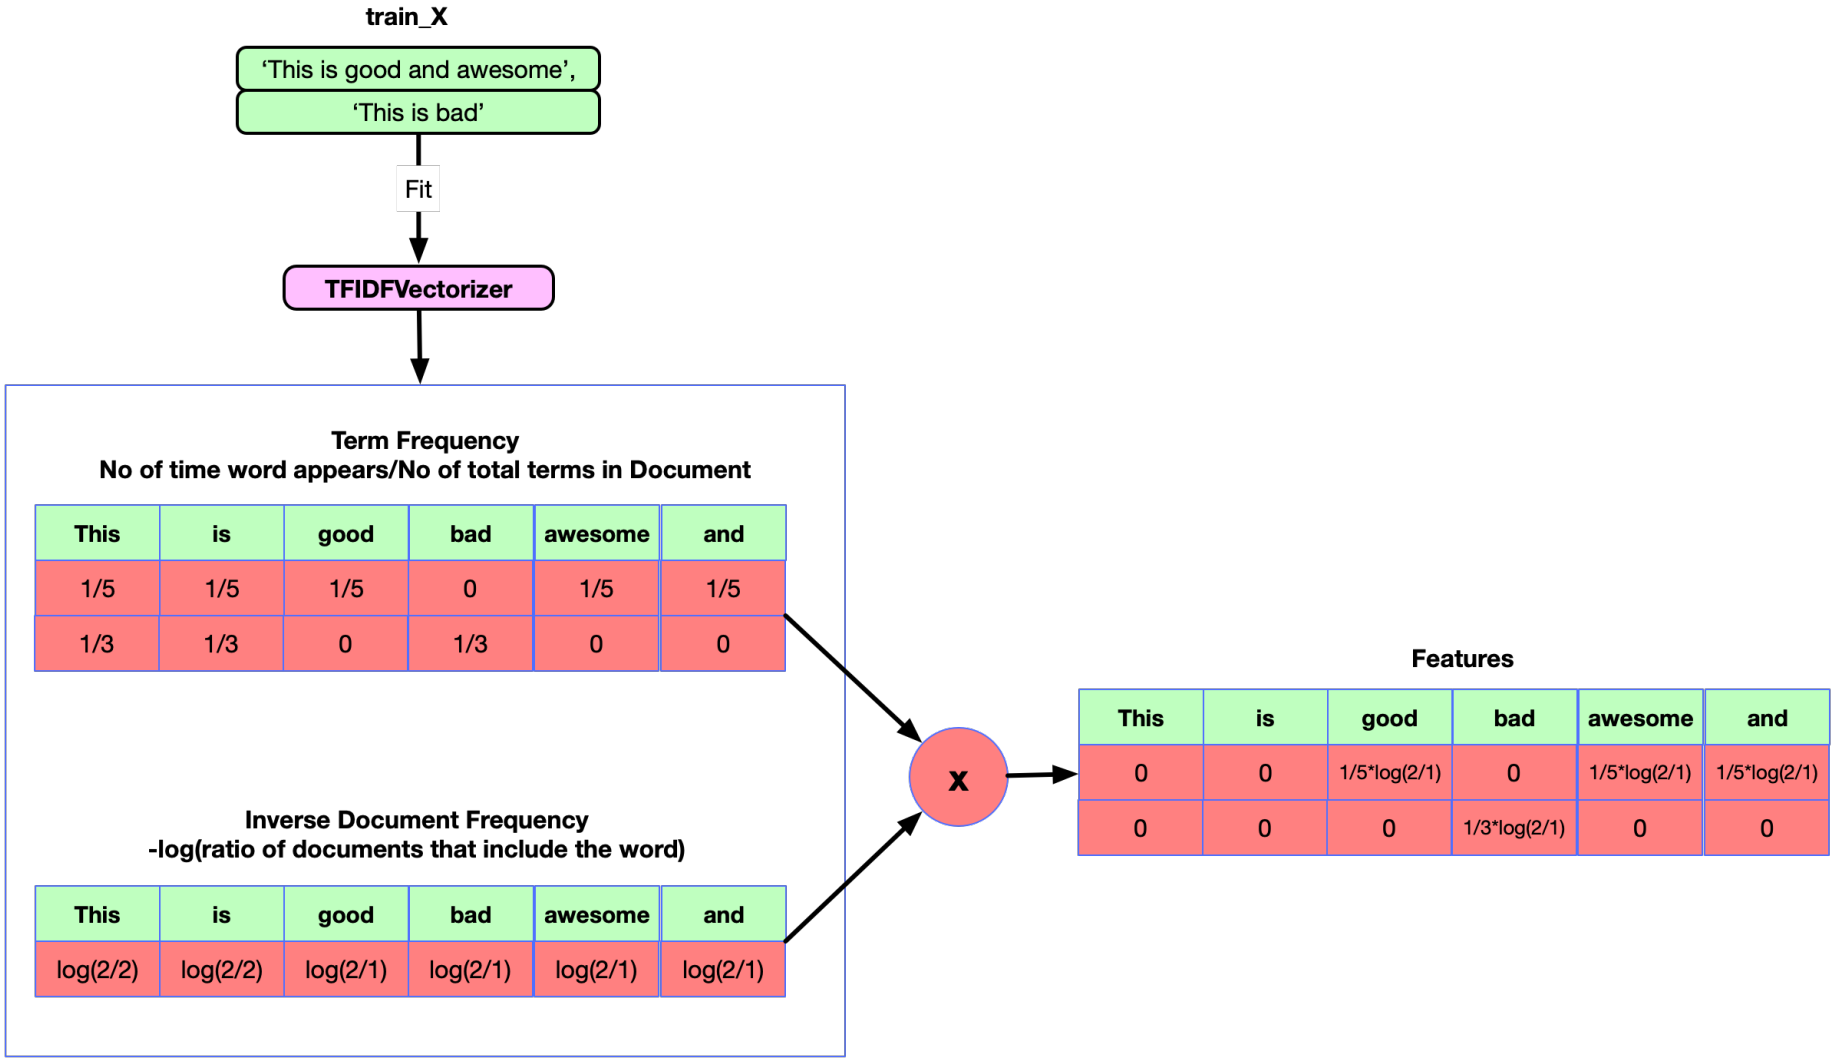
train_X
'This is good and awesome',
'This is bad'
Fit
TFIDFVectorizer
Term Frequency
No of time word appears/No of total terms in Document
This | is | good | bad | awesome | and
1/5 | 1/5 | 1/5 | 0 | 1/5 | 1/5
1/3 | 1/3 | 0 | 1/3 | 0 | 0
Inverse Document Frequency
-log(ratio of documents that include the word)
This | is | good | bad | awesome | and
log(2/2) | log(2/2) | log(2/1) | log(2/1) | log(2/1) | log(2/1)
X
Features
This | is | good | bad | awesome | and
0 | 0 | 1/5*log(2/1) | 0 | 1/5*log(2/1) | 1/5*log(2/1)
0 | 0 | 0 | 1/3*log(2/1) | 0 | 0

# N-gram

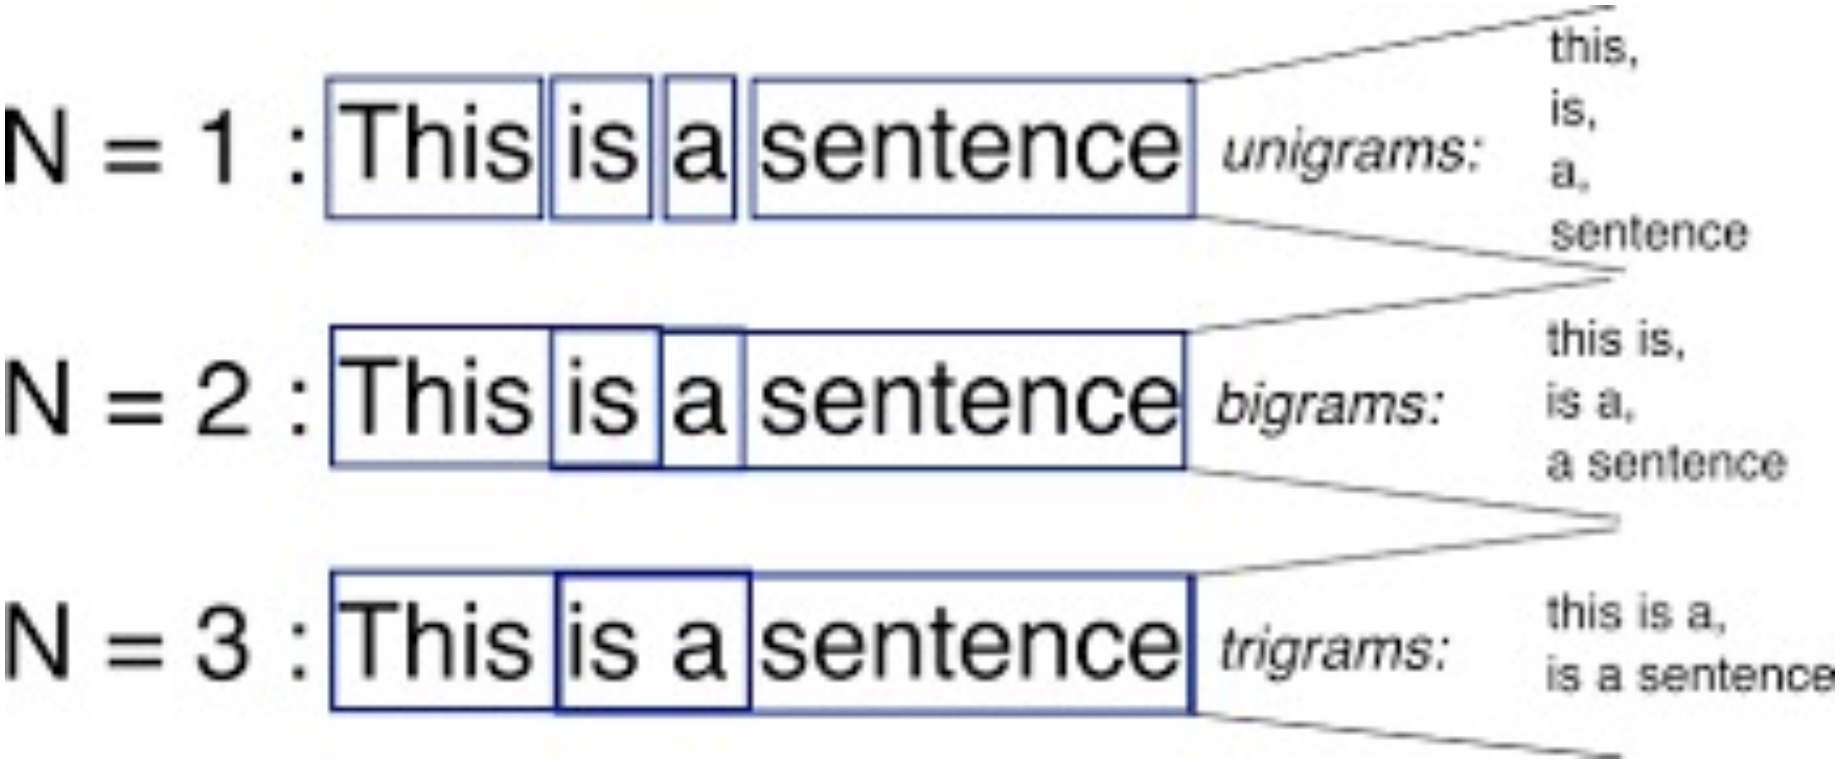
N = 1 : This is a sentence
unigrams:
this,
is,
a,
sentence
N = 2 : This is a sentence
bigrams:
this is,
is a,
a sentence
N = 3 : This is a sentence
trigrams:
this is a,
is a sentence

# N-grams

- *N*-gram of size
  - 1 is referred to as a "unigram";
  - 2 is a "bigram" (or, less commonly, a "digram
  - 3 is a "trigram"
- N-grams in NKTK:
  - from nltk import ngrams
  - For the word "hello"
    - set(ngrams("hello", 2))
      - {('e', 'l'), ('h', 'e'), ('l', 'l'), ('l', 'o')}
  - For the sentence "The cow jumps over the moon"
    - set(ngrams(nltk.word_tokenize("The cow jumps over the moon"), 2))
      - {('The', 'cow'), ('cow', 'jumps'), ('jumps', 'over'), ('over', 'the'), ('the', 'moon')})

# Lab

- Observe the impact of the preprocessing for sentiment analysis task with a Logistic Regression Classifier
  1. Use TfidfVectorizer and a classifier in order to build a pipeline with default parameters
  2. Modifier the pipeline in order to observe the impact of the following parameters on the « output » of tfidf vectorization and on the sentiment analysis task
     1. binary and use_idf (combination of this two parameters)
     2. lowercase
     3. stop_words (try None, 'english' and your own list)
     4. ngram_range (try (1,1) and (3,3) and (1,3)
     5. max_features (try None, 100, 1000)
     6. max_df and min_df (try 1, 50, 1.0, 0.5)
     7. Build you how preprocessor with stemming or lemming

# A approfondir en TD

- Transformation pipeline
  - With or without sklearn pipelines
- Train / Val / Test separation
- BOW representation, without and with ngrams
- Hyper-parameters search
- Confusion matrix and metrics for classification
  - ACC / Recall / Precision / F1
  - ROC-AUC curves
- Balanced and unbalanced classification problems
  - Modifier le target pour avoir 2 classes non balancée en regroupant les classe 2-4 et en laissant 1 tout seul
  - What impact on the metrics